



# Efficient Graph Analytics on Real World Graphs

Joana M. F. da Trindade

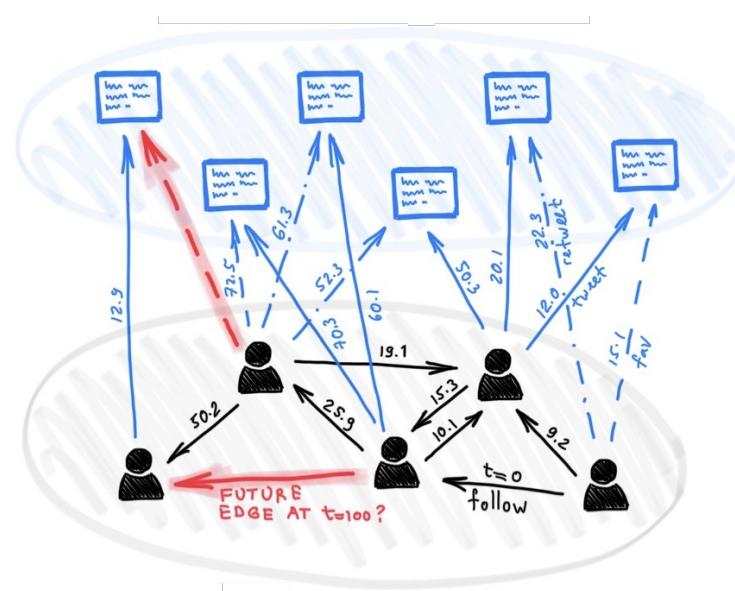
[jmf@csail.mit.edu](mailto:jmf@csail.mit.edu)

# “Graphs are everywhere” they say

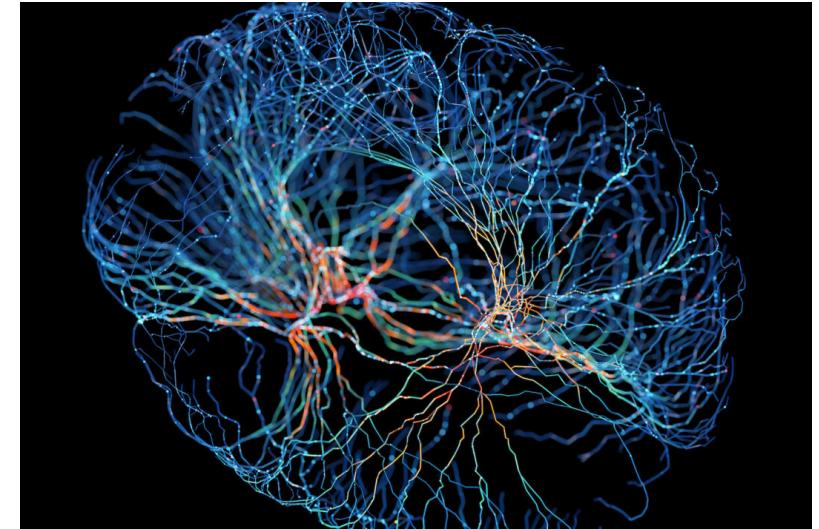


e.g.,

- hyperlink 2012: 128B edges
- **MSFT data infra: 16B edges / week**

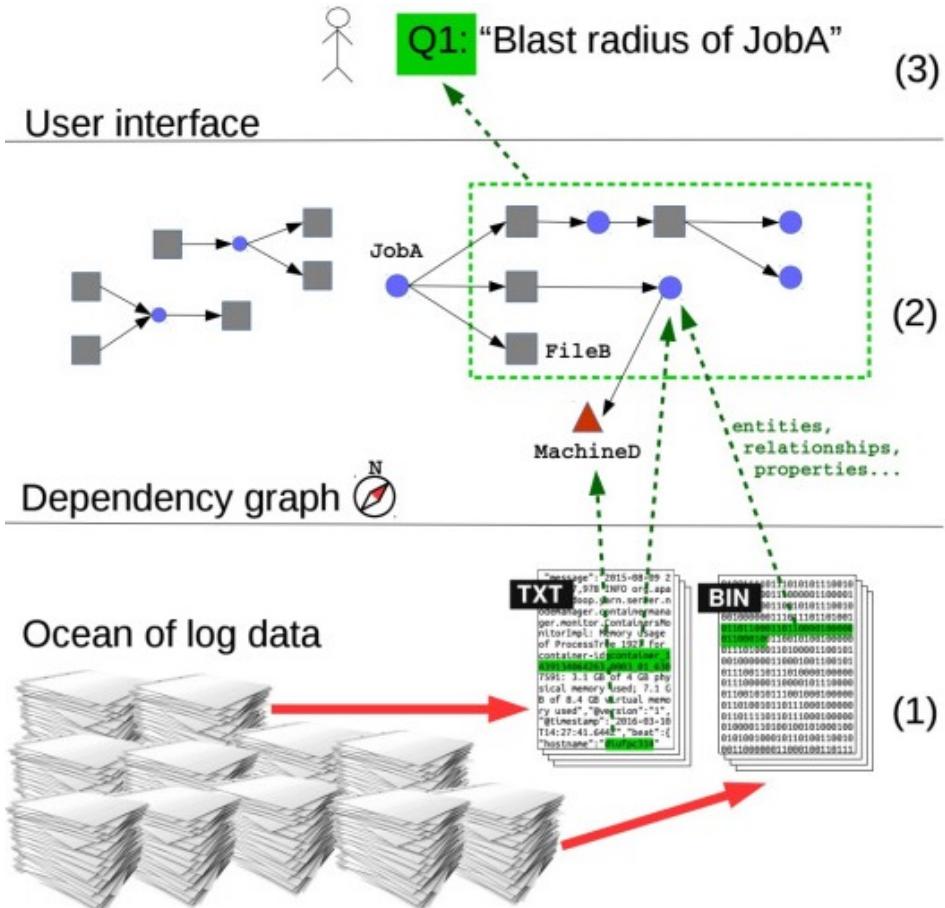


“Deep learning on dynamic graphs” ([blog.twitter.com](http://blog.twitter.com))



“Brain Network Mapping Study Challenges Basis for Psychiatric Distinctions” ([genengnews.com](http://genengnews.com))

# Cloud computing companies extract large graphs from production systems logs every day



Queries include:

- k-hop paths (dependency analysis)
- subgraph traversals
- temporal paths

Graphs are large and skewed:

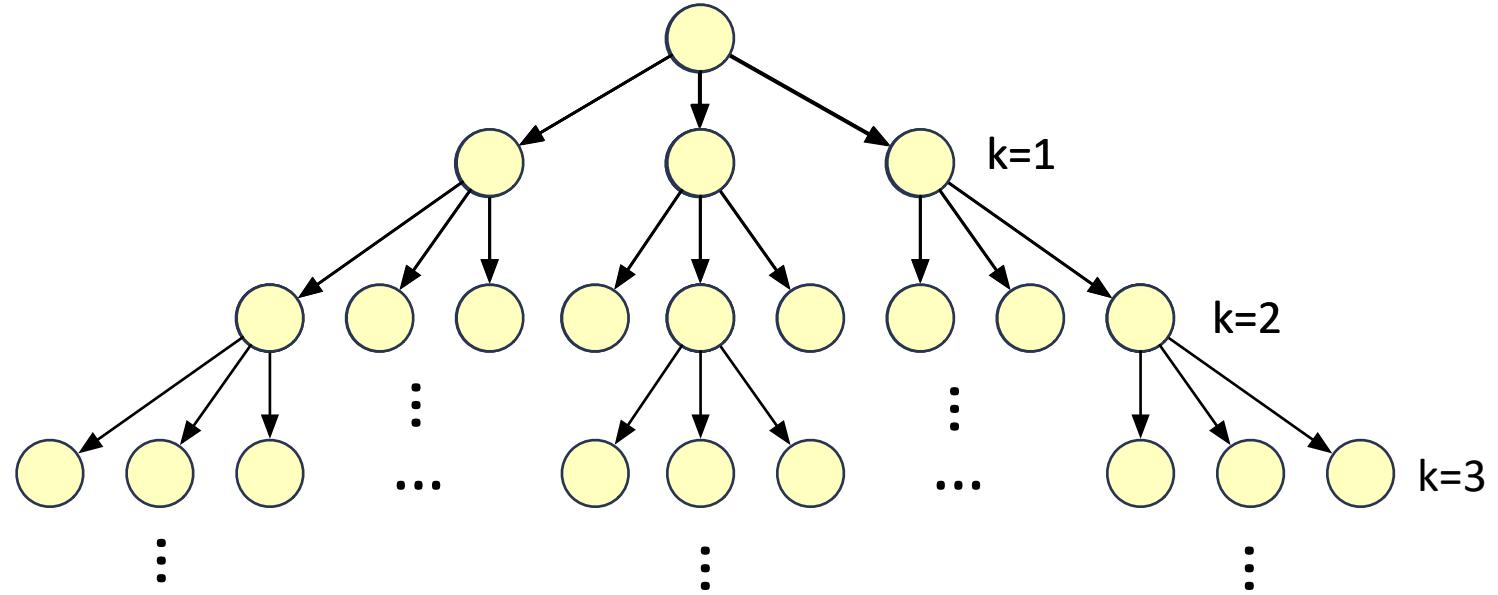
- billions of edges every week
- both degree and temporal skew

Problem: on large graphs,  $k$ -hop path queries can be expensive even for small values of  $k$

Worst-case **complexity**

is proportional to:

$$\left( \max_{v \in V} \deg(v) \right)^k$$



Weekly data lineage graph of Microsoft production cluster had 3.2 Billion vertices, 16.4 Billion edges, and on-disk footprint of 10+ Terabytes.

Problem: on large graphs,  $k$ -hop path queries can be expensive even for small values of  $k$

Distributed relational engine at Microsoft took *hours* to answer it

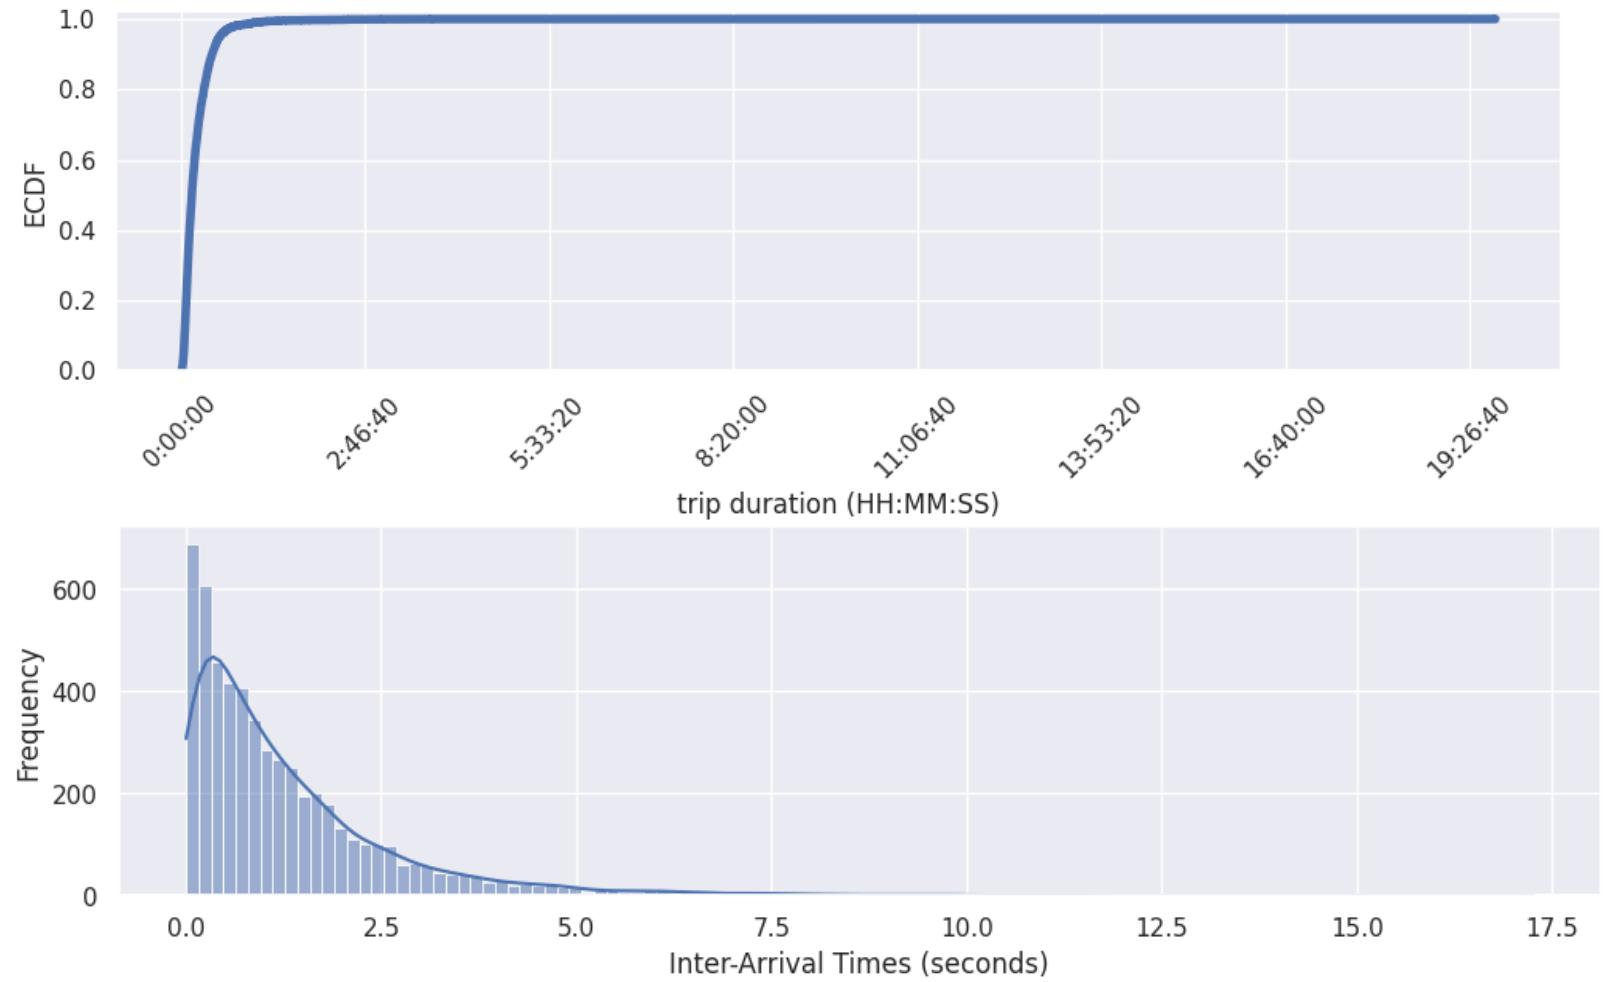
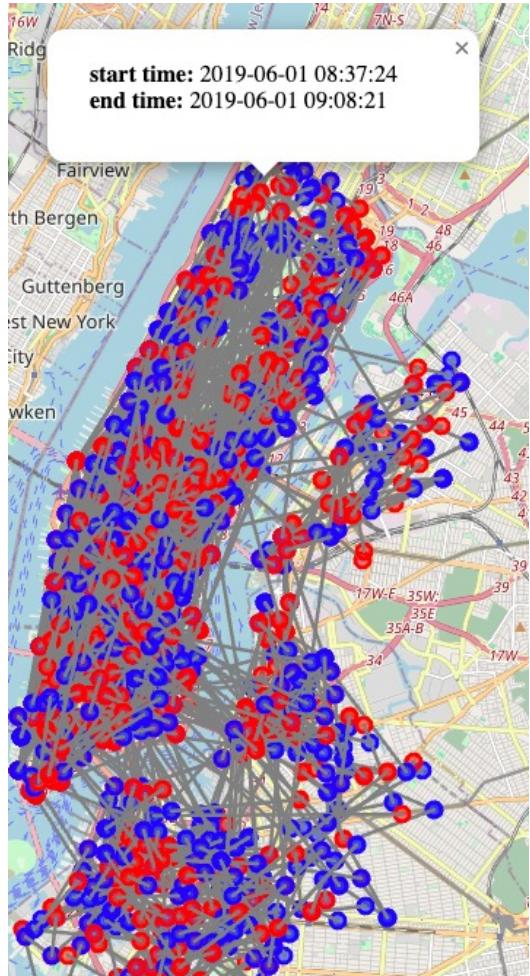
Real-world graphs are skewed; **relational engines produce inefficient plans (recursive self-joins)**

Existing graph databases limited in performance

Distributed engines (e.g., Pregel and Flink) incur lots of messaging overhead

To be used for operational decisions, these queries require short response times (seconds to mins)

# Problem: one size does not fit all



# Problem: one size does *not* fit all

Real-world graph data is skewed in several dimensions

**Degree skew**: caused by **preferential attachment** (“rich get richer”)

**Temporal skew**: human generated event **durations and IAT** follow power-law

Graph query optimizers treat every vertex the same

Assume i.i.d. and/or fail to capture temporal correlations

Temporal graph data offers crucial insights, yet remains unexplored

How can we leverage temporal dynamics for query optimization?

# Key Contributions of this Thesis

How can we speed up expensive graph queries (e.g.,  $k$ -hop)?

- Introduce class of sparser graph views

- Infer materialized graph views that are beneficial for queries

- Rewrite queries in terms of those views

How can we leverage temporal dynamics for query optimization?

- Memory-efficient data representations for temporal graphs

- Selective Indexing: treat large vertices differently than others

- Leverage the above to implement efficient algorithms over temporal graphs

# Thesis Outline

Kaskade: **Graph Views** for Efficient Graph Analytics

*J M F da Trindade, K Karanasos, C Curino, S Madden, J Shun*

*ICDE 2020*

Kairos: Efficient **Temporal Graph Analytics** on a **Single Machine**

*J M F da Trindade, J Shun, S Madden, N Tatbul*

*Under submission to VLDB 2024*

# Kaskade: Key Technical Contributions

## Graph views

Sparser views such as *connectors* and *summarizers* (e.g., filters and aggregates)

## Constraint-based view enumeration

Mine constraints from graph's schema and data's structural properties

Rewrite queries in terms of beneficial views

## Cost model for graph views

View size estimation leverages structural properties

Only most beneficial views are materialized

# Kaskade: Key Technical Contributions

## Graph views

Sparser views such as *connectors* and *summarizers* (e.g., filters and aggregates)

## Constraint-based view enumeration

Mine constraints from graph's schema and data's structural properties

Rewrite queries in terms of beneficial views

~~Cost model for graph views~~

see thesis for details

~~View size estimation leverages structural properties~~

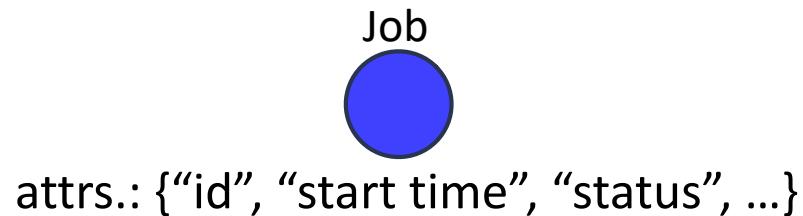
~~Only most beneficial views are materialized~~

# What is the Graph Data Model in Databases?

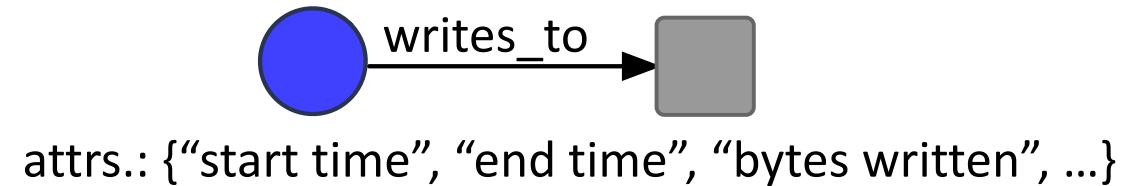
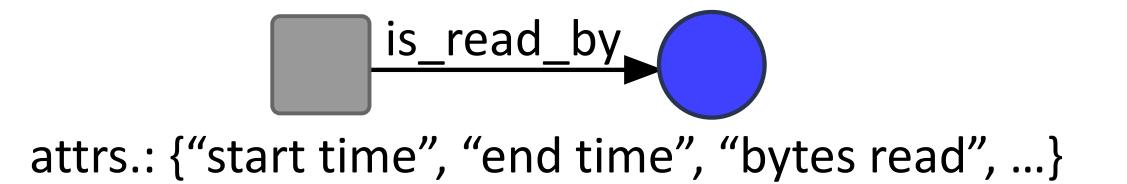
**Schema** defines **types** of vertices and edges.

Each vertex or edge type has a **name**, and a set of **attributes**.

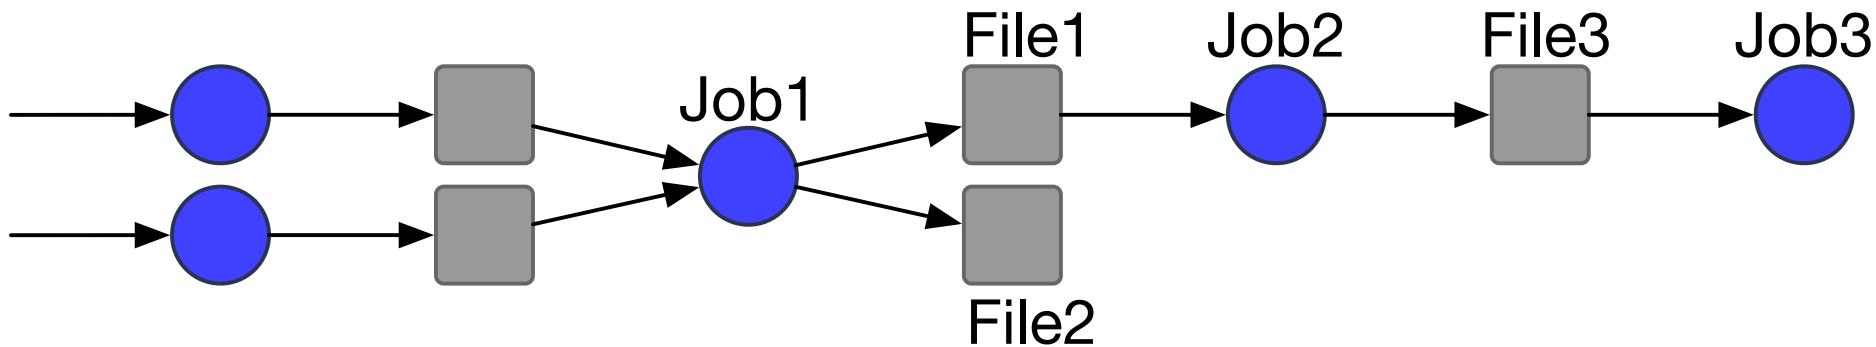
Vertex Types:



Edge Types:



# What is a Graph Query?



*Q1: "Which files has Job1 written to?"*

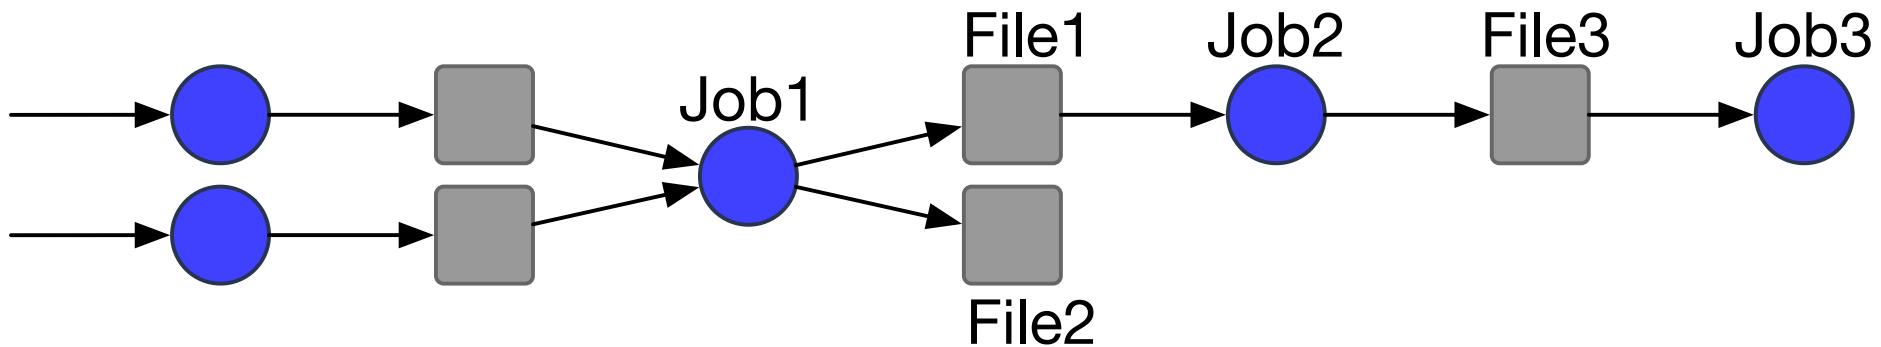
*Q2: "Which jobs have read File1?"*

*Q3: "Which jobs may have indirectly consumed data from File1?"*

*Q4: "Which jobs depend on Job1?"*

*k-hop path queries*

# What is a Graph View?



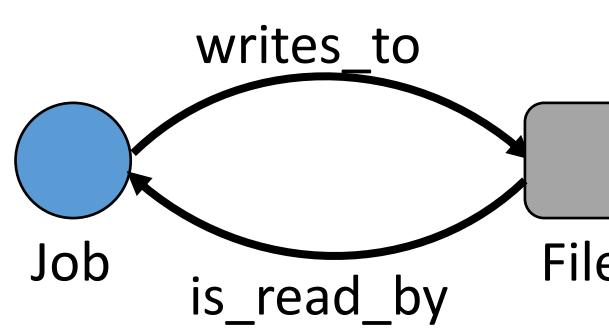
V1: “Which files has Job1 written to?”



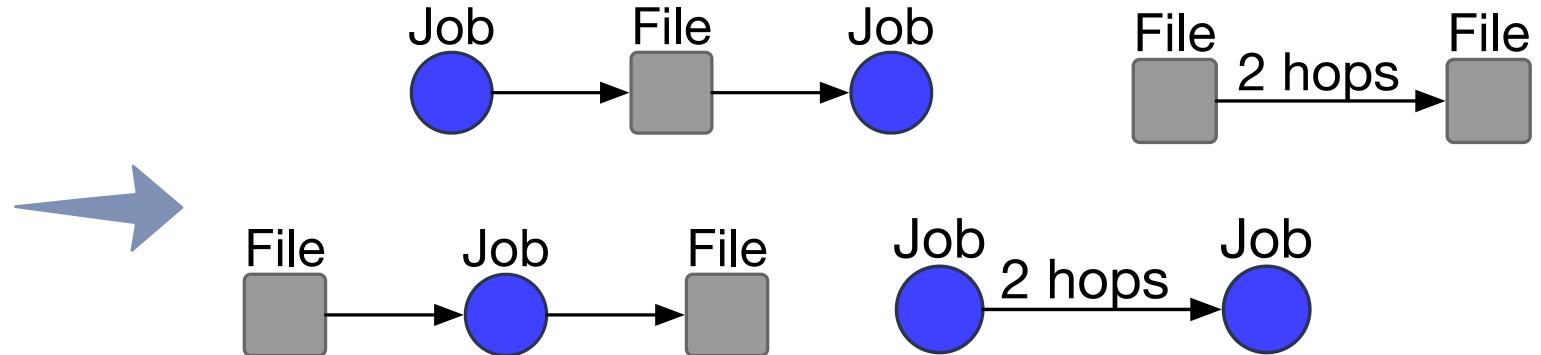
V2: “Which jobs depend on Job1?”



# Key Insight: exploit regular structure in graphs to improve query performance

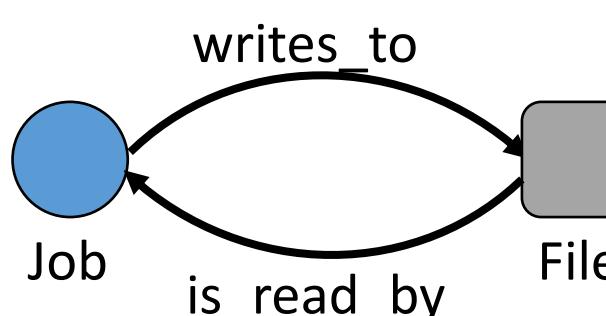


*Schema for data lineage graph.*

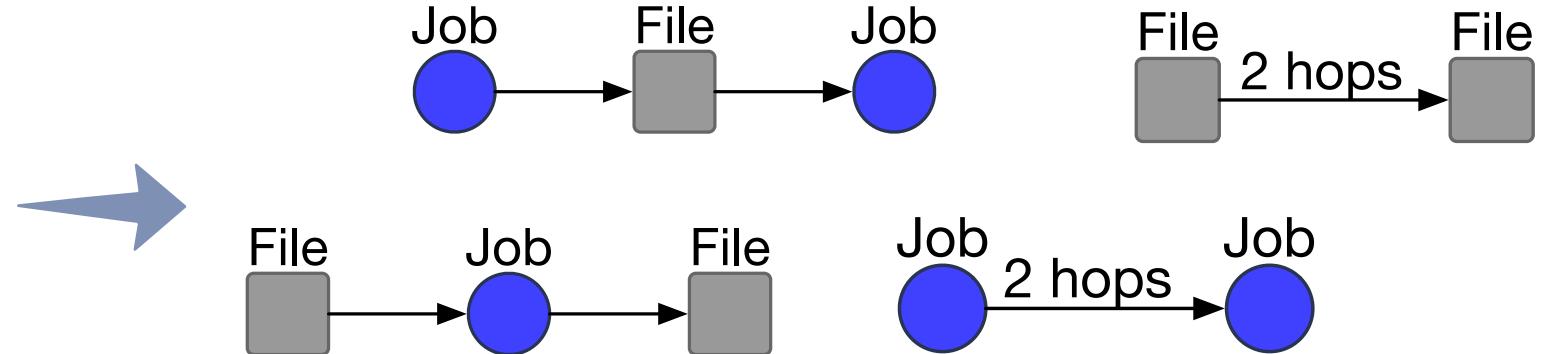


Schema dictates **regularity**: can use it to infer smaller **graph views**.

# Key Insight: exploit regular structure in graphs to improve query performance



*Schema for data lineage graph.*



Schema dictates **regularity**: can use it to infer smaller **graph views**.

Statistic	Example
Cardinality	- No. vertices per Type - No. edges per type
Degree Distribution	- Max out-deg per vertex type - Median out-deg per vertex type

*Example graph data statistics*

Graph data **statistics** tell us when these views have a lot **fewer vertices and edges**, leading to much faster queries.

# Problem Statement

## Inputs

- Graph queries
- Graph schema
- Graph data statistics

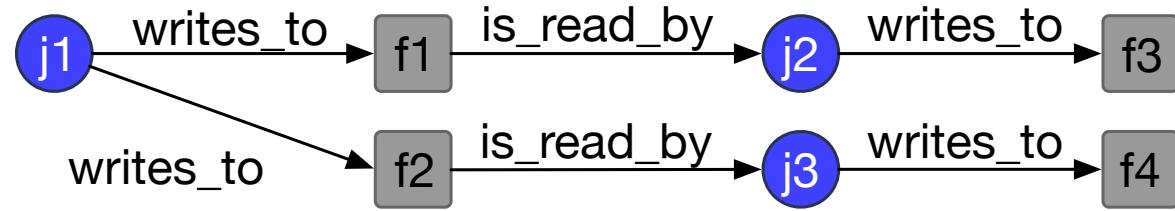
## Goals

Find **graph views** that can be used to answer queries more efficiently.

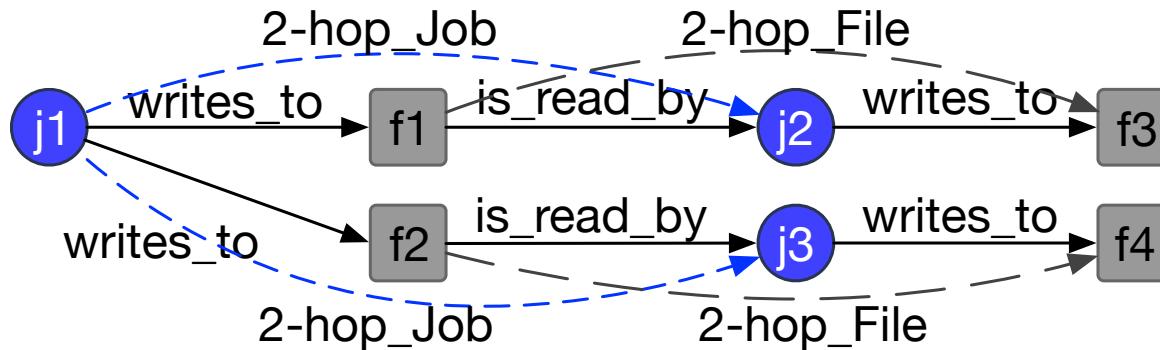
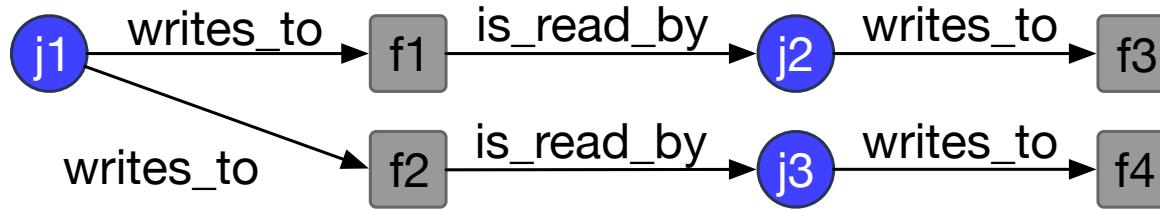
Translate original queries to equivalent version (*query rewriting*).

# Kaskade: Graph Views for Efficient Graph Analytics

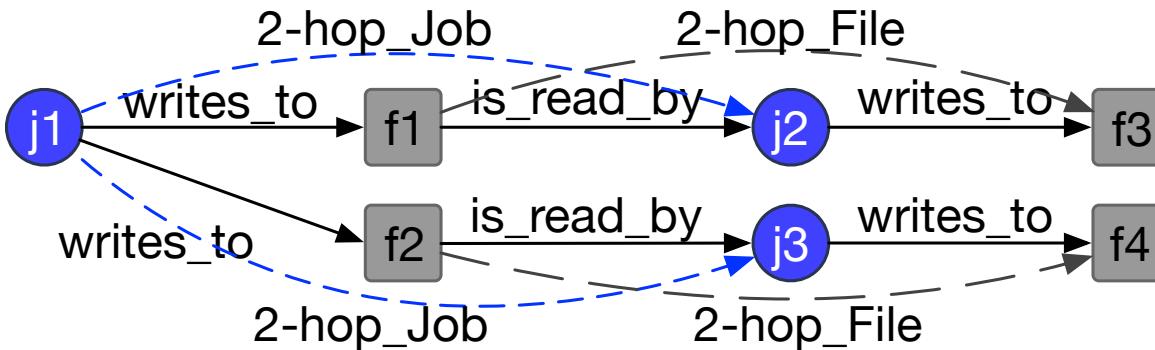
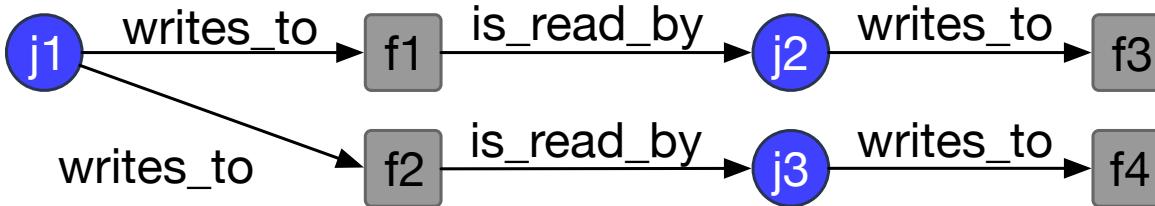
# Graph view example: 2-hop connectors on data lineage graph



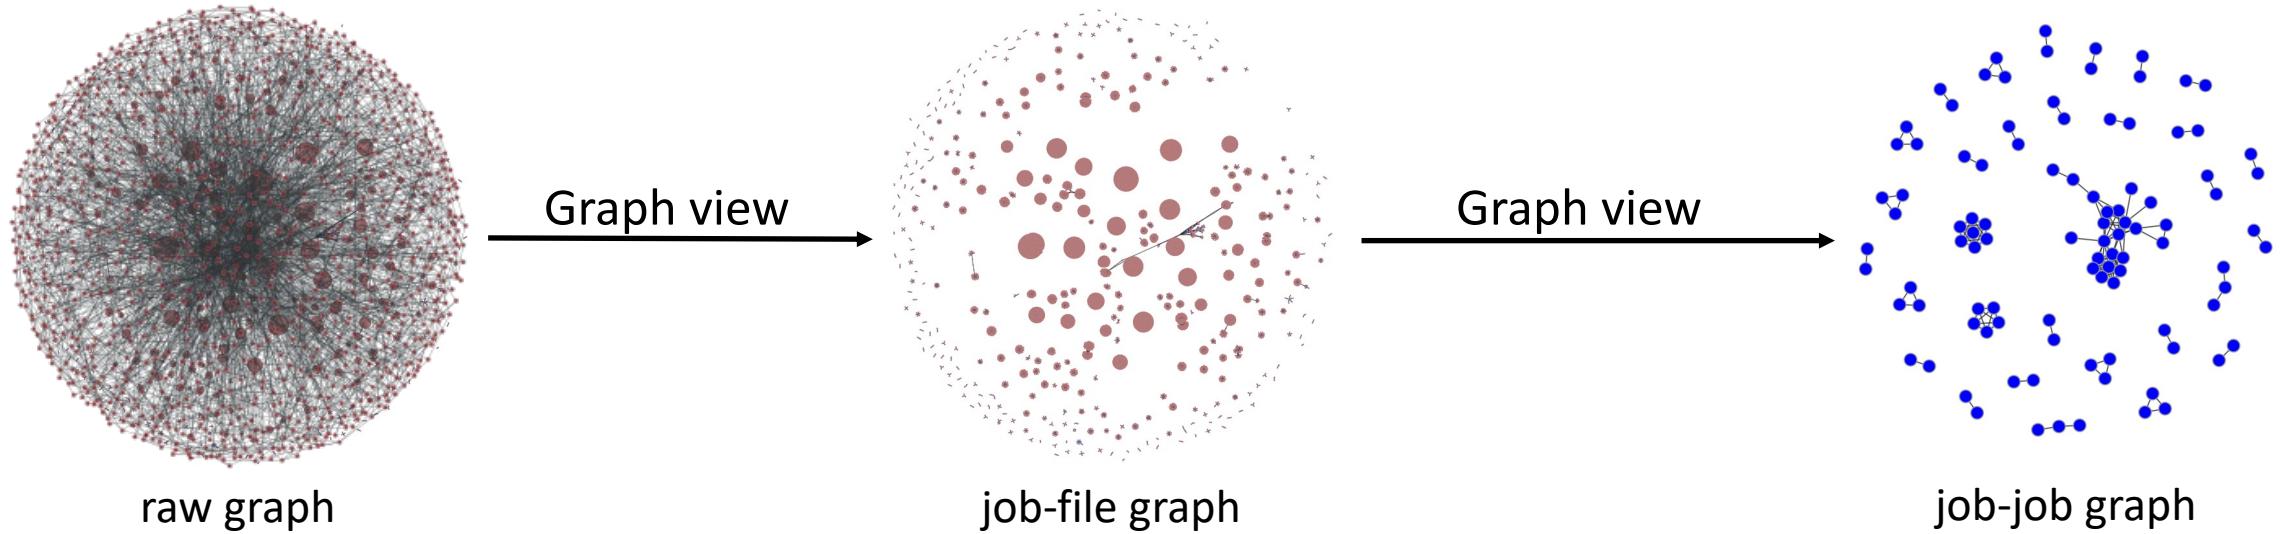
# Graph view example: 2-hop connectors on data lineage graph



# Graph view example: 2-hop connectors on data lineage graph



Example: input data size reduction when using graph views to answer “blast radius” query



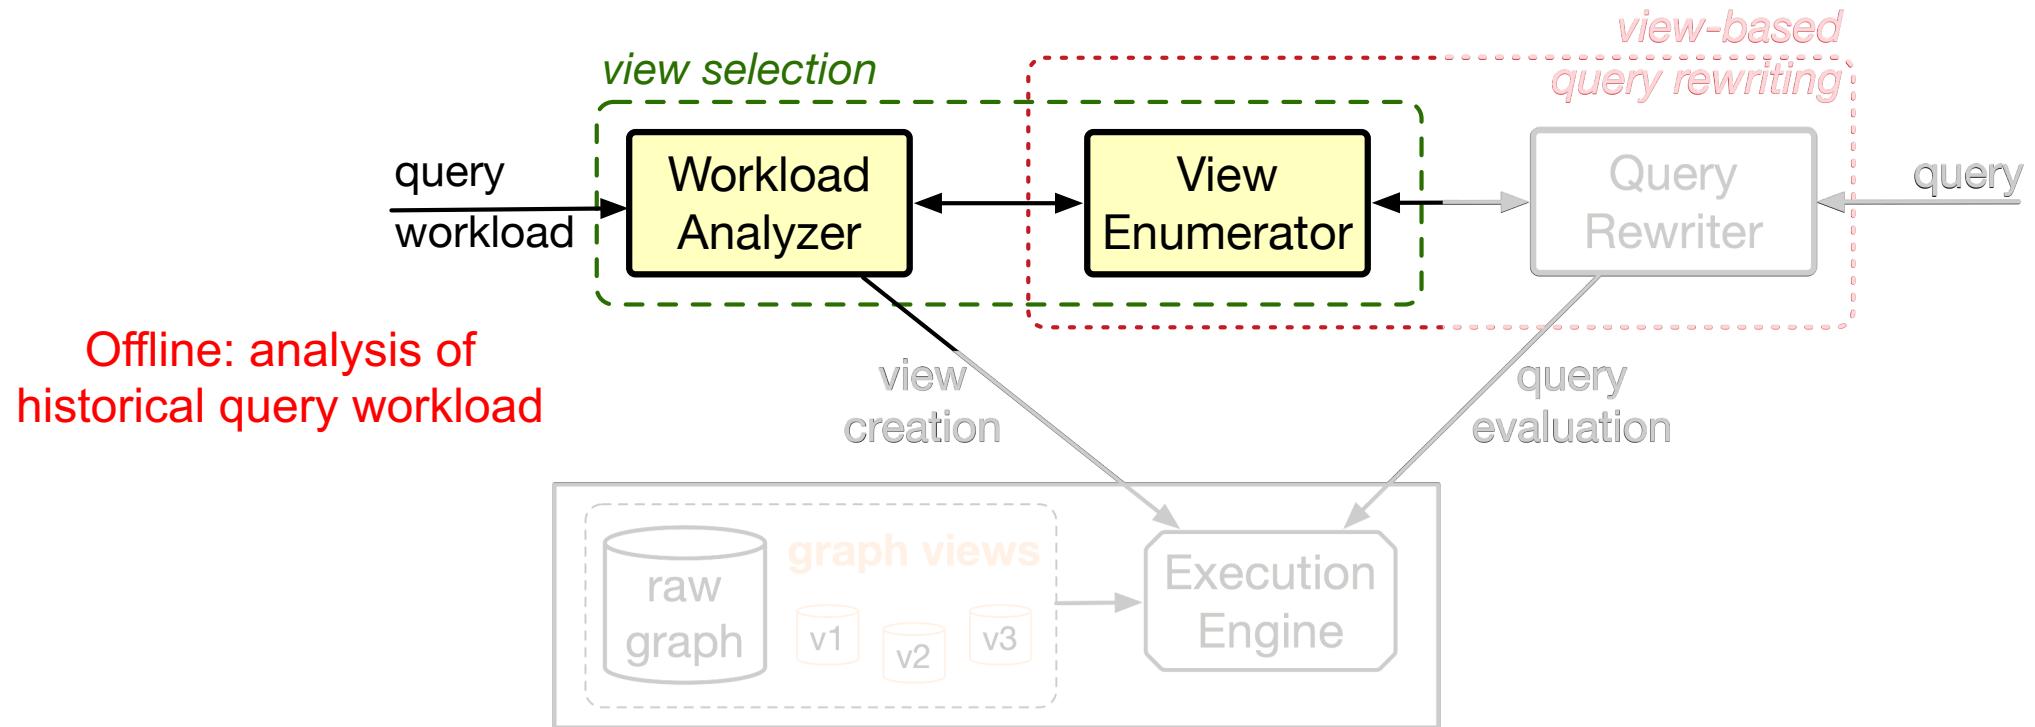
Size:  
 $O(B)$  vertices  
 $O(B)$  edges  
 $O(TB)$  on disk

Size:  
 $O(M)$  vertices  
 $O(M)$  edges  
 $O(GB)$  on disk

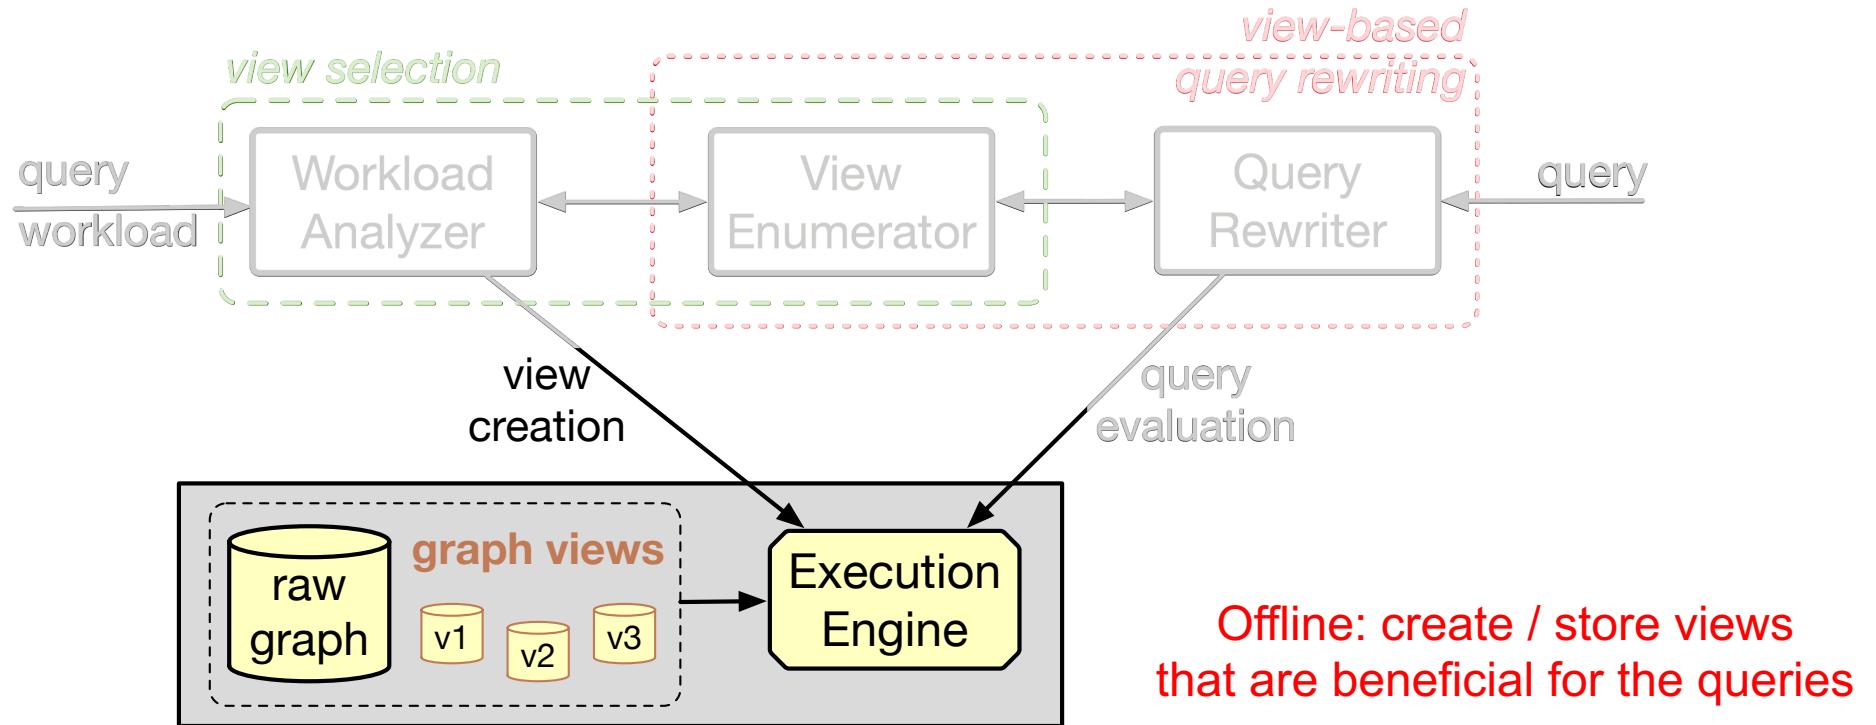
Size:  
 $O(K)$  vertices  
 $O(K)$  edges  
 $O(MB)$  on disk

# Kaskade: Graph Views for Efficient Graph Analytics

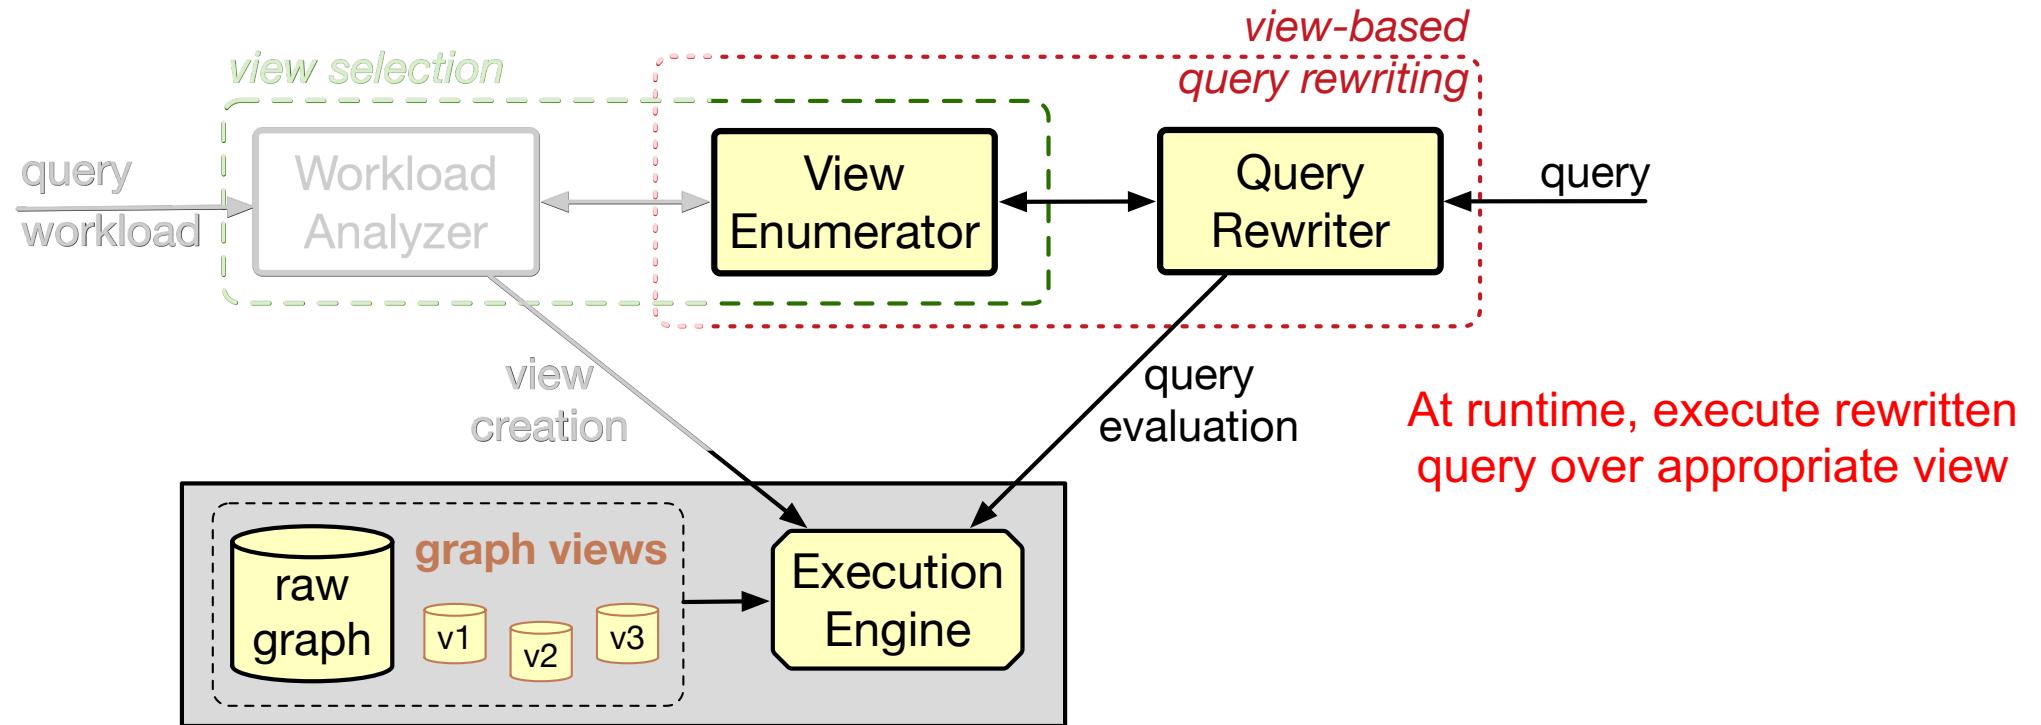
# Kaskade: a system for graph query optimization using graph views



# Kaskade: a system for graph query optimization using graph views



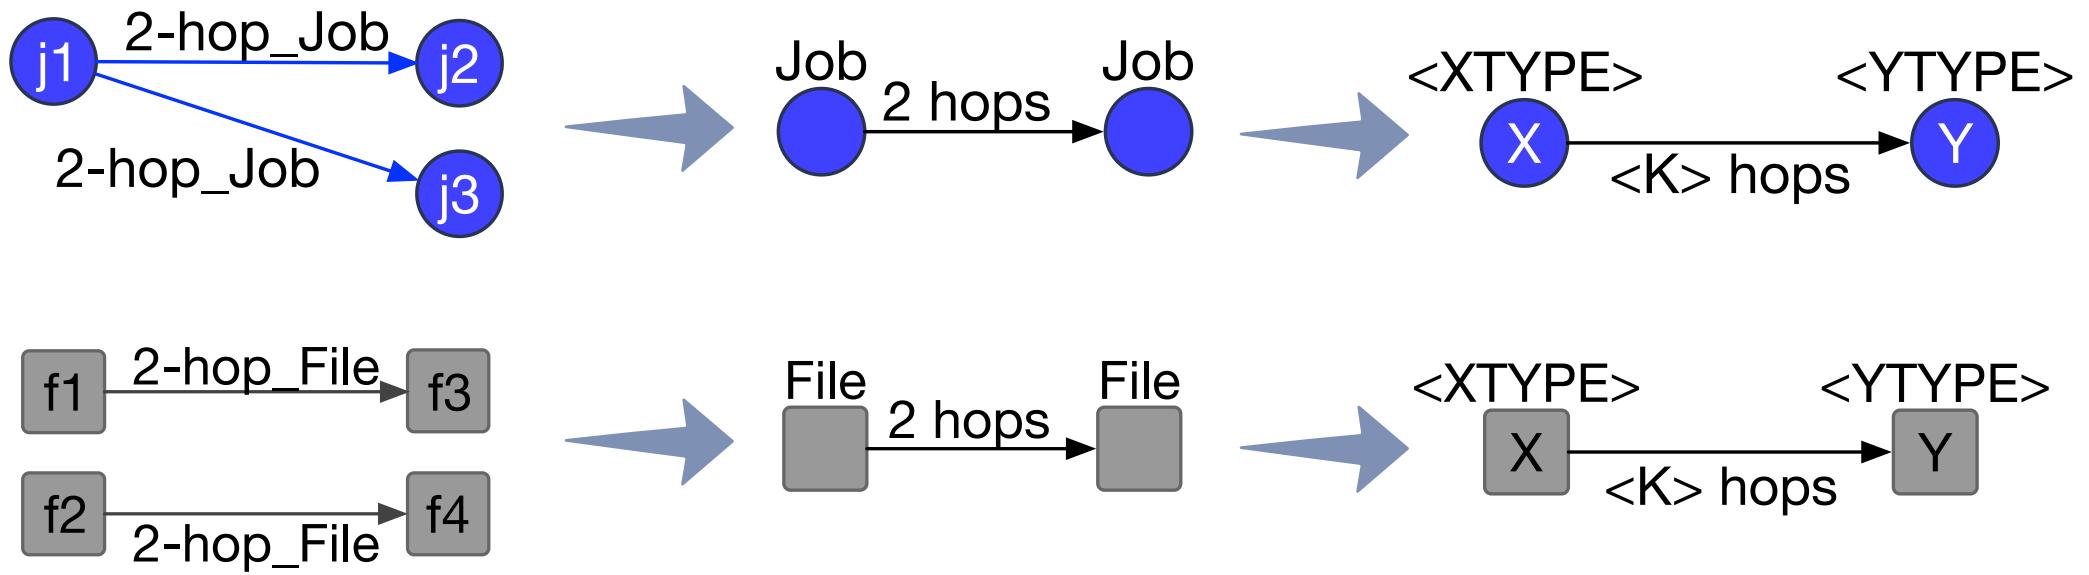
# Kaskade: a system for graph query optimization using graph views



Key contribution: exploit graph regularity to find efficient views for faster query execution

Graph view templates define the possible views we consider

# Graph view templates define the possible views we consider



*Concrete 2-hop connector views*

*are instantiations of*

*$k$ -hop connector view template.*

**View template:** `kHopConnector(X, Y, XTYPE, YTYPE, K)` with `X=Y`

Key contribution: exploit graph regularity to find efficient views for faster query execution

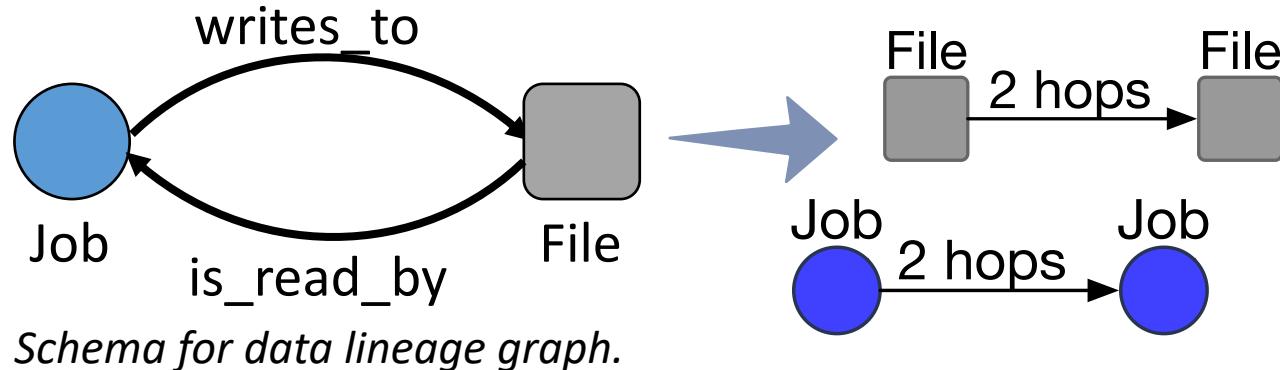
Graph view templates define the possible views we consider

Constraints are

- regular patterns imposed by the graph's schema (a graph)
- regular patterns imposed by the query (a graph)

# Constraints are regular patterns imposed by the graph's schema and query

Schema is a graph which imposes constraints on what views can be enumerated.



Query is a graph which imposes constraints on what views can be enumerated.

```
MATCH (j1:Job) -> (f1:File)  
      (f1:File) - [r*0..k] -> (f2:File)  
      (f2:File) -> (j2:Job)  
RETURN j1, j2
```

*Pattern matches all jobs  
up to k hops away*

Key contribution: exploit graph regularity to find efficient views for faster query execution

Graph view templates define the possible views we consider

Constraints are

- regular patterns imposed by the graph's schema (a graph)
- regular patterns imposed by the query (a graph)

Constraint-based view enumeration

- exploits **regularity** on graph's schema and data's structural properties.
- **finds views** that can correctly answer queries.
- rewrites queries in terms of these views.

# Problem formulation: finding graph views as a Boolean satisfiability problem

Graph view enumerators (*templates*) are encoded as logical clauses that are **satisfiable iff view is feasible** given input constraints.

e.g., a k-hop Job-Job connector view can be used iff:

- A k-hop path exists between Job vertices in the graph's schema
- A k-hop path exists between Job vertices in the query
- Only Job vertices are projected out in the query's schema

# Problem formulation: finding graph views as a Boolean satisfiability problem

- A k-hop path exists between Job vertices in the query
- A k-hop path exists between Job vertices in the graph's schema

```
% k-hop connector between nodes X and Y.  
kHopConnector(X, Y, XTYPE, YTYPE, K) :-  
    % query constraints  
    queryVertexType(X, XTYPE),  
    queryVertexType(Y, YTYPE),  
    queryKHopPath(X, Y, K),  
    % schema constraints  
    schemaKHopPath(XTYPE, YTYPE, K).
```

*Prolog definition for k-hop connector view template*

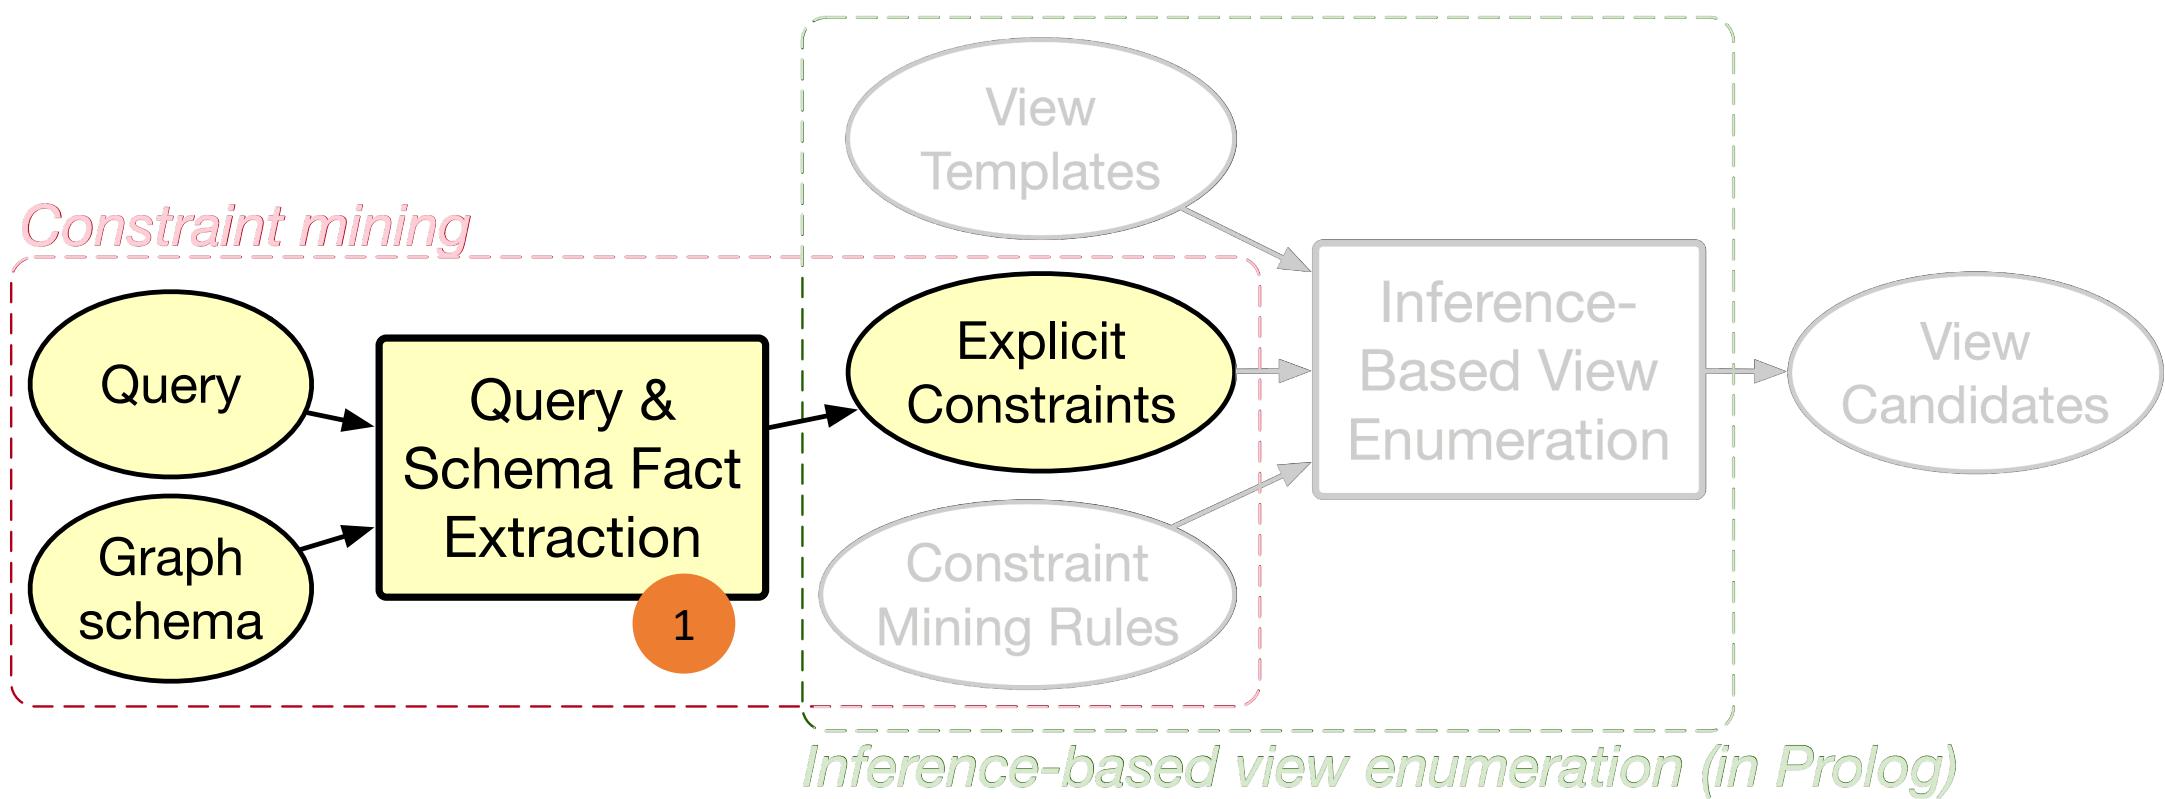
# Problem formulation: finding graph views as a Boolean satisfiability problem

- A k-hop path exists between Job vertices in the query
- A k-hop path exists between Job vertices in the graph's schema

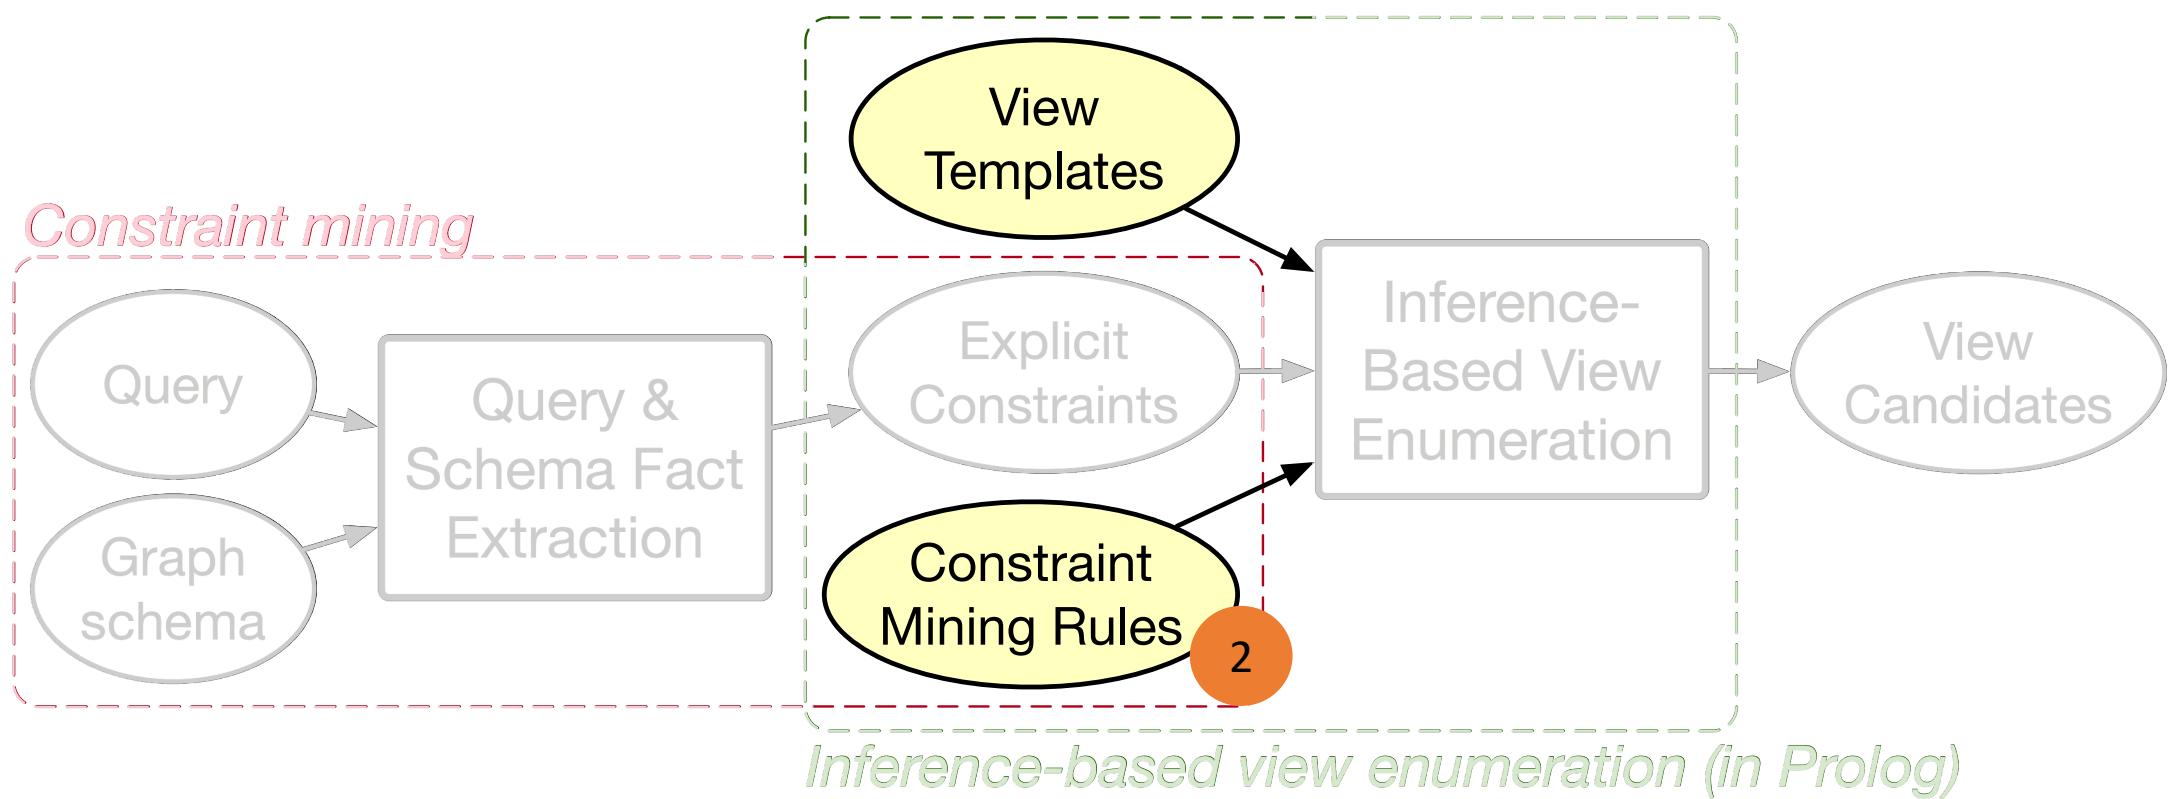
```
% k-hop connector between nodes X and Y.  
kHopConnector(X, Y, XTYPE, YTYPE, K) :-  
    % query constraints  
    queryVertexType(X, XTYPE),  
    queryVertexType(Y, YTYPE),  
    queryKHopPath(X, Y, K),  
    % schema constraints  
    schemaKHopPath(XTYPE, TYPE, K).
```

Prolog definition for k-hop connector view template

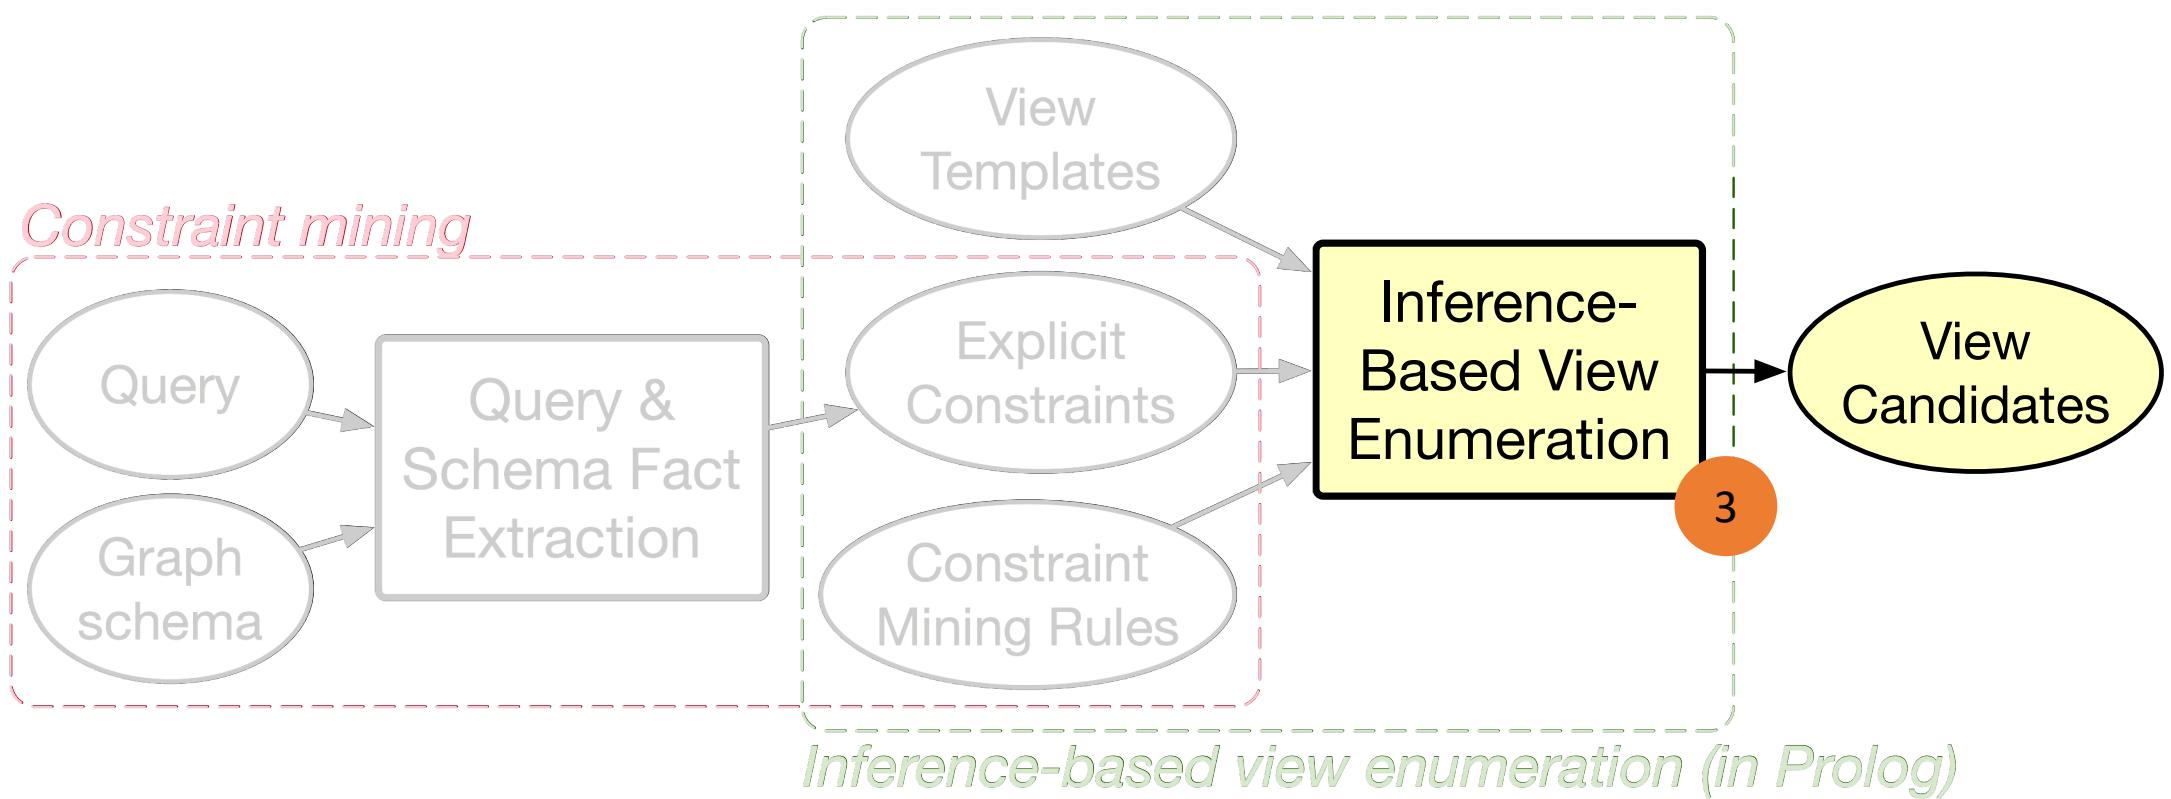
# Recap: Constraint-based view enumeration



# Recap: Constraint-based view enumeration



# Recap: Constraint-based view enumeration



# Constraint-based view enumeration: Why do we need a cost model?

Combinatorial problem: **constraints reduce search space** of possible views, but can still enumerate **more views than can fit in memory**.

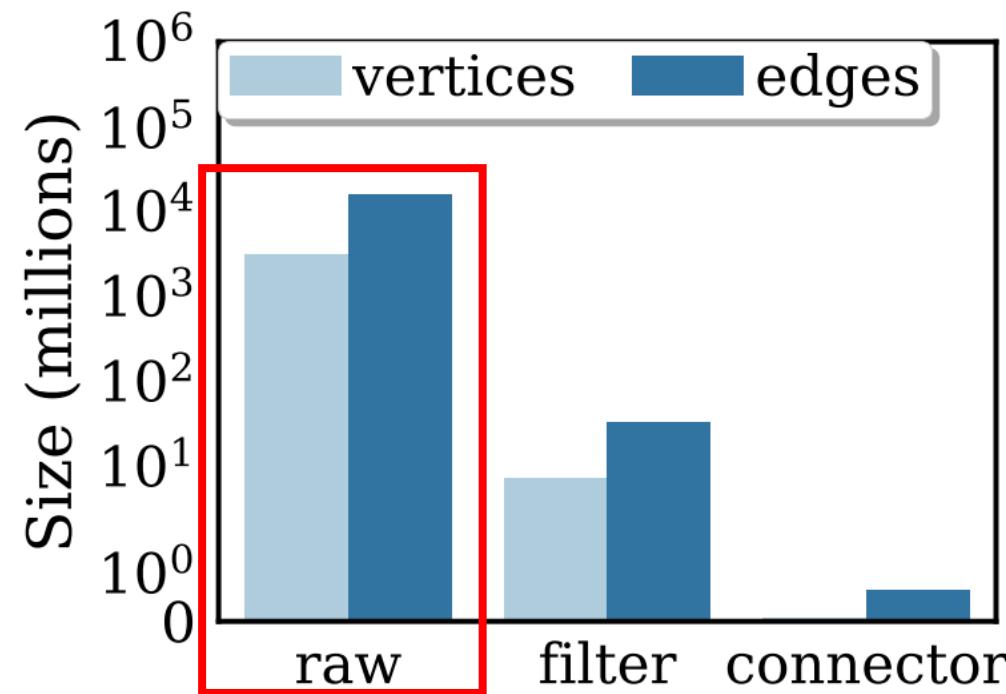
Cost model helps us **estimate how beneficial** it is to materialize a view.

Views are costed in terms of estimated size (I/O dominates).

**Optimizer picks best view** subject to resource budget (memory).

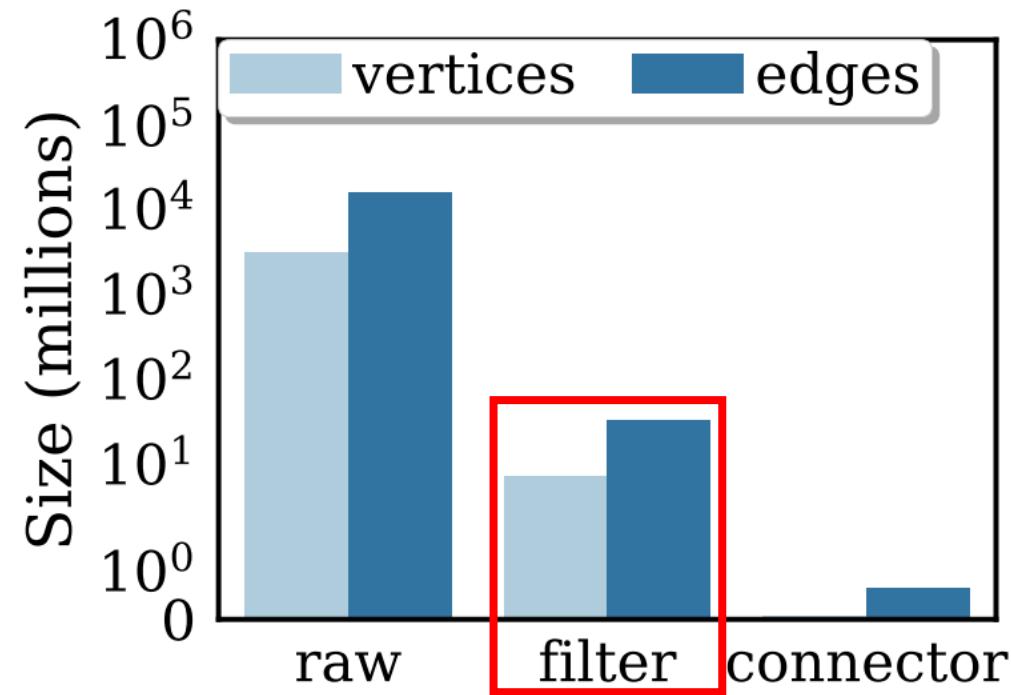
# Experimental Results

# Example gains in size reduction



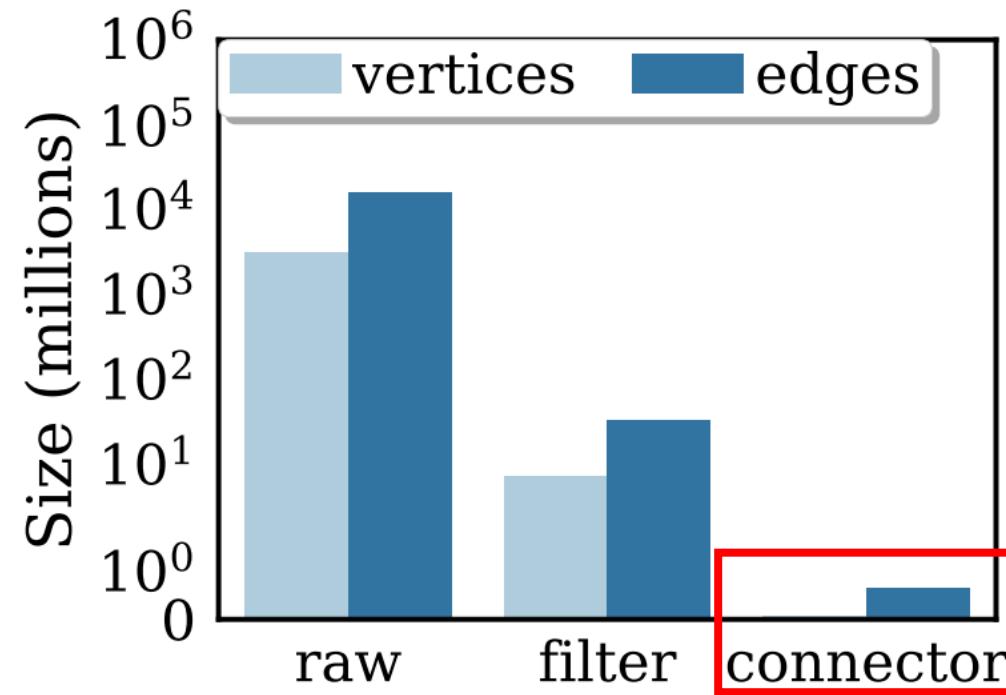
*Original data lineage raw graph has billions of vertices and edges.*

# Example gains in size reduction



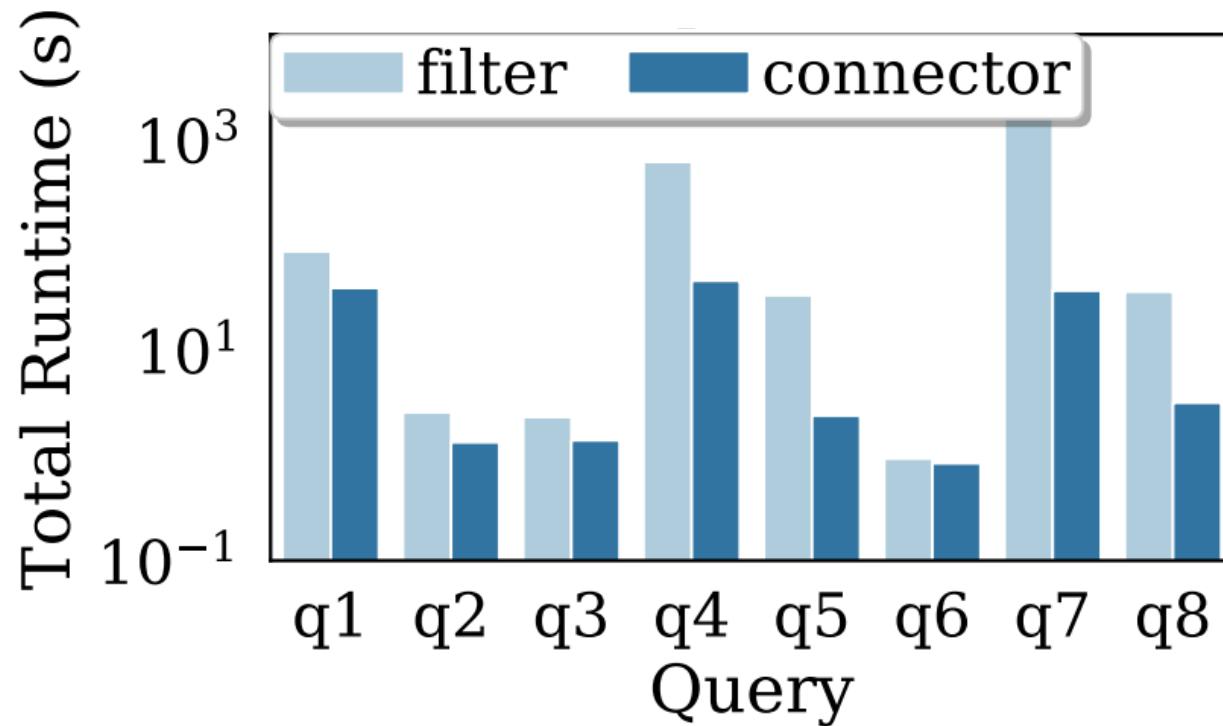
*Graph view that filters out unnecessary vertex/edge types has millions of edges.*

# Example gains in size reduction



*A 2-hop job-job connector graph view further reduces size.*

# Example gains in query runtime



*Runtime gains compound: orders of magnitude faster on smallest view.*

# Kaskade: Key Technical Contributions

## Graph view selection for query rewriting

Introduced class of graph views that are beneficial for different queries

Constraint-based view enumeration is guided by queries and structural properties of graph

## Significant performance gains

Up to 50x faster response times, orders of magnitude reduction in size

Rewriting techniques are engine-agnostic: only rely on fundamental graph transformations

# Kaskade: Misc. Contributions

Wrote graph views extraction pipeline, made available in prod

Pipeline served as basis for other papers from Microsoft folks, e.g.,

## **Peering through the Dark: An Owl's View of Inter-job Dependencies and Jobs' Impact in Shared Clusters**

Andrew Chung  
Carnegie Mellon University

Konstantinos Karanasos  
Microsoft

Carlo Curino  
Microsoft

Panagiotis Garefalakis  
Imperial College London

Subru Krishnan  
Microsoft

Gregory R. Ganger  
Carnegie Mellon University

SIGMOD 2019

## **Unearthing inter-job dependencies for better cluster scheduling**

Andrew Chung, Carnegie Mellon University; Subru Krishnan, Konstantinos Karanasos, and Carlo Curino, Microsoft; Gregory R. Ganger, Carnegie Mellon University

OSDI 2020

# Thesis Outline

Kaskade: **Graph Views** for Efficient Graph Analytics  
*J M F da Trindade, K Karanasos, C Curino, S Madden, J Shun*  
*ICDE 2020*

Kairos: Efficient **Temporal Graph Analytics** on a **Single Machine**  
*J M F da Trindade, J Shun, S Madden, N Tatbul*  
*Under submission to VLDB 2024*

# Kairos: Key Technical Contributions

Efficient **parallel implementation of algorithms over temporal graphs**, e.g.  
minimal paths (earliest arr., latest dep., fastest, shortest duration)  
centrality (betweenness using fastest or shortest duration)  
k-core, ranking, and connectivity

Efficient data layouts

**T-CSR: compressed sparse row** for temporal graphs

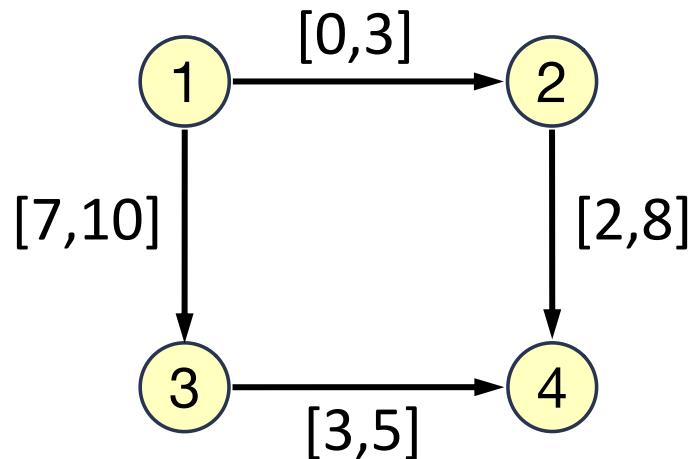
**TGER**: highly optimized **parallel index** for temporal edges

**Selective Indexing**

chooses best access method to retrieve vertex's nbors at runtime

# What is a temporal graph?

$$G = (V, E, \tau)$$



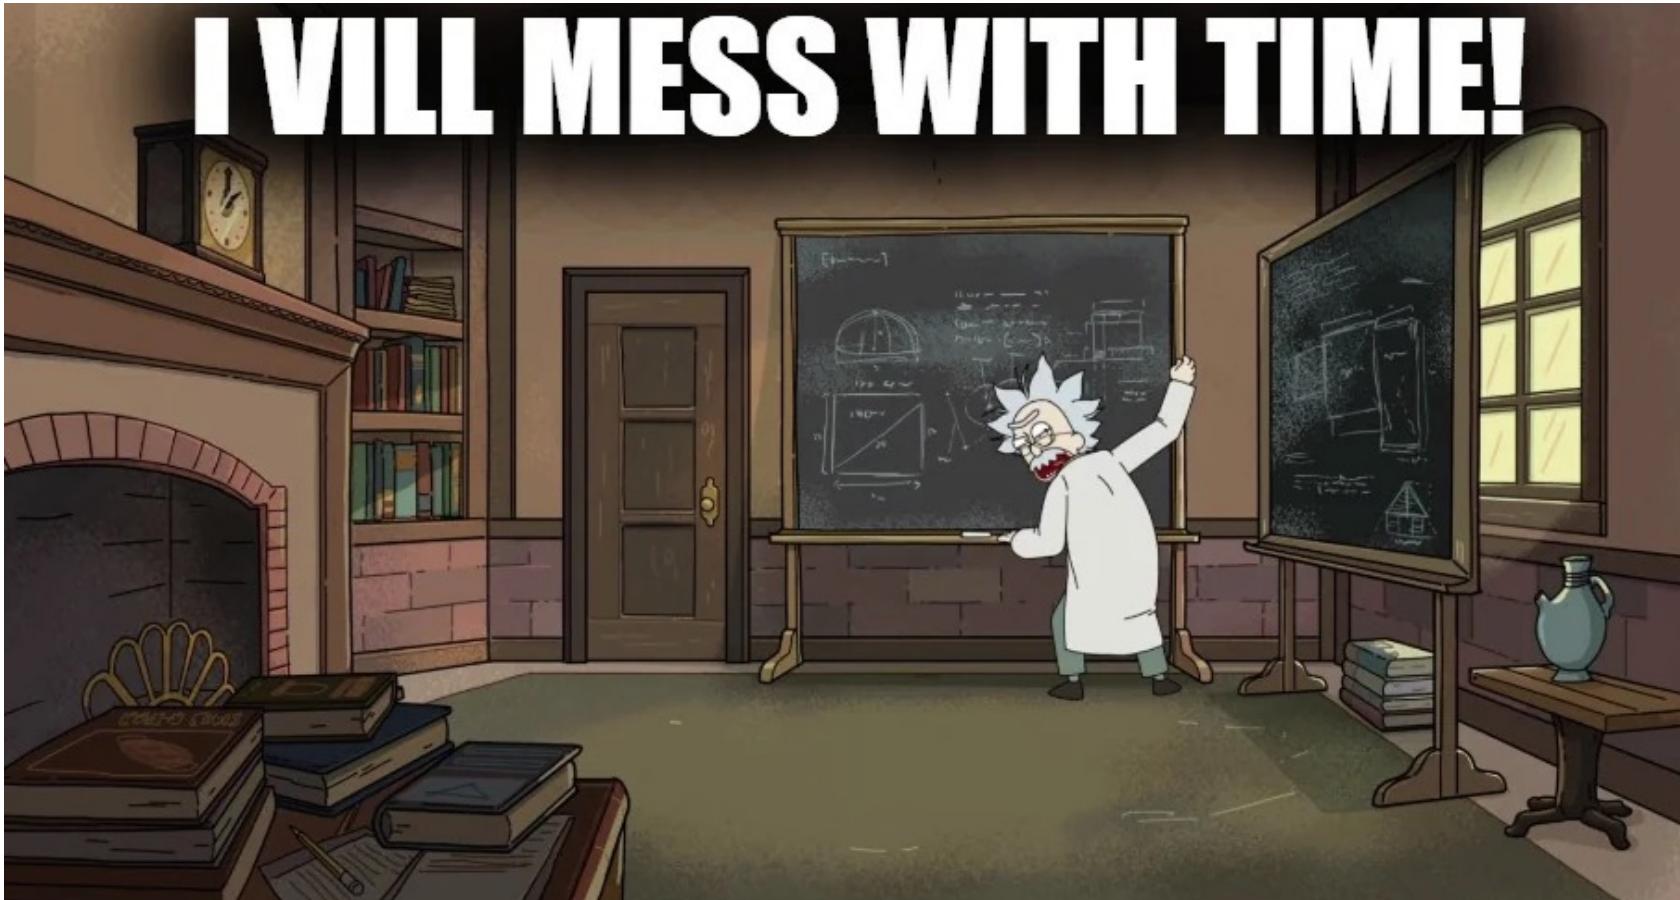
$$V = \{1, 2, 3, 4\}$$

$$E = \{(1,2), (1,3), (2, 4), (3, 4)\}$$

$$\tau: E \rightarrow \mathbb{N} \times \mathbb{N}$$

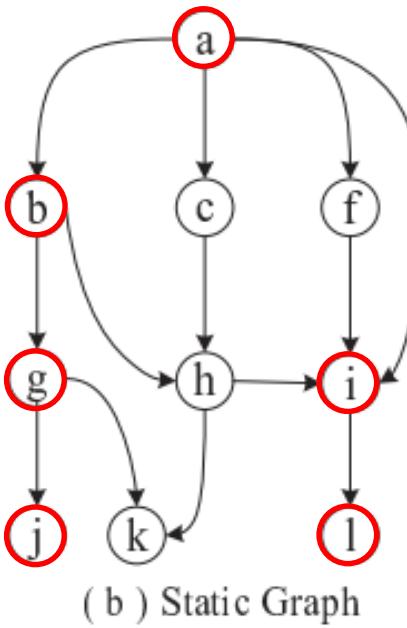
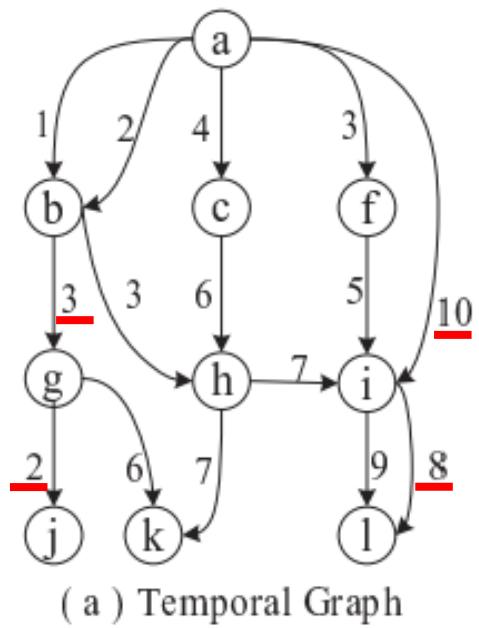
AKA as “temporal network”: edges only valid during [start time, end time] interval

# But really, why temporal graphs?



Rick and Morty – A Rickle In Time (Season 2, Episode 1)

Why do we care about temporal graphs? Certain questions answered incorrectly without time info



*Paths may be invalid or suboptimal when temporal info is ignored*

- $\langle a, b, g, j \rangle$  is not valid
- $\langle a, i, l \rangle$  is not shortest

# Why do we care about temporal graphs? Certain questions answered incorrectly without time info

## Temporal Networks

Petter Holme<sup>1,2,3</sup> and Jari Saramäki<sup>4</sup>

<sup>1</sup>*IceLab, Department of Physics, Umeå University, 901 87 Umeå, Sweden*

<sup>2</sup>*Department of Energy Science, Sungkyunkwan University, Suwon 440–746, Korea*

<sup>3</sup>*Department of Sociology, Stockholm University, 106 91 Stockholm, Sweden*

<sup>4</sup>*Department of Biomedical Engineering and Computational Science,  
School of Science, Aalto University, 00076 Aalto, Espoo, Finland*

communications  
physics

Use of temporal contact graphs to understand  
the evolution of COVID-19 through contact  
tracing data

Mincheng Wu<sup>1,2,8</sup>, Chao Li<sup>1,2,8</sup>, Zhangchong Shen<sup>2</sup>, Shibo He<sup>1,2,✉</sup>, Lingling Tang<sup>3</sup>, Jie Zheng<sup>4</sup>, Yi Fang<sup>5</sup>,  
Kehan Li<sup>6,2</sup>, Yanggang Cheng<sup>2</sup>, Zhiguo Shi<sup>6</sup>, Guoping Sheng<sup>3</sup>, Yu Liu<sup>5</sup>, Jinxing Zhu<sup>5</sup>, Xinjiang Ye<sup>5</sup>, Jinlai Chen<sup>5</sup>,  
Wenrong Chen<sup>5</sup>, Lanjuan Li<sup>7,✉</sup>, Youxian Sun<sup>1,2</sup> & Jiming Chen<sup>1,2,✉</sup>

---

## Temporal Graph Benchmark for Machine Learning on Temporal Graphs

---

Shenyang Huang<sup>1,2,\*</sup> Farimah Poursaei<sup>1,2,\*</sup> Jacob Danovitch<sup>1,2</sup> Matthias Fey<sup>3</sup>  
Weihua Hu<sup>3</sup> Emanuele Rossi<sup>4</sup> Jure Leskovec<sup>7</sup> Michael Bronstein<sup>8</sup>  
Guillaume Rabusseau<sup>1,5,6</sup> Reihaneh Rabbany<sup>1,2,6</sup>

<sup>1</sup>Mila - Quebec AI Institute, <sup>2</sup>School of Computer Science, McGill University  
<sup>3</sup>Kumo.AI, <sup>4</sup>Imperial College London, <sup>5</sup>DIRO, Université de Montréal  
<sup>6</sup>CIFAR AI Chair, <sup>7</sup>Stanford University, <sup>8</sup>University of Oxford

## You Are What You Do: Hunting Stealthy Malware via Data Provenance Analysis

Qi Wang<sup>1</sup>, Wajih Ul Hassan<sup>1</sup>, Ding Li<sup>2</sup>, Kangkook Jee<sup>3</sup>, Xiao Yu<sup>2</sup>  
Kexuan Zou<sup>1</sup>, Junghwan Rhee<sup>2</sup>, Zhengzhang Chen<sup>2</sup>, Wei Cheng<sup>2</sup>, Carl A. Gunter<sup>1</sup>, Haifeng Chen<sup>2</sup>  
<sup>1</sup>University of Illinois Urbana-Champaign    <sup>2</sup>NEC Laboratories America, Inc.    <sup>3</sup>University of Texas at Dallas

{qiwang11, whassan3, kzou3, cgunter}@illinois.edu    kangkook.jee@utdallas.edu  
{dingli, xiao, rhee, zchen, weicheng, haifeng}@nec-labs.com

# Kairos: Key Technical Contributions

Efficient **parallel implementation of algorithms over temporal graphs**, e.g.  
minimal paths (earliest arr., latest dep., fastest, shortest duration)  
centrality (betweenness using fastest or shortest duration)  
k-core, ranking, and connectivity

# Categories of temporal graph analytics tasks

## Temporal Minimal Paths

- Provenance: tracking information flow in machine generated graphs
- Indoor routing: temporal shortest paths for robot navigation
- Epidemiology: identifying infection transmission paths for contact tracing

## Temporal Connectivity

- Social networks: analyzing community evolution

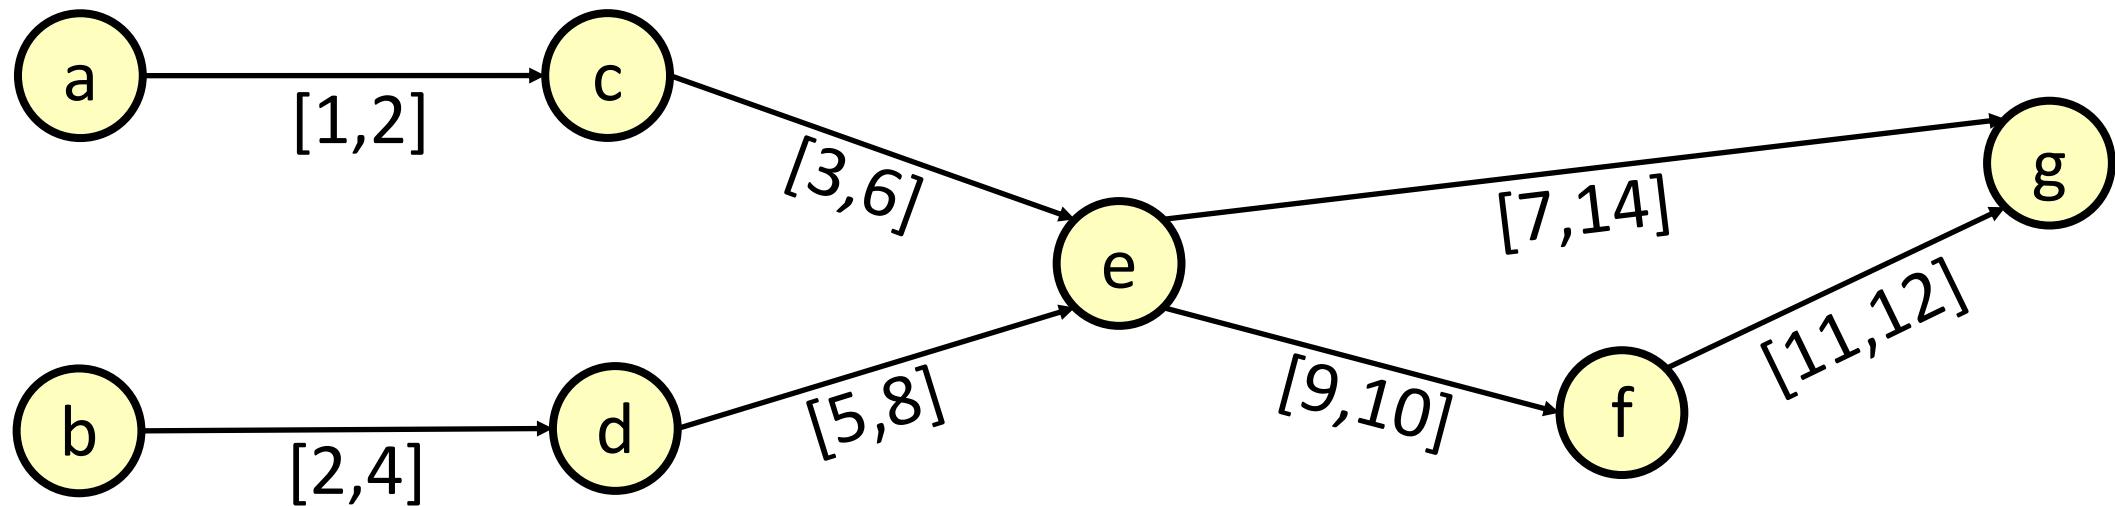
## Temporal Centrality

- Epidemiology: identifying critical nodes in the spread of infections

## Time-constrained reachability

- Social networks: analyzing influence propagation

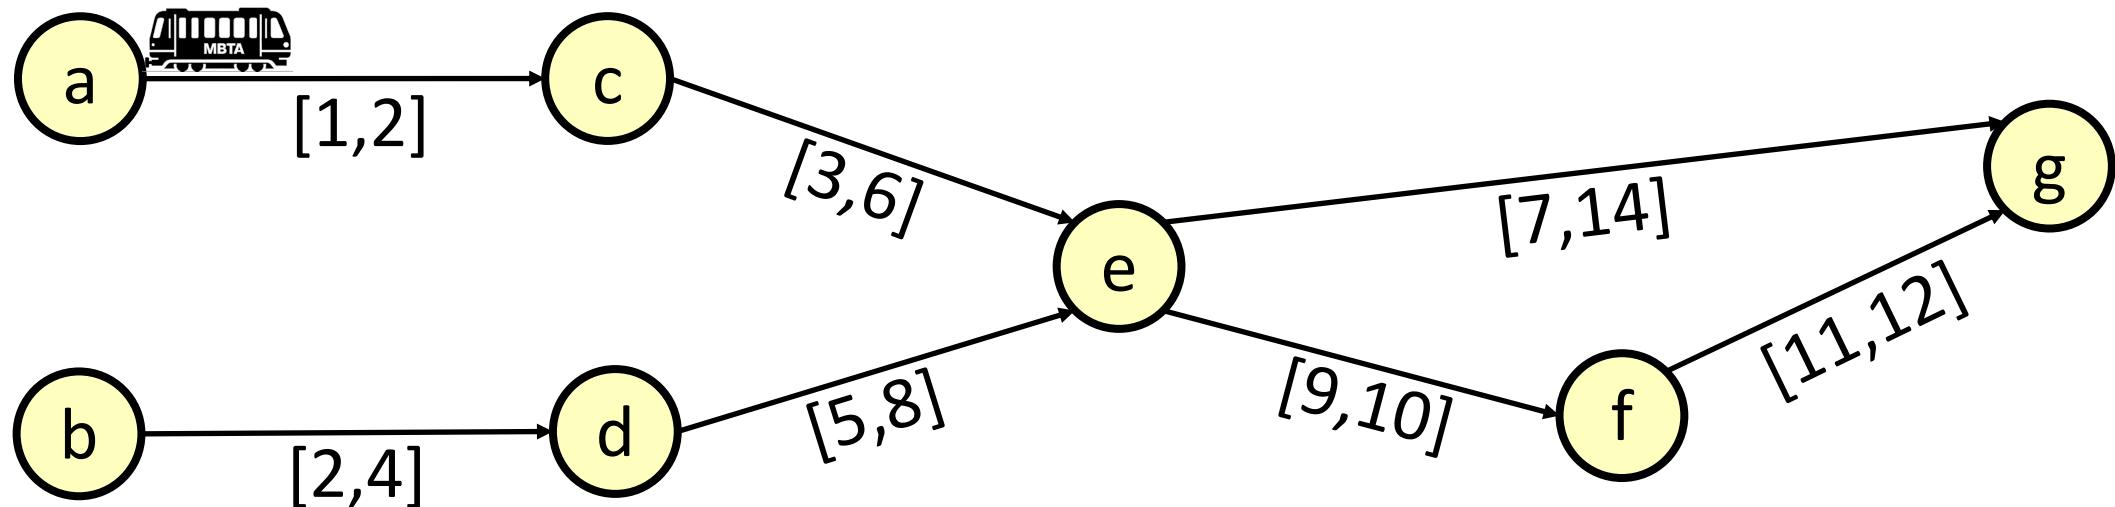
# Example T. Minimal Path: Earliest Arrival



Input: **graph G**, source vertex  $u$  in  $V$ , time interval  $[t_s, t_e]$

Output: earliest arrival from  $u$  to every vertex  $v$  in  $V$  within  $[t_s, t_e]$

# Example T. Minimal Path: Earliest Arrival

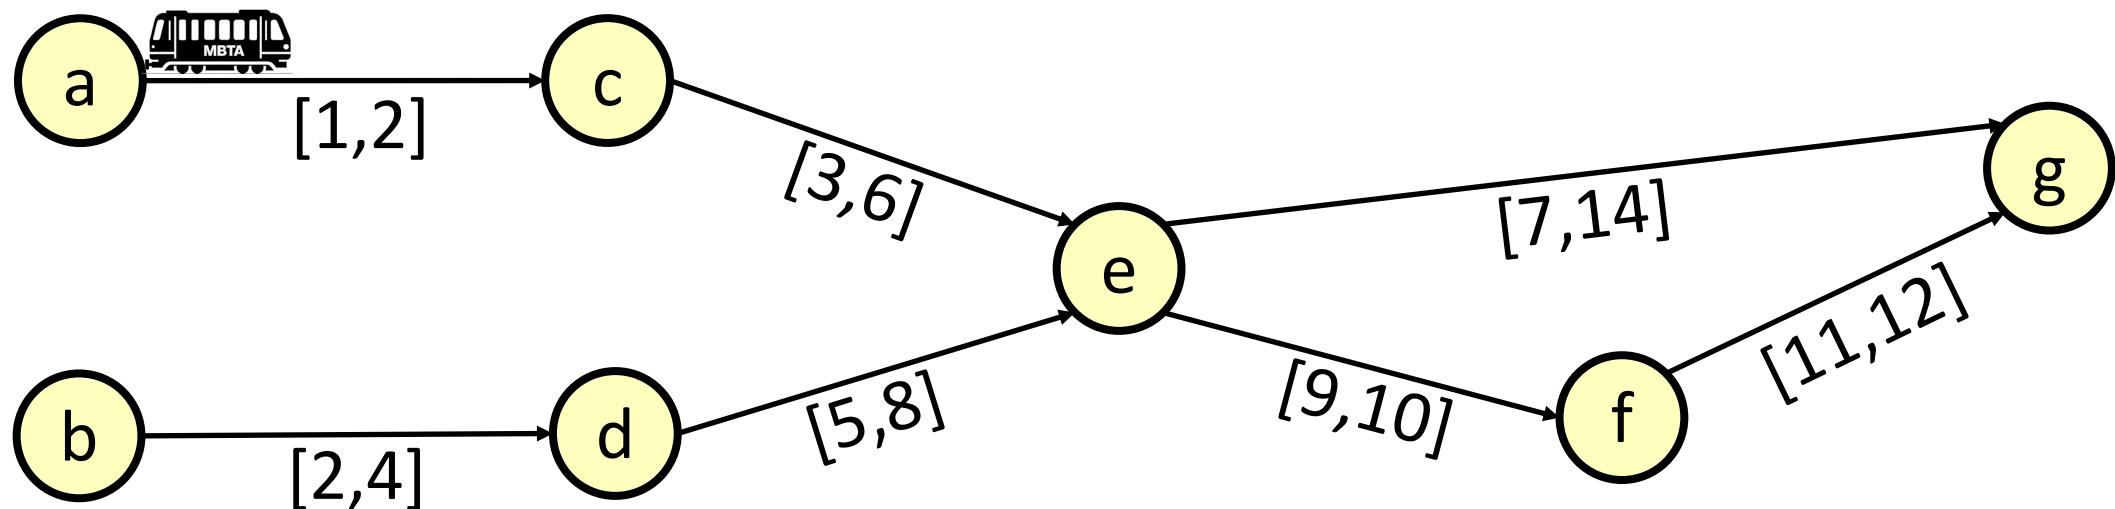


Input: graph  $G$ , **source vertex  $u$  in  $V$** , time interval  $[t_s, t_e]$

Output: earliest arrival from  $u$  to every vertex  $v$  in  $V$  within  $[t_s, t_e]$

**$u = a$**

# Example T. Minimal Path: Earliest Arrival



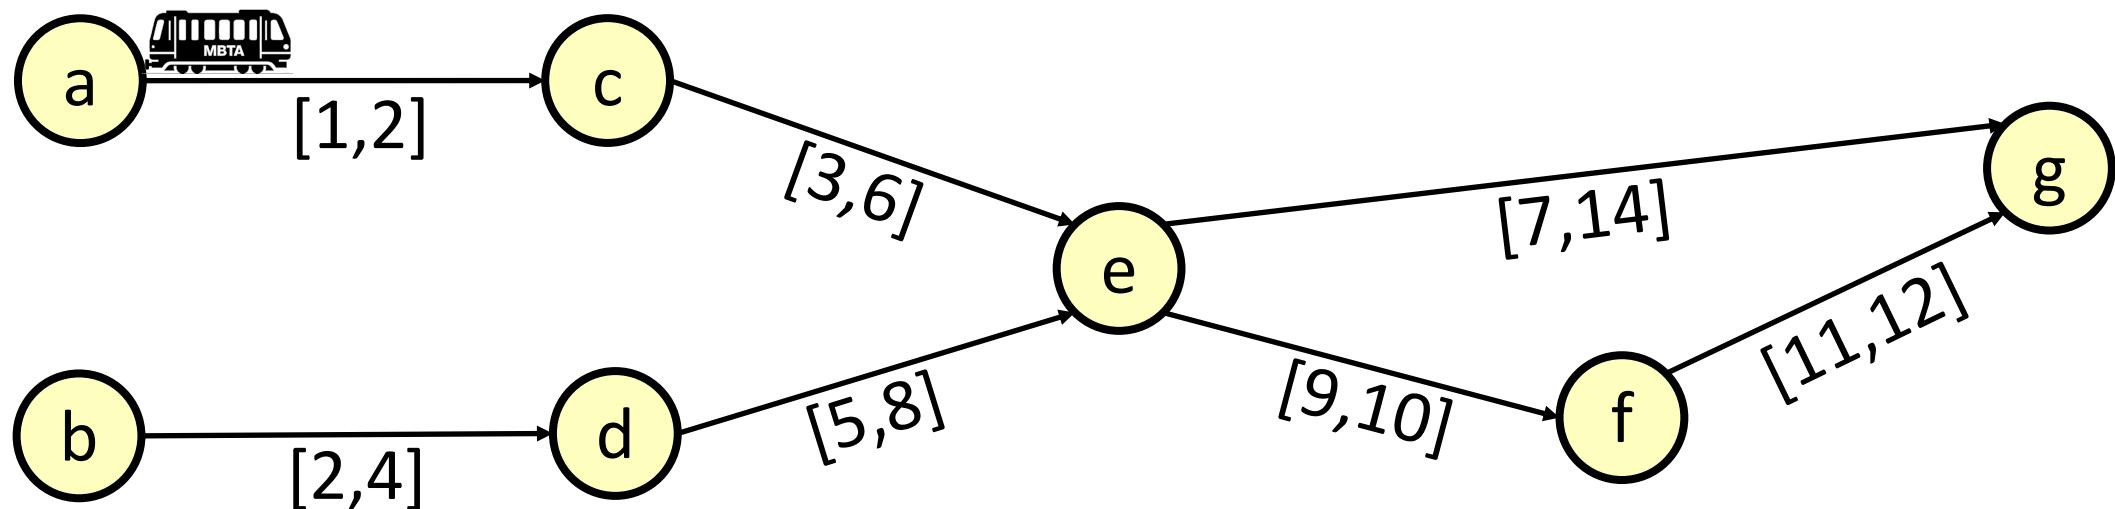
Input: graph  $G$ , source vertex  $u$  in  $V$ , **time interval  $[t_s, t_e]$**

Output: earliest arrival from  $u$  to every vertex  $v$  in  $V$  within  $[t_s, t_e]$

$u = a$

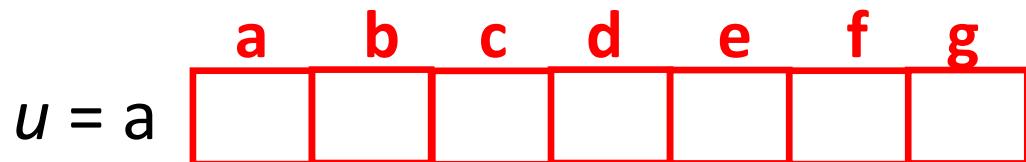
$$[t_s, t_e] = [1, 15]$$

# Example T. Minimal Path: Earliest Arrival



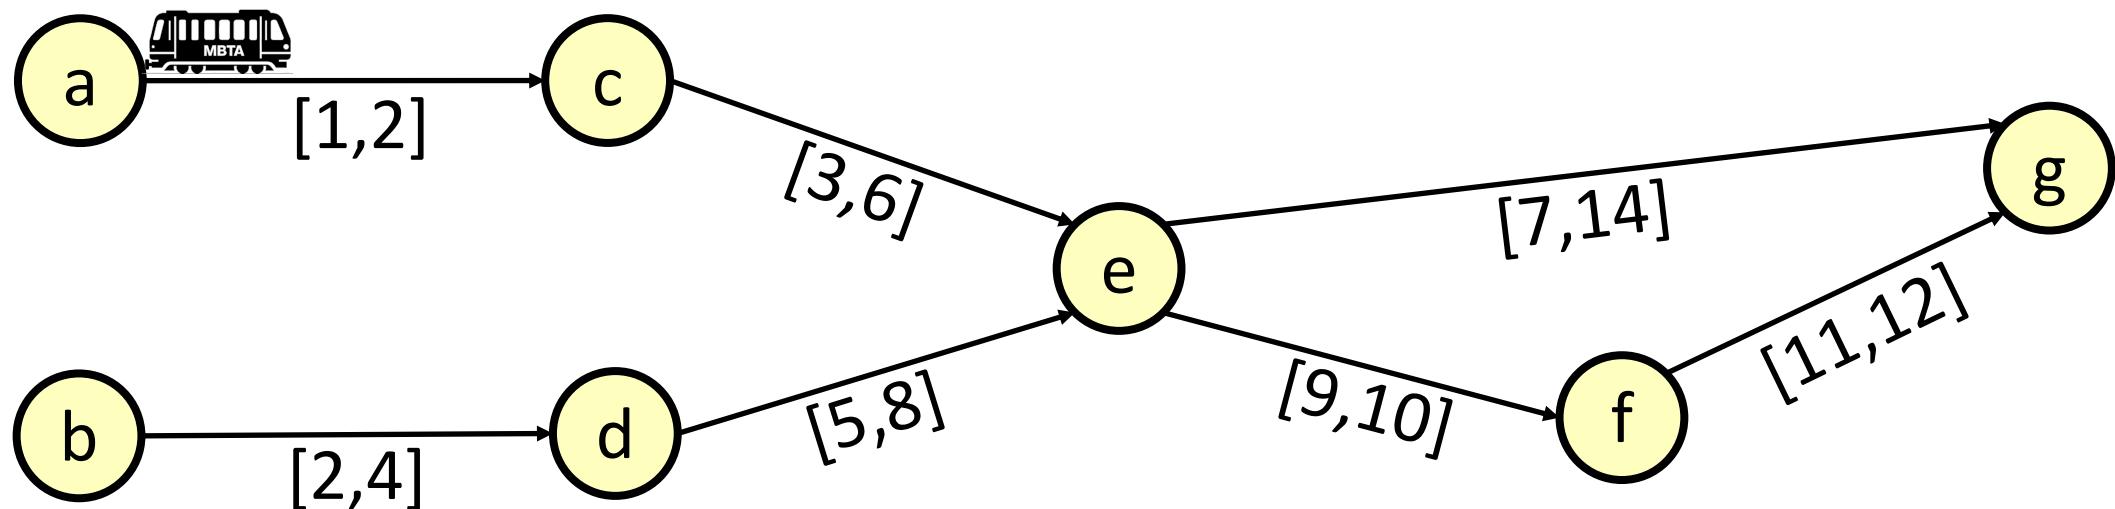
Input: graph  $G$ , source vertex  $u$  in  $V$ , time interval  $[t_s, t_e]$

Output: earliest arrival from  $u$  to every vertex  $v$  in  $V$  within  $[t_s, t_e]$



$$[t_s, t_e] = [1, 15]$$

# Example T. Minimal Path: Earliest Arrival



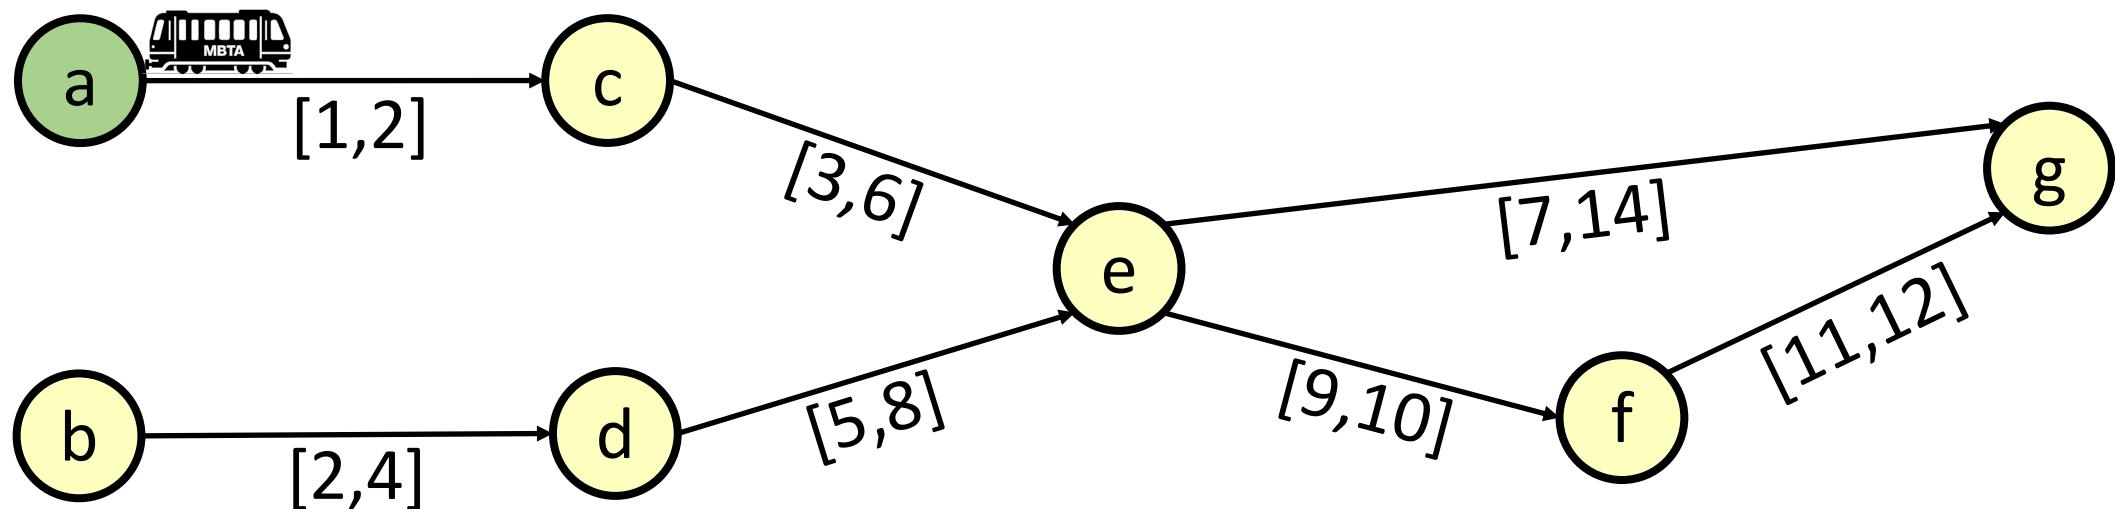
Input: graph  $G$ , source vertex  $u$  in  $V$ , time interval  $[t_s, t_e]$

Output: earliest arrival from  $u$  to every vertex  $v$  in  $V$  within  $[t_s, t_e]$

	a	b	c	d	e	f	g
$u = a$	$t_s$						

$$[t_s, t_e] = [1, 15]$$

# Example T. Minimal Path: Earliest Arrival



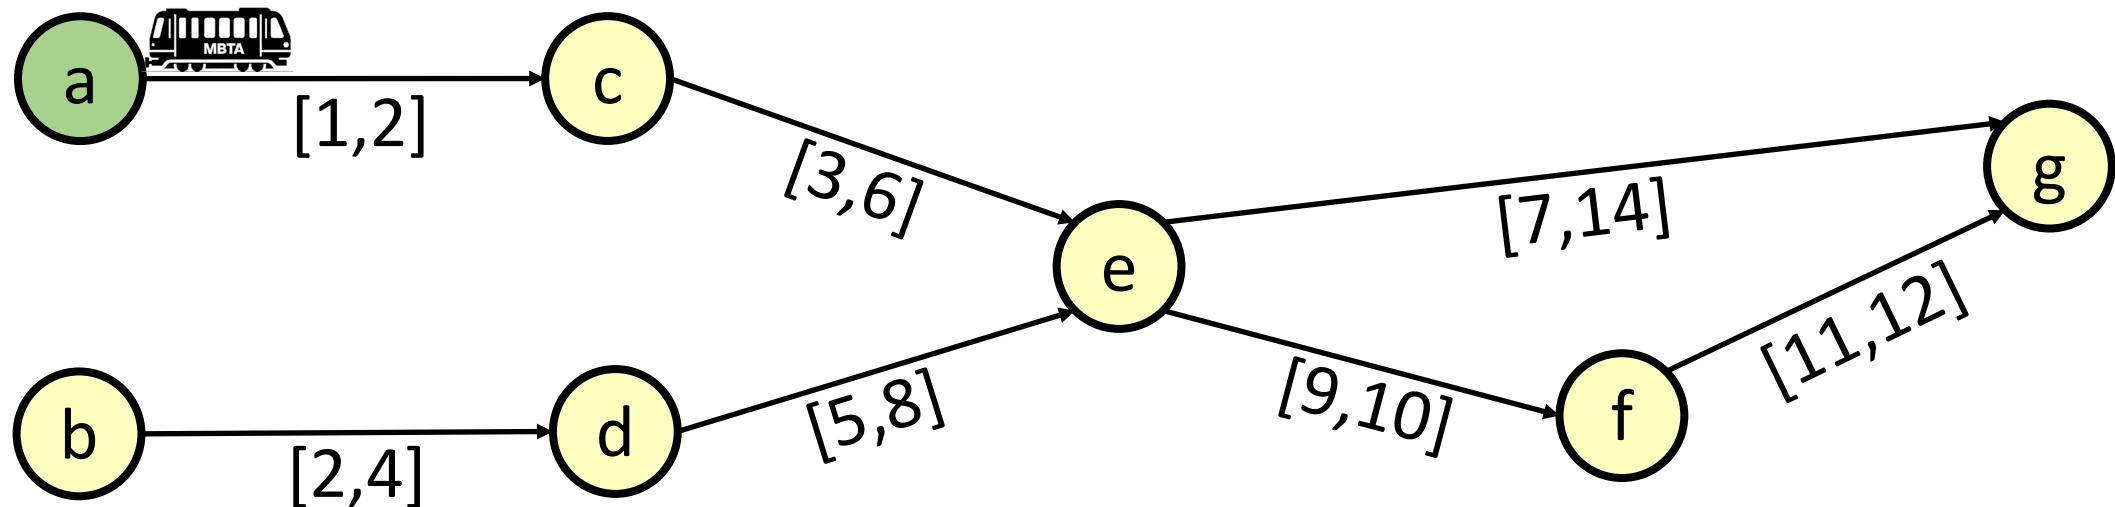
Input: graph  $G$ , source vertex  $u$  in  $V$ , time interval  $[t_s, t_e]$

Output: earliest arrival from  $u$  to every vertex  $v$  in  $V$  within  $[t_s, t_e]$

<b>a</b>	b	c	d	e	f	g
<b>u = a</b>	1					

$$[t_s, t_e] = [1, 15]$$

# Example T. Minimal Path: Earliest Arrival



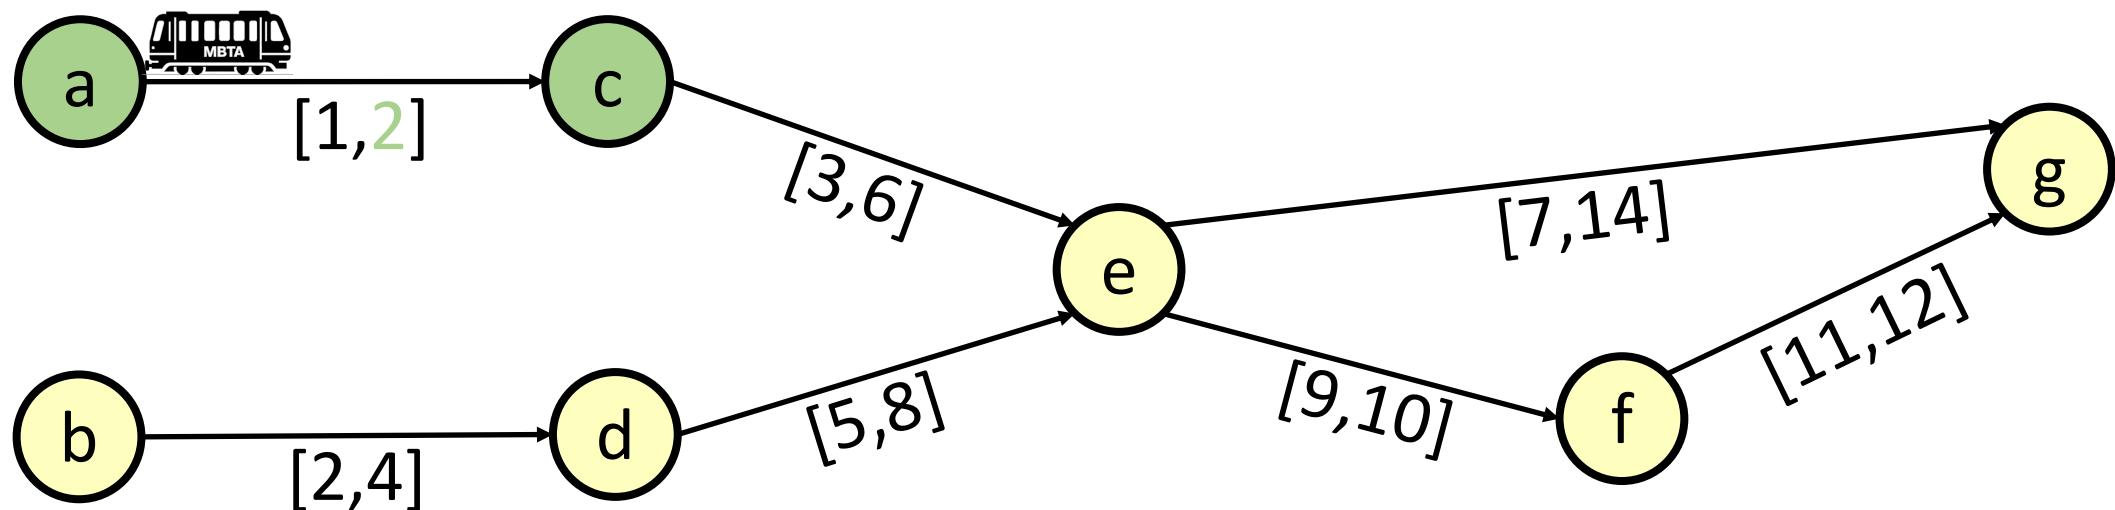
Input: graph  $G$ , source vertex  $u$  in  $V$ , time interval  $[t_s, t_e]$

Output: earliest arrival from  $u$  to every vertex  $v$  in  $V$  within  $[t_s, t_e]$

$u = a$	a	b	c	d	e	f	g
	1	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$

$$[t_s, t_e] = [1, 15]$$

# Example T. Minimal Path: Earliest Arrival

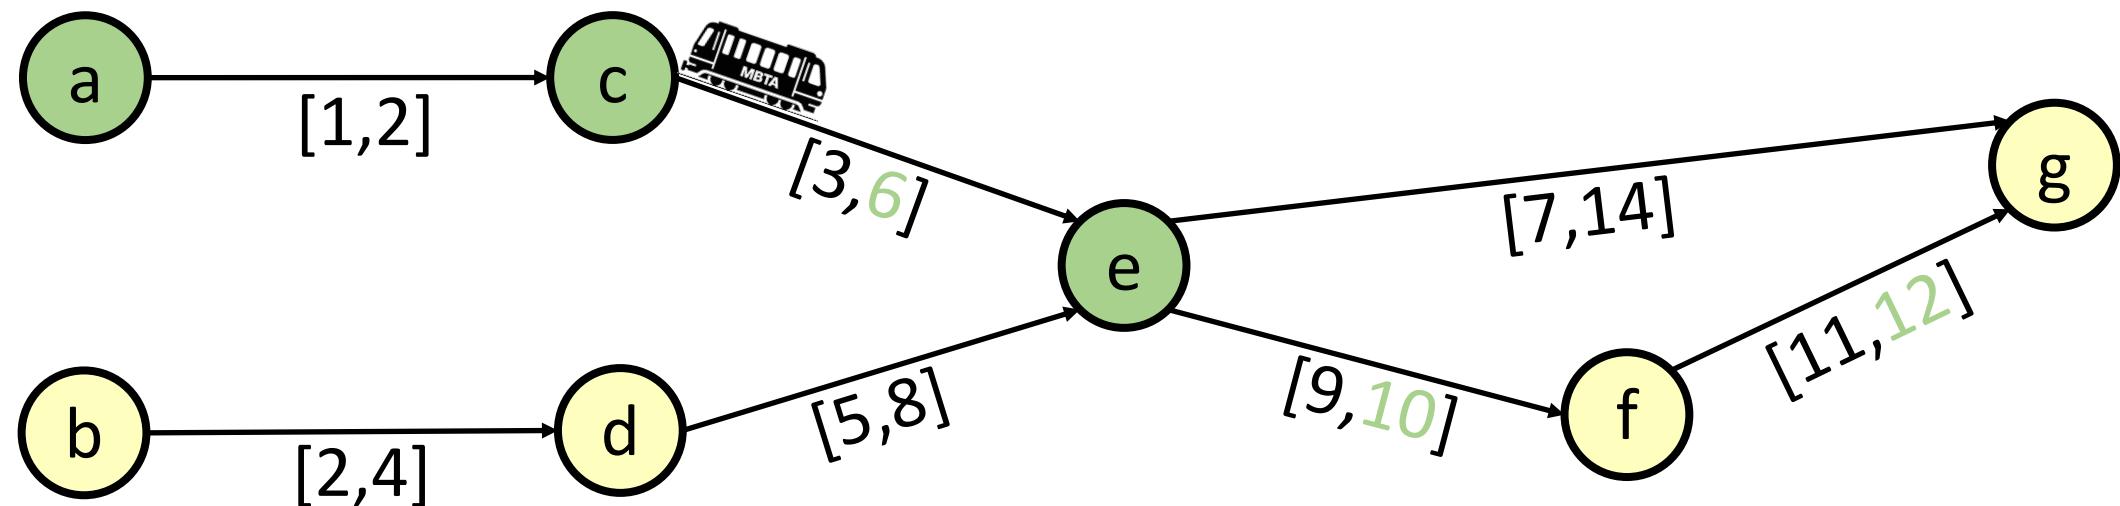


$u = a$

a	b	c	d	e	f	g
1	$\infty$	2				

$$[t_s, t_e] = [1, 15]$$

# Example T. Minimal Path: Earliest Arrival

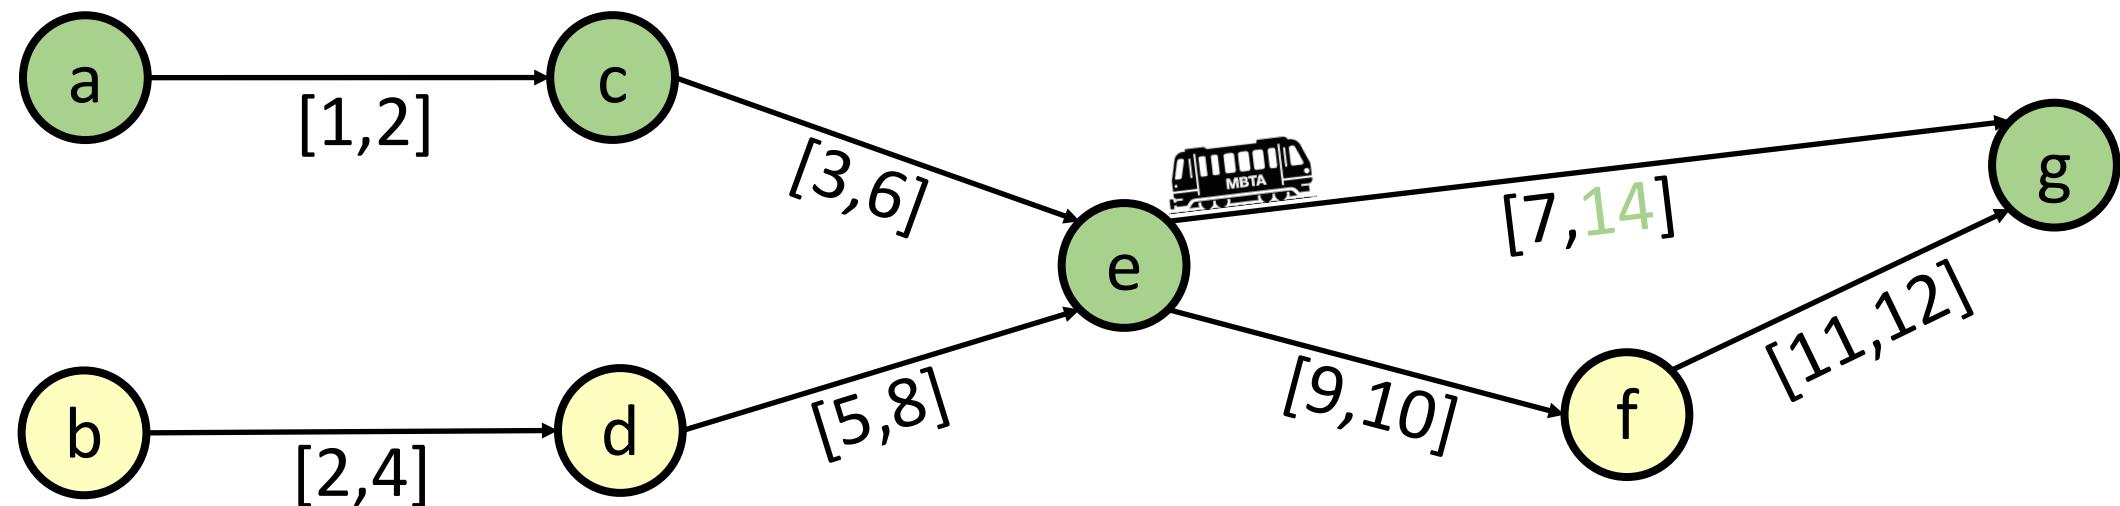


$u = a$

a	b	c	d	e	f	g
1	$\infty$	2	$\infty$	6		

$$[t_s, t_e] = [1, 15]$$

# Example T. Minimal Path: Earliest Arrival

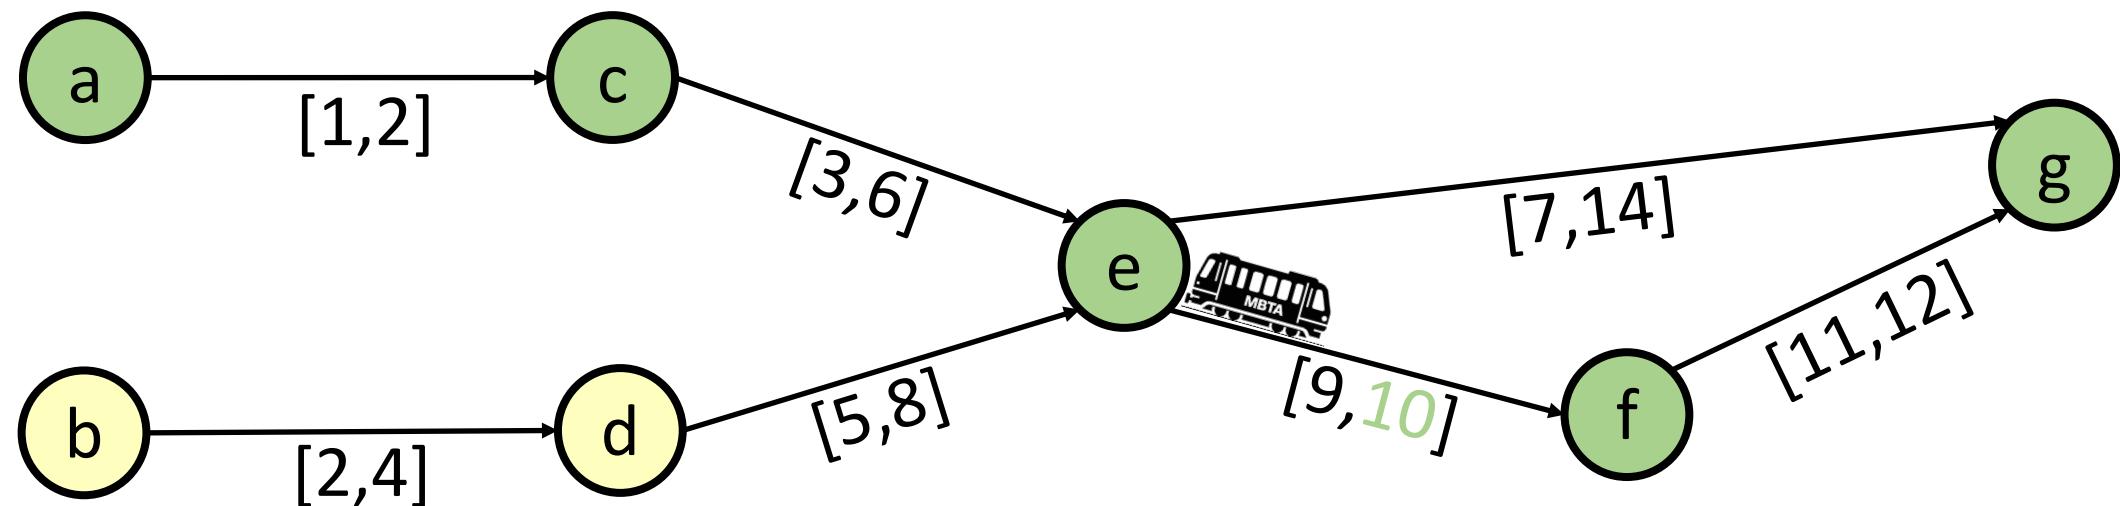


$u = a$

a	b	c	d	e	f	g
1	$\infty$	2	$\infty$	6	$\infty$	14

$$[t_s, t_e] = [1, 15]$$

# Example T. Minimal Path: Earliest Arrival

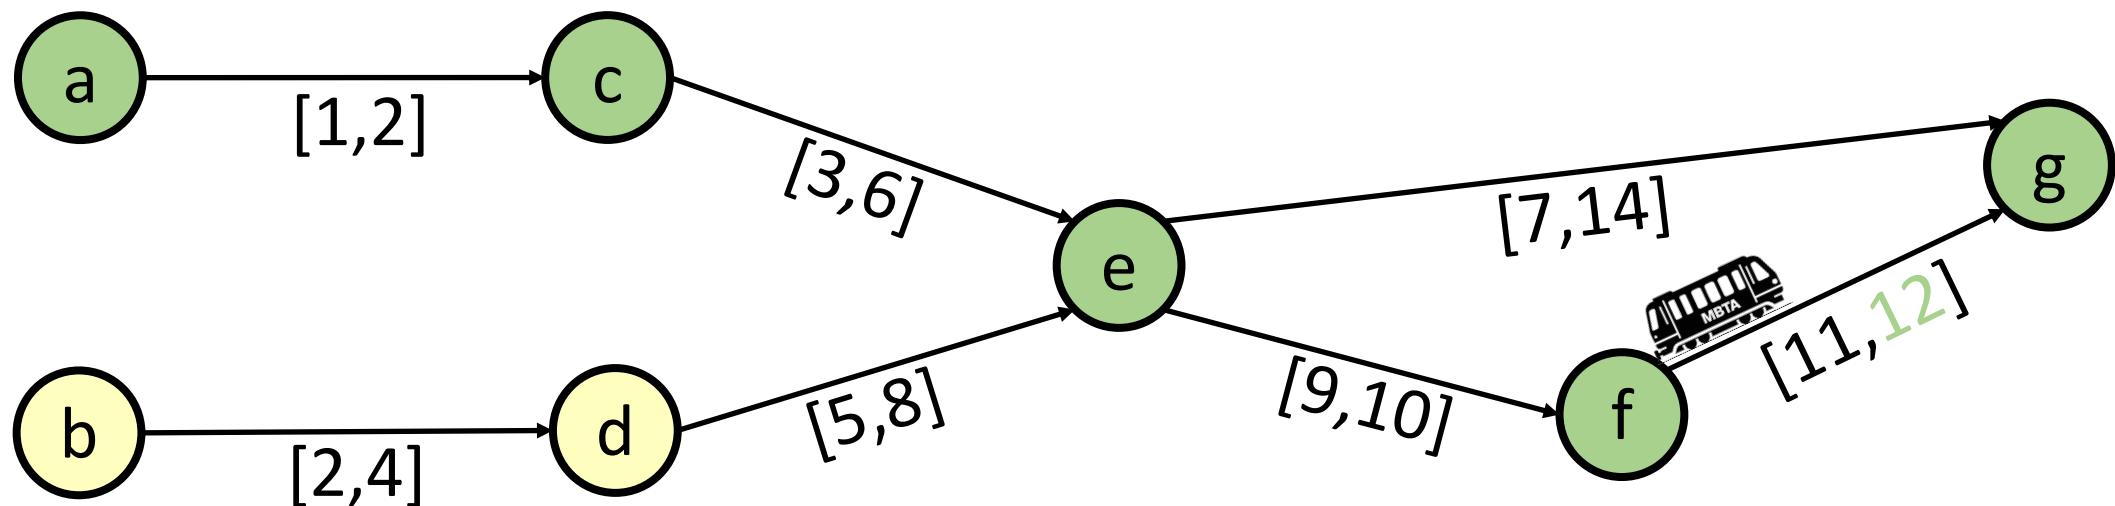


$u = a$

a	b	c	d	e	f	g
1	$\infty$	2	$\infty$	6	10	14

$$[t_s, t_e] = [1, 15]$$

# Example T. Minimal Path: Earliest Arrival

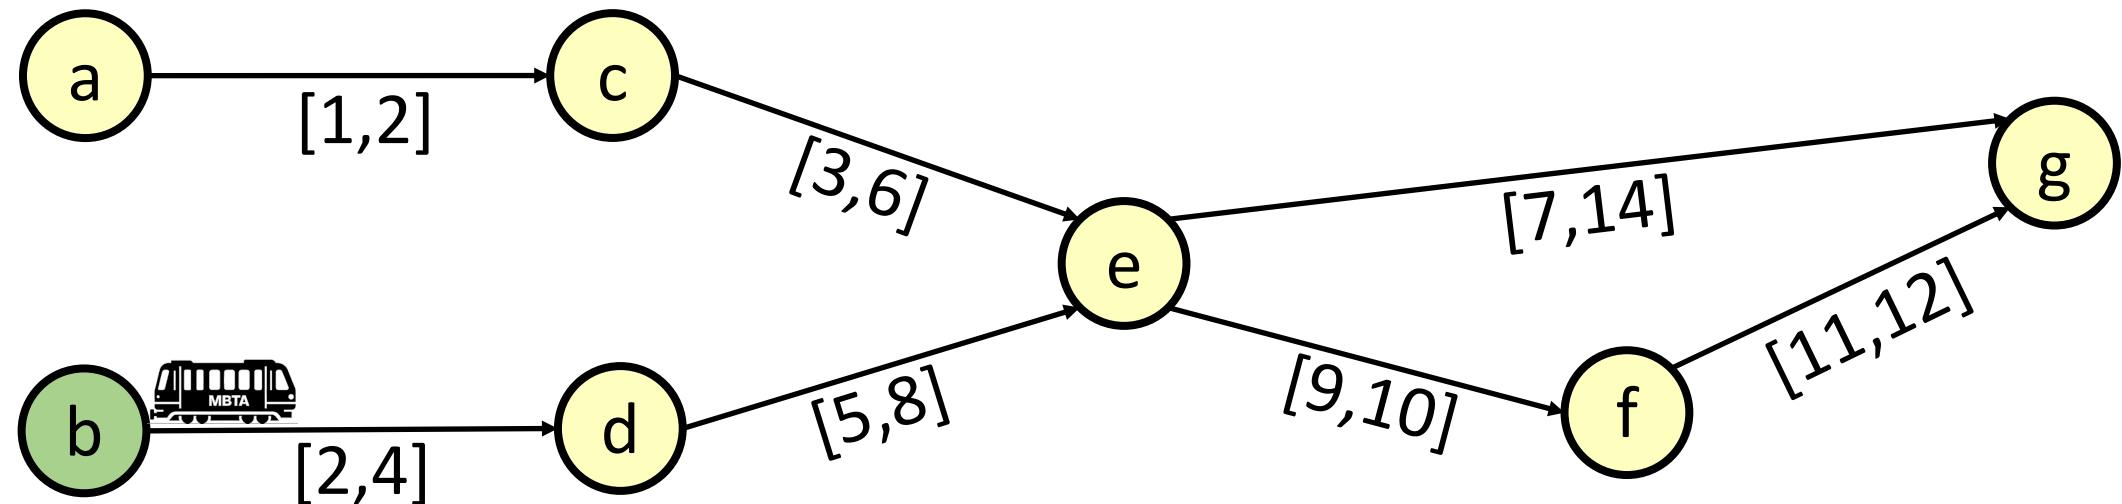


$u = a$

a	b	c	d	e	f	g
1	$\infty$	2	$\infty$	6	10	12

$$[t_s, t_e] = [1, 15]$$

# Example T. Minimal Path: Earliest Arrival



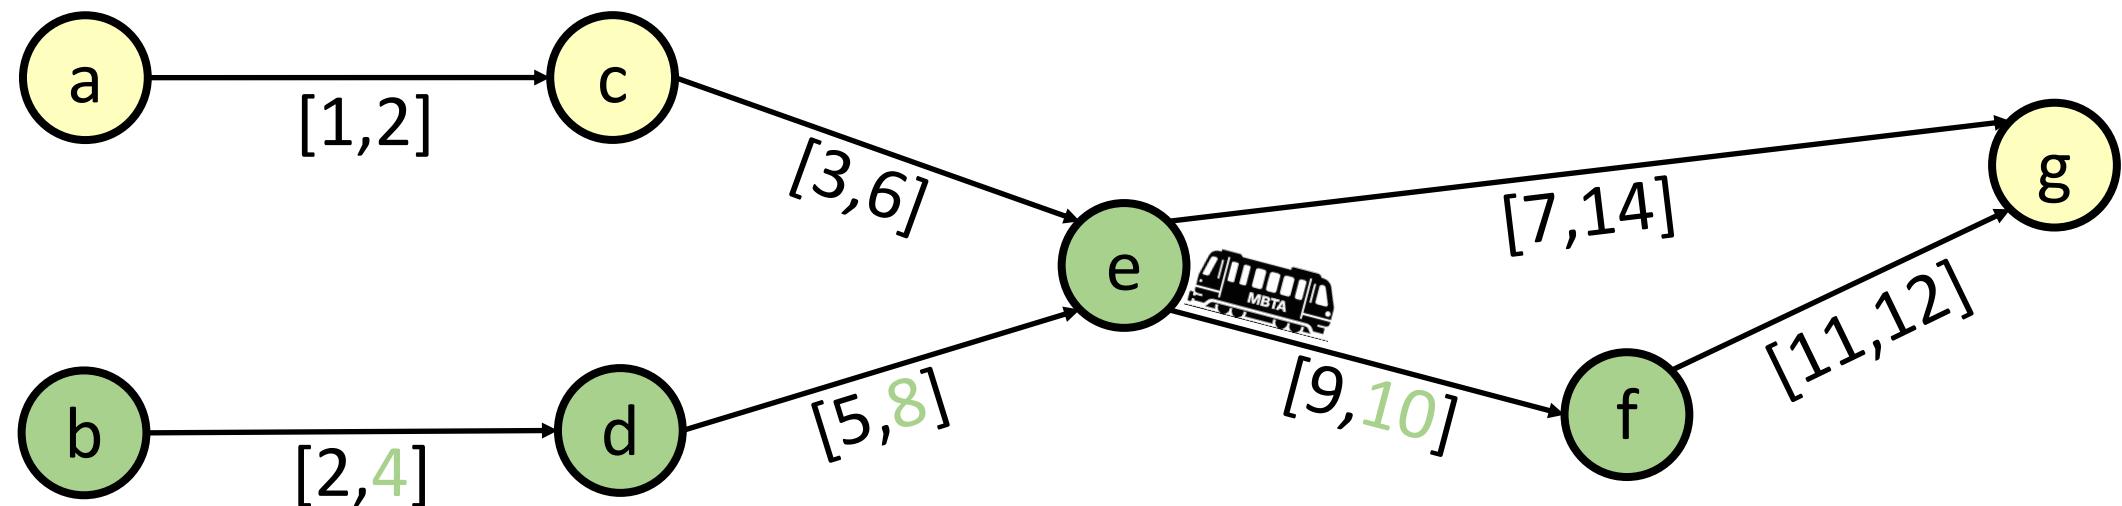
	a	b	c	d	e	f	g
$u = a$	1	$\infty$	2	$\infty$	6	10	12

$$[t_s, t_e] = [1, 15]$$

	a	b	c	d	e	f	g
$u = b$	$\infty$	2	$\infty$	$\infty$	$\infty$	$\infty$	$\infty$

$$[\mathbf{t}_s, \mathbf{t}_e] = [\mathbf{2}, 11]$$

# Example T. Minimal Path: Earliest Arrival



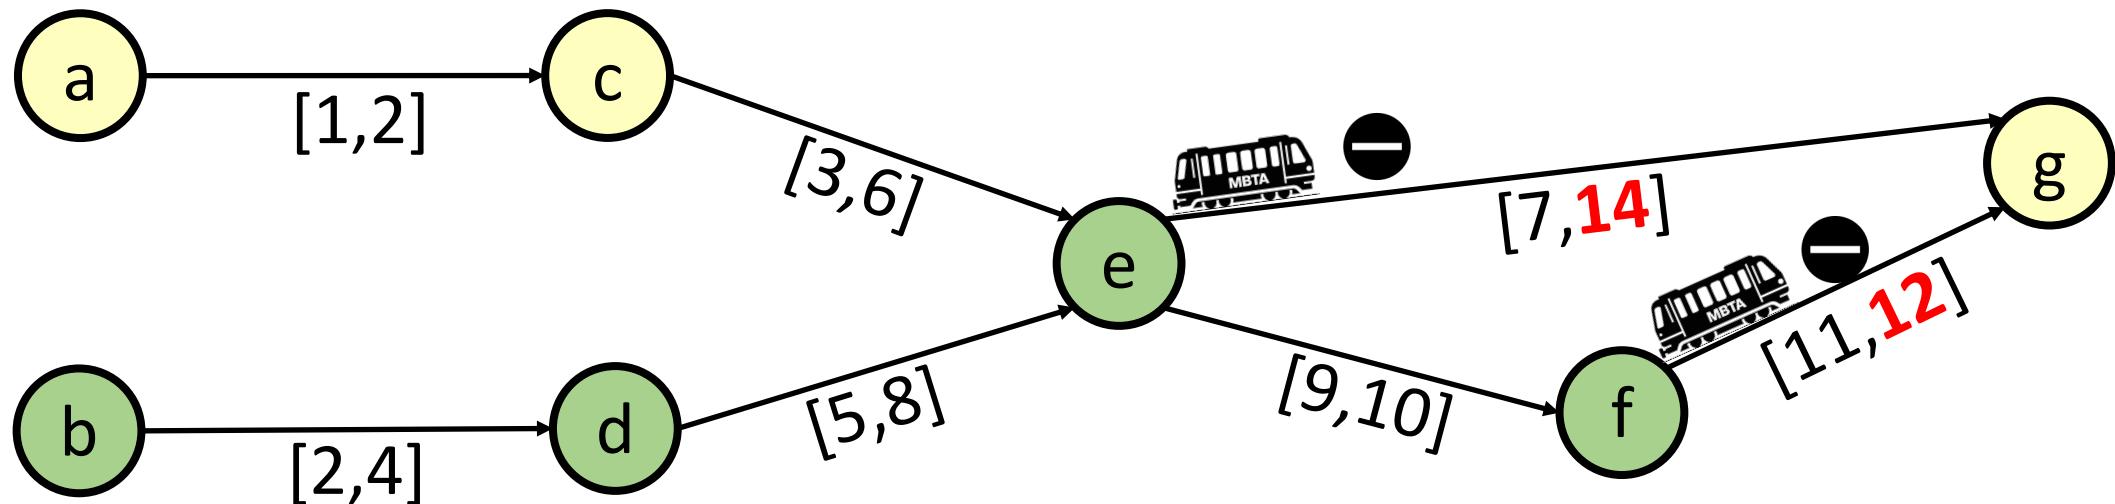
$u = a$	a	b	c	d	e	f	g
	1	$\infty$	2	$\infty$	6	10	12

$$[t_s, t_e] = [1, 15]$$

$u = b$	a	b	c	d	e	f	g
	$\infty$	2	$\infty$	4	8	10	$\infty$

$$[t_s, \mathbf{t}_e] = [2, 11]$$

# Example T. Minimal Path: Earliest Arrival



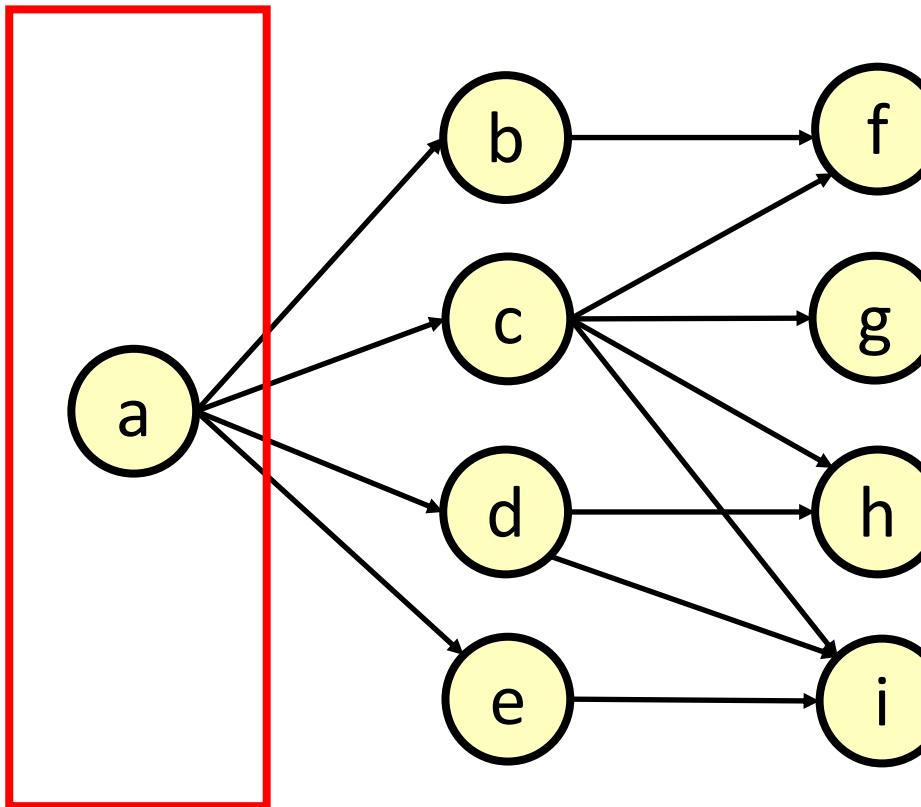
	a	b	c	d	e	f	g
$u = a$	1	$\infty$	2	$\infty$	6	10	12

$$[t_s, t_e] = [1, 15]$$

	a	b	c	d	e	f	g
$u = b$	$\infty$	2	$\infty$	4	8	10	$\infty$

$$[t_s, \mathbf{t}_e] = [2, 11]$$

# Efficient parallel implementations



...

Extends **Ligra** to temporal setting:

- can process **each frontier in parallel**

...

**temporalEdgeMap** (vertexSet, ...):

- applies fn to all vertices in each parallel frontier

...

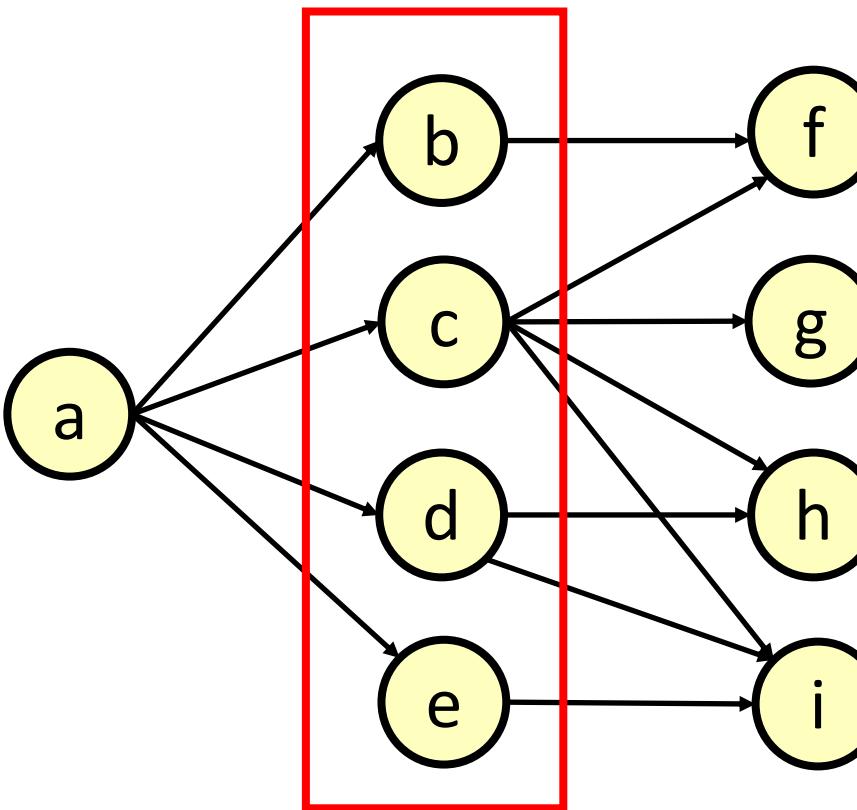
**update (temporalEdge)**:

- applies fn to all nbors in each parallel frontier

...

see thesis for details

# Efficient parallel implementations



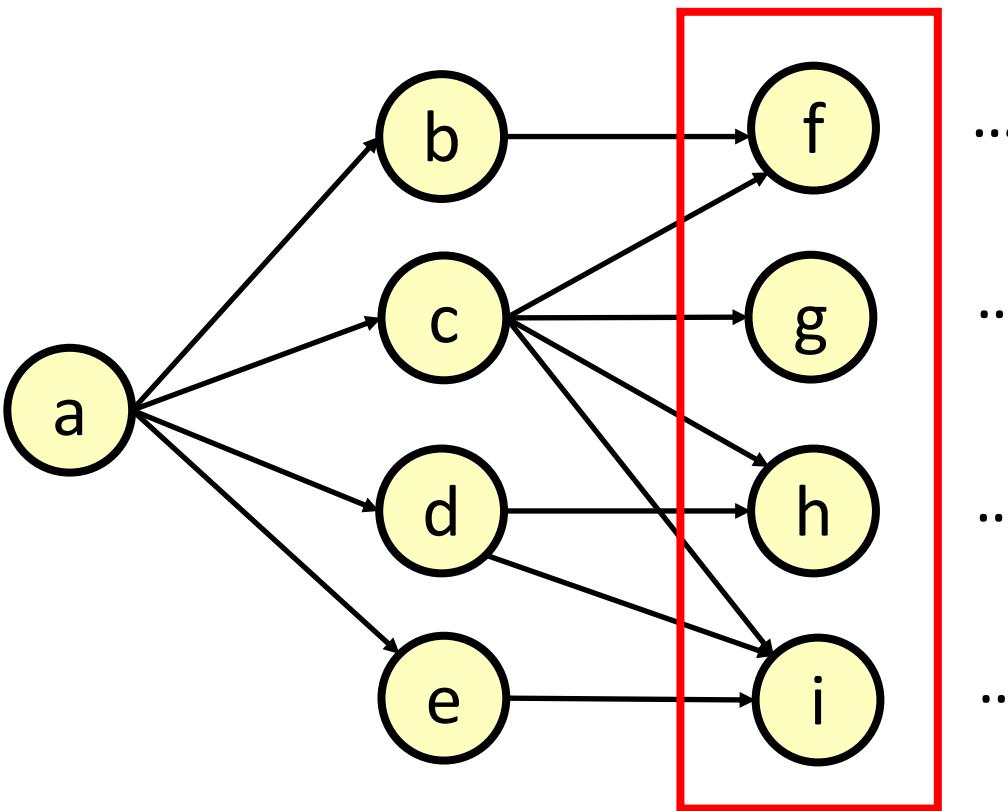
Extends **Ligra** to temporal setting:  
- can process **each frontier in parallel**

**temporalEdgeMap** (vertexSet, ...):  
- applies fn to all vertices in each parallel frontier

**update (temporalEdge)**:  
- applies fn to all nbors in each parallel frontier

see thesis for details

# Efficient parallel implementations



Extends **Ligra** to temporal setting:

- can process **each frontier in parallel**

**temporalEdgeMap** (vertexSet, ...):

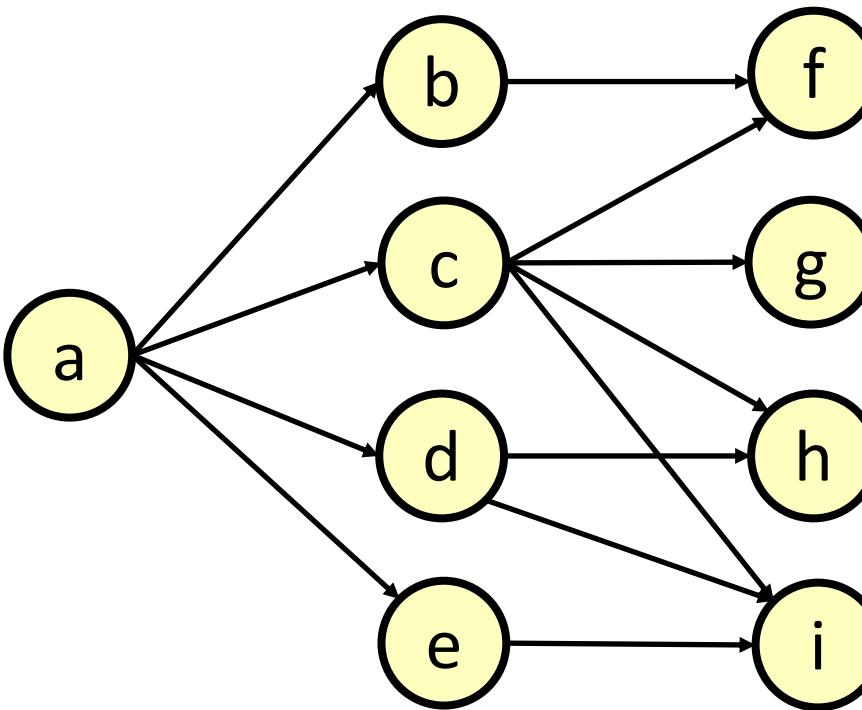
- applies fn to all vertices in each parallel frontier

**update (temporalEdge)**:

- applies fn to all nbors in each parallel frontier

see thesis for details

# Efficient parallel implementations



...

Extends **Ligra** to temporal setting:  
- can process **each frontier in parallel**

...

**temporalEdgeMap** (vertexSet, ...):  
- applies fn to all vertices in each parallel frontier

...

**update (temporalEdge)**:  
- applies fn to all nbors in each parallel frontier

...

see thesis for details

Julian Shun et al, “Ligra: A Lightweight Graph Processing for Shared Memory.”

Julian Shun et al, “Smaller and Faster: Parallel Processing of Compressed Graphs with Ligra+”

# Kairos API extends Ligra to temporal setting

Interface	Description
VertexSet	Represents a subset of vertices $V' \subseteq V$
TemporalEdgeSet	Represents a subset of temporaledges $E' \subseteq E$ .
OrderingPredicate( $A$ : temporaledge, $B$ : temporaledge, $T$ : OrderingPredicate-Type): bool	Evaluates the order between two temporal edges based on one of the three ordering predicate types: Succeeds, StrictlySucceeds, or Overlaps
TemporalGraph( $V$ : VertexSet, $E$ : TemporalEdgeSet, $P$ : OrderingPredicate)	Constructs a temporal graph using given input vertices, temporaledges, and ordering predicate.
VertexMap( $U$ : VertexSet, $G$ : TemporalGraph, $F$ : vertex $\rightarrow$ bool): VertexSet	Applies $F(u)$ for each $u \in U$ ; returns a VertexSet{ $u \in U \mid F(u)=\text{true}$ }.
TemporalEdgeMap( $U$ :TemporalEdgeSet, $G$ :TemporalGraph, $F$ :temporaledge $\rightarrow$ bool): TemporalEdgeSet	Applies $F(u)$ for each $u \in U$ ; returns a TemporalEdgeSet{ $u \in U \mid F(u)=\text{true}$ }.

# Kairos: Key Technical Contributions

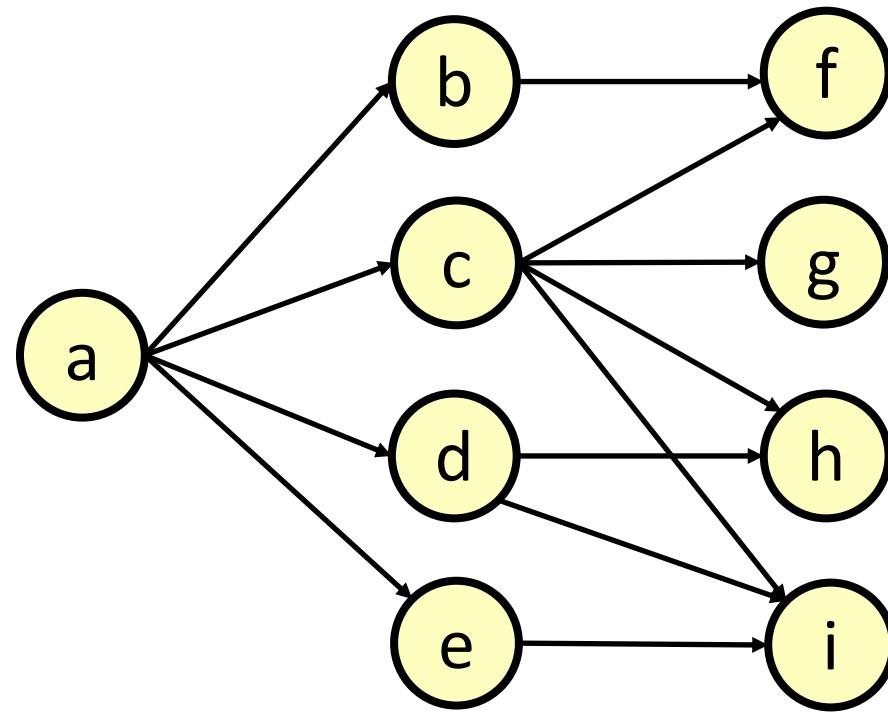
Efficient **parallel implementation of algorithms over temporal graphs**, e.g.  
minimal paths (earliest arr., latest dep., fastest, shortest duration)  
centrality (betweenness using fastest or shortest duration)  
k-core, ranking, and connectivity

Efficient data layouts

**T-CSR: compressed sparse row** for temporal graphs

**TGER**: highly optimized **parallel index** for temporal edges

# Problem: efficient graph data layout

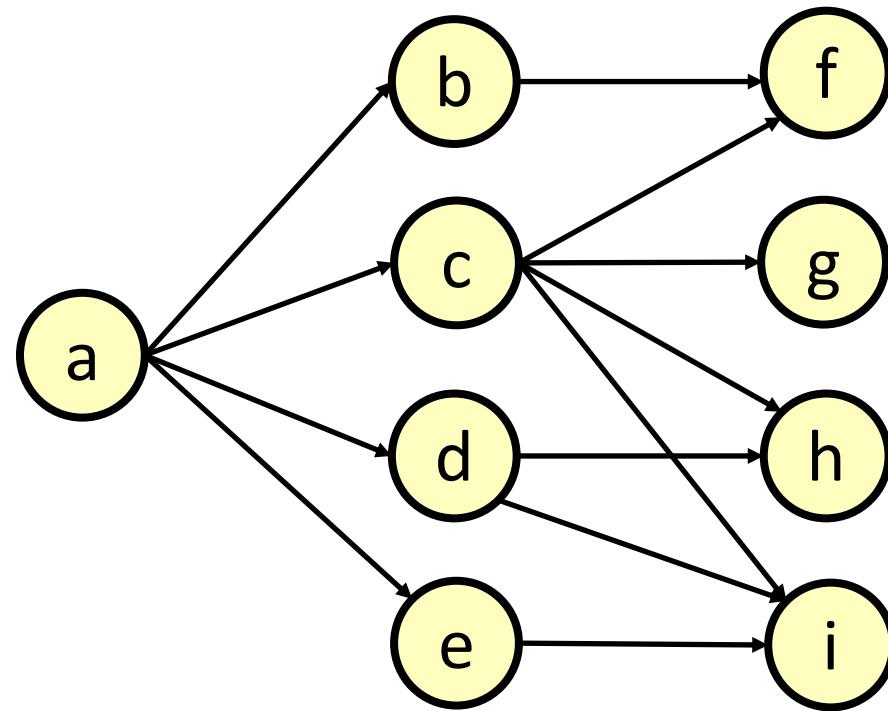


	a	b	c	d	e	f	g	h	i
a	0	1	1	1	1	0	0	0	0
b	0	0	0	0	0	1	0	0	0
c	0	0	0	0	0	1	1	1	1
d	0	0	0	0	0	0	0	1	1
e	0	0	0	0	0	0	0	0	1
f	0	0	0	0	0	0	0	0	0
g	0	0	0	0	0	0	0	0	0
h	0	0	0	0	0	0	0	0	0
i	0	0	0	0	0	0	0	0	0

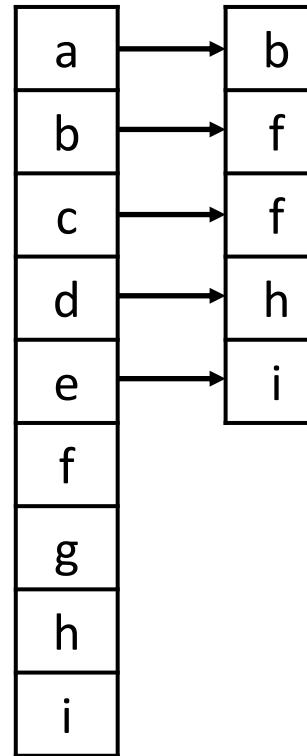
>85%  
zeros

*Adjacency matrix is not space efficient for sparse graphs*

# Problem: efficient graph data layout



*adjacency list:*



*pointer chasing*

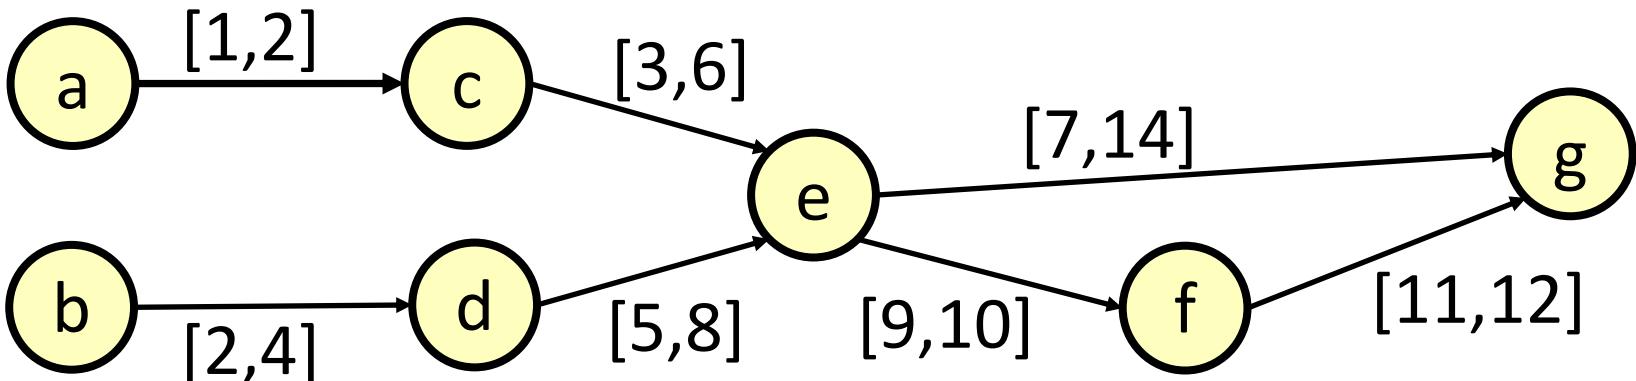
*need scan to find  
src vertex offset*

*edge list:*

(a,b)  
(a,c)  
(a,d)  
(a,e)  
(b,f)  
(c,f)  
(c,g)  
(c,h)  
(c,i)  
(d,h)  
(e,i)

*Adjacency list* and *edge list* use less space, but are not cache efficient

# Solution: T. Compressed Sparse Row (T-CSR)



*position where  
nbors start in the  
other arrays*

offsets:

a	b	c	d	e	f	g
0	1	2	3	4	6	6

dest. id:

--	--	--	--	--	--	--

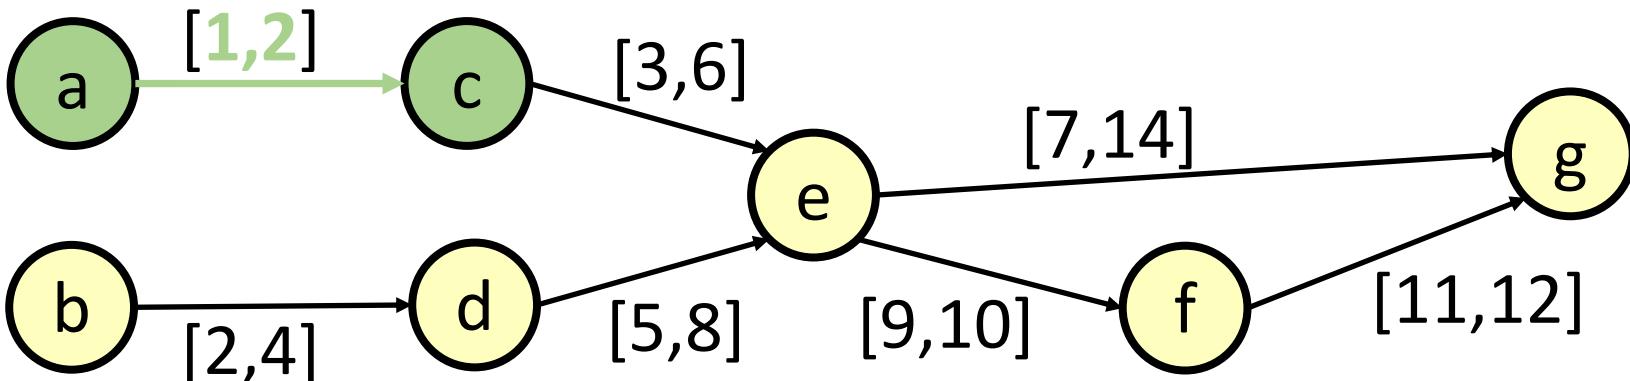
s. time:

--	--	--	--	--	--	--

e. time:

--	--	--	--	--	--	--

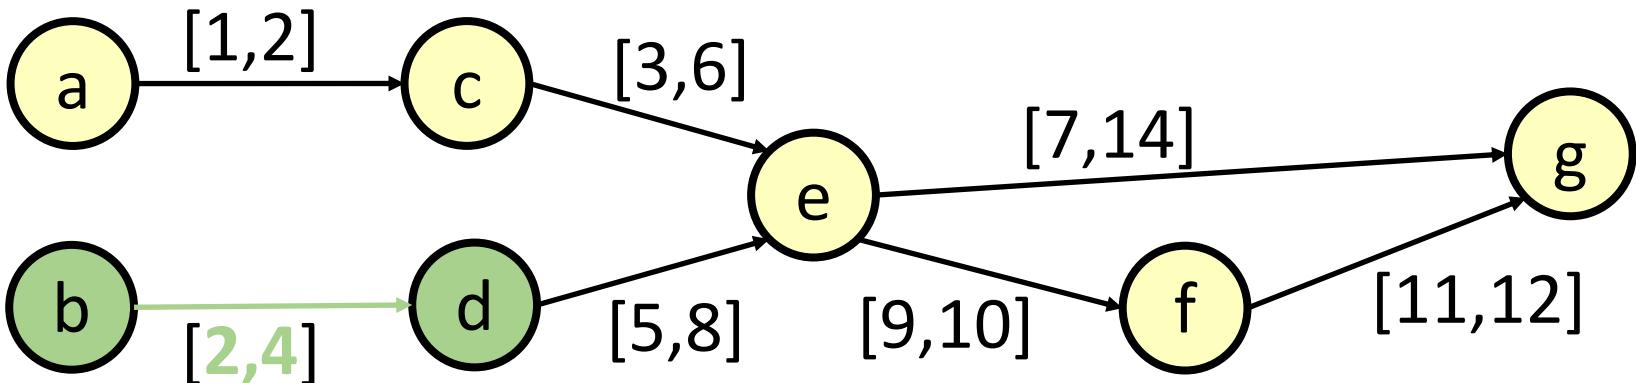
# Solution: T. Compressed Sparse Row (T-CSR)



	a	b	c	d	e	f	g
offsets:	0	1	2	3	4	6	6
dest. id:	c						

s. time:	1						
e. time:	2						

# Solution: T. Compressed Sparse Row (T-CSR)



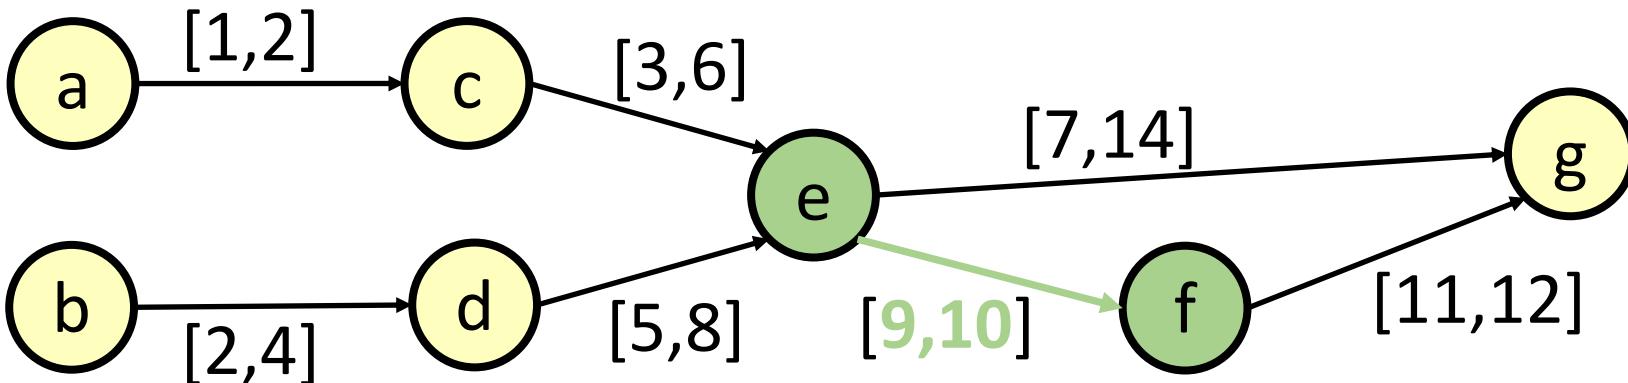
	a	b	c	d	e	f	g
offsets:	0	1	2	3	4	6	6
dest. id:	c	d					

A red arrow points to the second cell of the 'dest. id' row, indicating the start of the edge list for node b.

s. time:	1	2						
e. time:	2	4						

Two red arrows point to the second cell of each of the two time arrays, indicating the start of the edge list for node b.

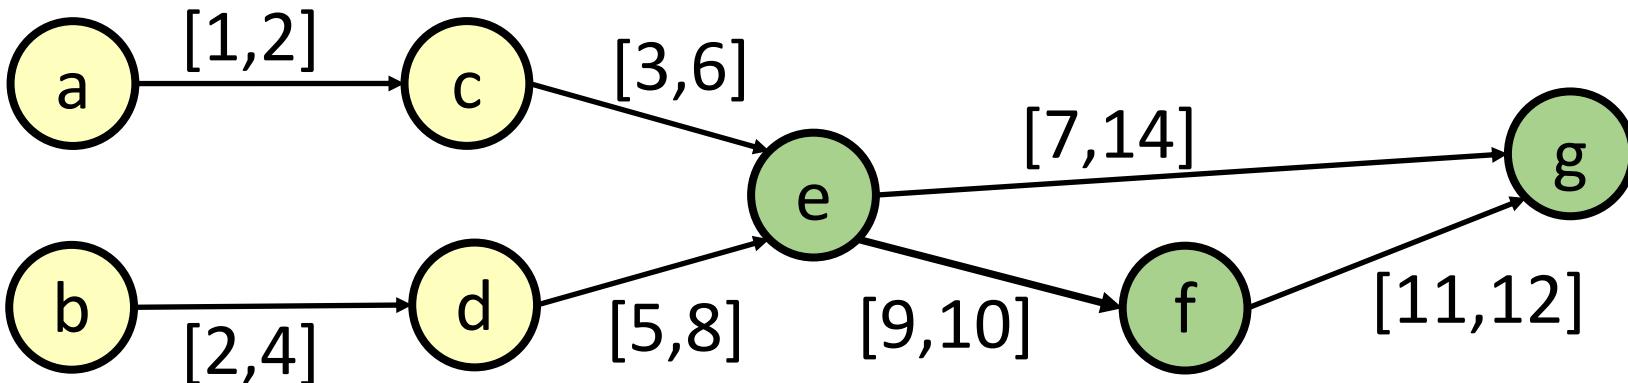
# Solution: T. Compressed Sparse Row (T-CSR)



	a	b	c	d	e	f	g
offsets:	0	1	2	3	4	6	6
dest. id:	c	d	e	e	f		

s. time:	1	2	3	5	9		
e. time:	2	4	6	8	10		

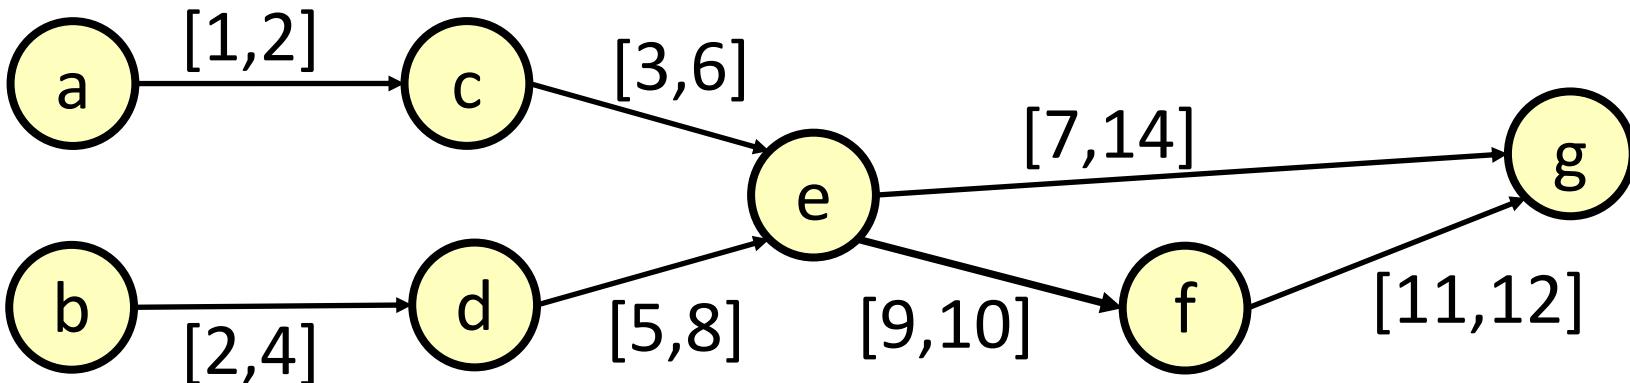
# Solution: T. Compressed Sparse Row (T-CSR)



	a	b	c	d	e	f	g
offsets:	0	1	2	3	4	6	6
dest. id:	c	d	e	e	f	g	

s. time:	1	2	3	5	9	7	
e. time:	2	4	6	8	10	14	

# Solution: T. Compressed Sparse Row (T-CSR)



	a	b	c	d	e	f	g
offsets:	0	1	2	3	4	6	6

dest. id:	c	d	e	e	f	g	g
-----------	---	---	---	---	---	---	---

s. time:	1	2	3	5	9	7	11
----------	---	---	---	---	---	---	----

e. time:	2	4	6	8	10	14	12
----------	---	---	---	---	----	----	----

# Problem: T-CSR still requires linear scan

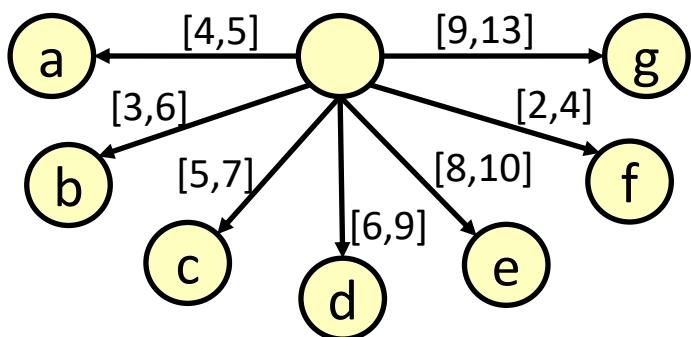
Parallel scan starting at vertex offset is cache efficient, but still  $O(\text{degree})$

In real-world graphs, single vertex degree can be  $O(\text{Millions})$

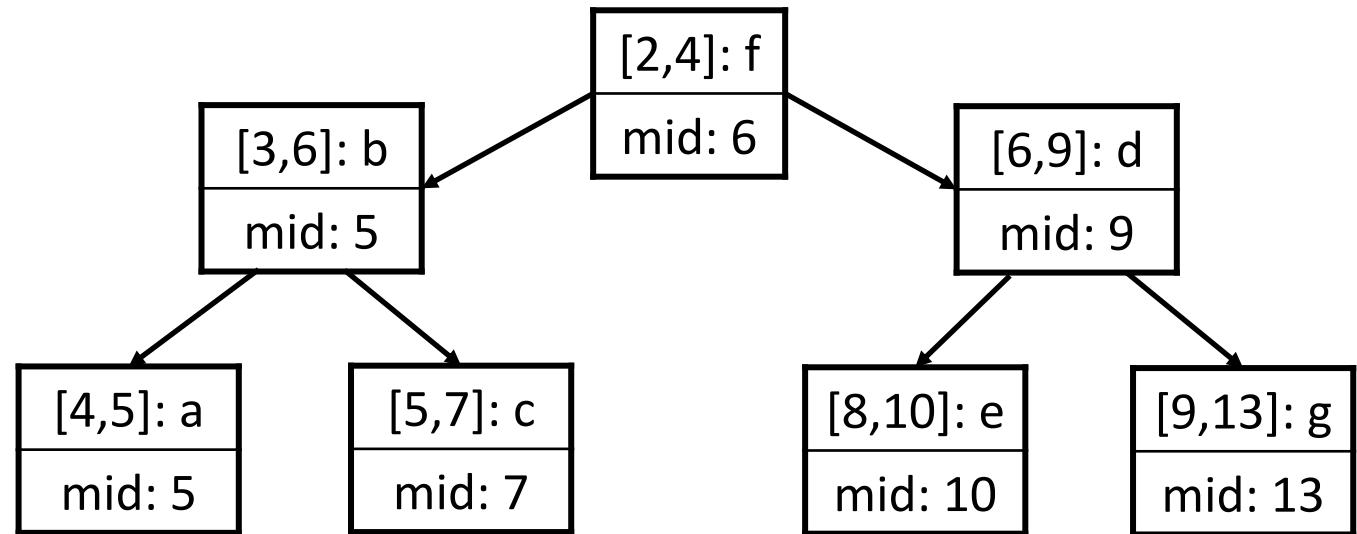
Name	$ V $	$ E $	$\max_{v \in V} \deg^+(v)$	$\max_{v \in V} \deg^-(v)$	$\text{avg}_{v \in V} \deg(v)$
bitcoin [37]	$4.80 \times 10^7$	$1.13 \times 10^8$	$2.66 \times 10^6$	$2.53 \times 10^6$	4
netflow [38]	$3.72 \times 10^8$	$1.20 \times 10^9$	$5.78 \times 10^5$	$1.82 \times 10^8$	12
reddit-reply [37]	$1.17 \times 10^7$	$6.46 \times 10^8$	$3.92 \times 10^5$	$9.94 \times 10^5$	110
stackoverflow [39]	$6.02 \times 10^6$	$6.34 \times 10^7$	$1.01 \times 10^5$	93143	20
transportation [40]	41794	$7.93 \times 10^7$	50881	50625	3798
twitter-cache [41]	94226	$8.55 \times 10^8$	$5.21 \times 10^6$	$2.31 \times 10^5$	36328
synthetic	$10^7$	$10^9$	$5.65 \times 10^7$	$3.08 \times 10^7$	99

# Solution: T. Graph Edge Registry (TGER)

input vertex out-neighbors:

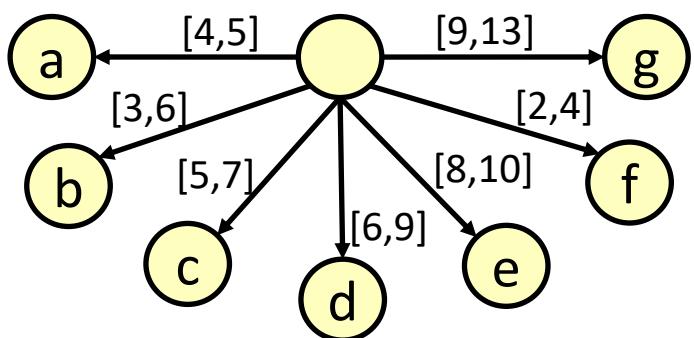


Corresponding TGER:

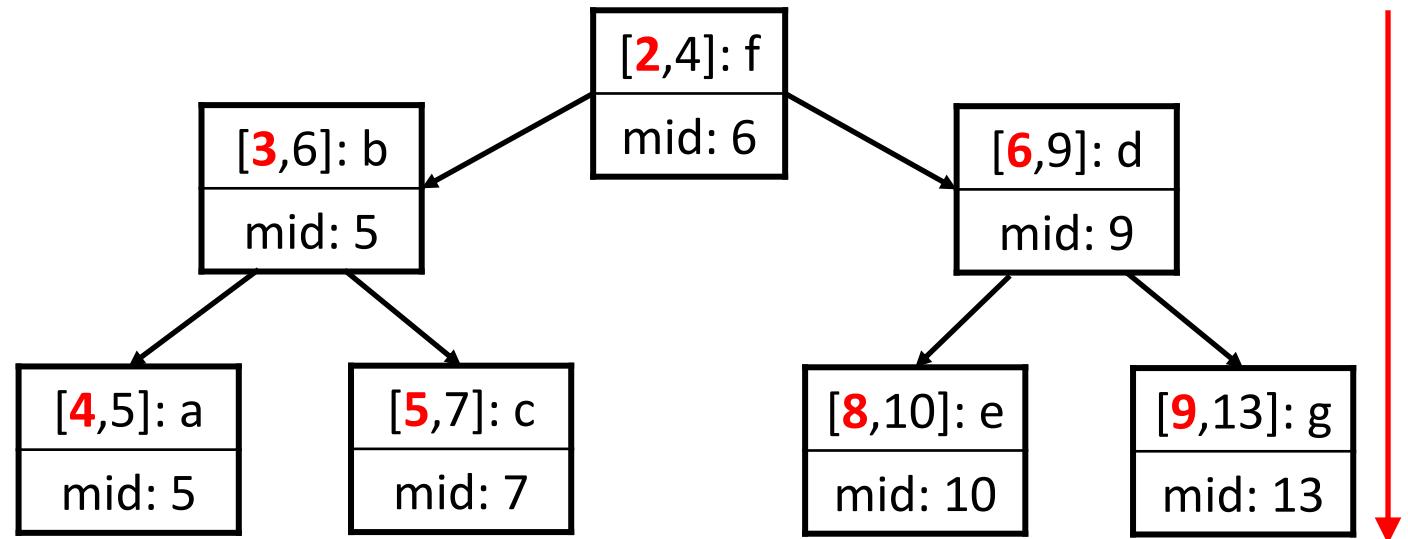


# Solution: T. Graph Edge Registry (TGER)

input vertex out-neighbors:



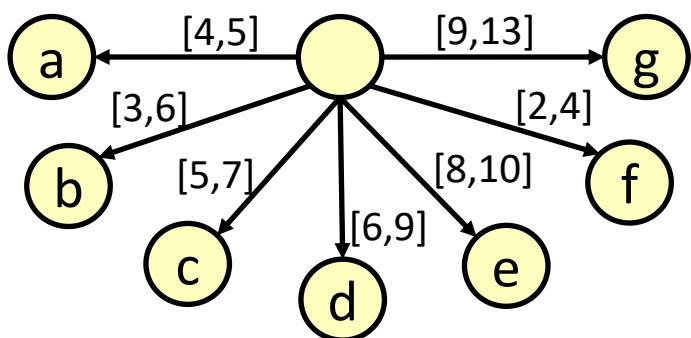
Corresponding TGER:



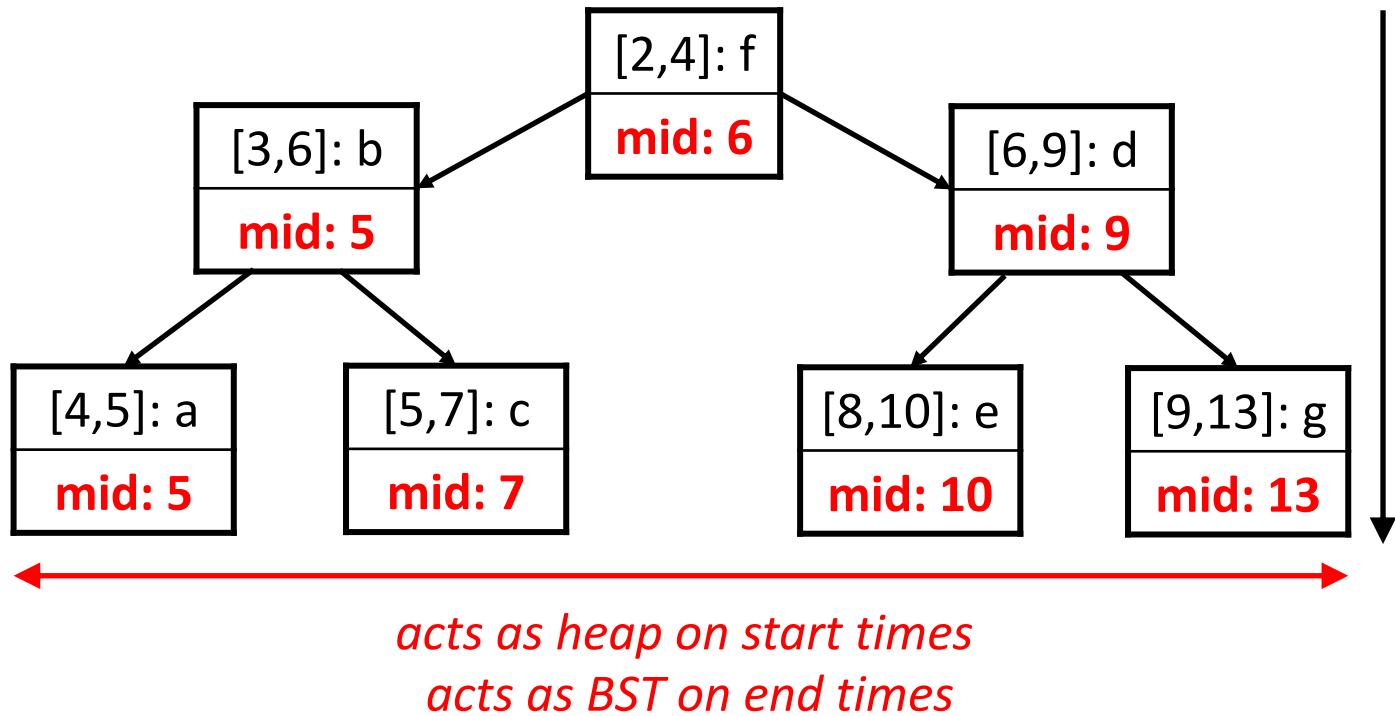
*acts as heap on start times*

# Solution: T. Graph Edge Registry (TGER)

input vertex out-neighbors:

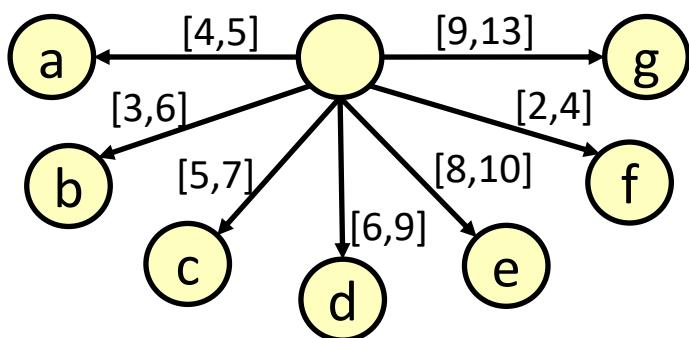


Corresponding TGER:

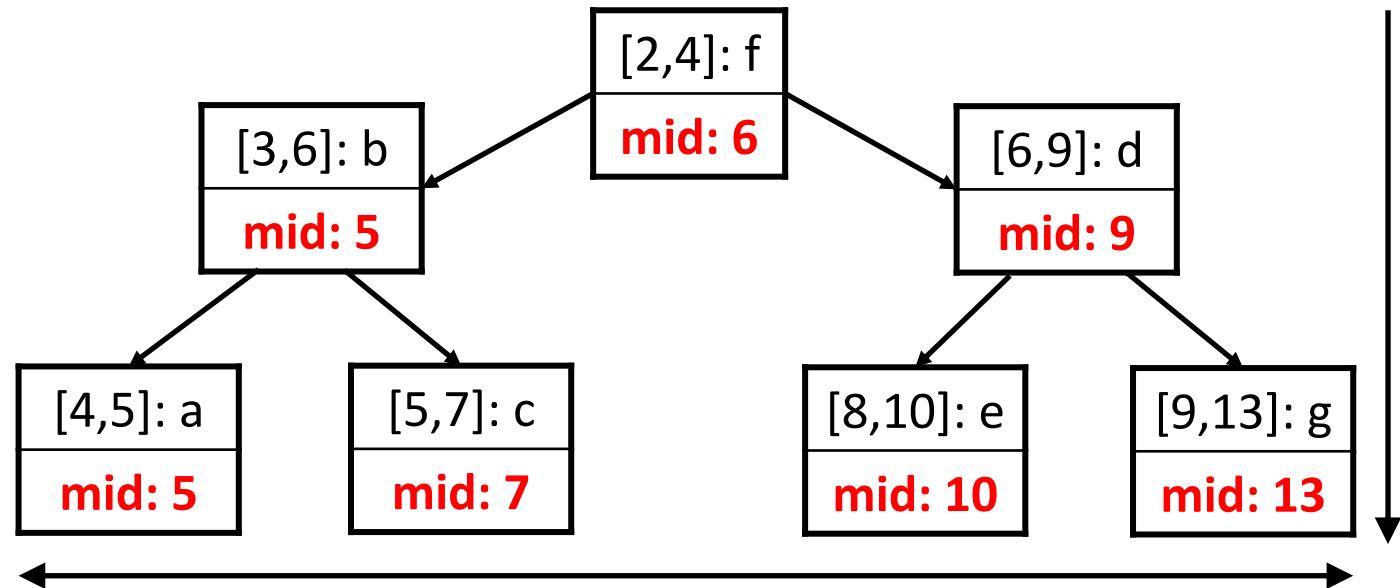


# Solution: T. Graph Edge Registry (TGER)

input vertex out-neighbors:



Corresponding TGER:



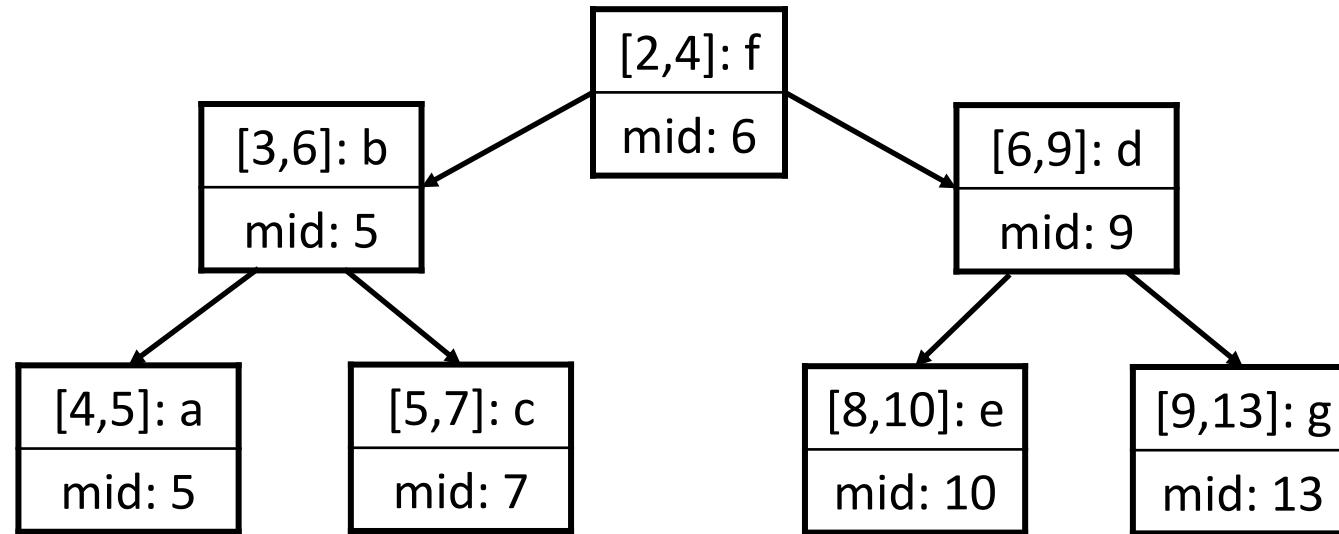
parallel construction (see thesis for details):

- parallel sort edges by *start time*, root is first element
- recurse: split subtrees by left and right of median *end time*

*acts as heap on start times*

*acts as BST on end times*

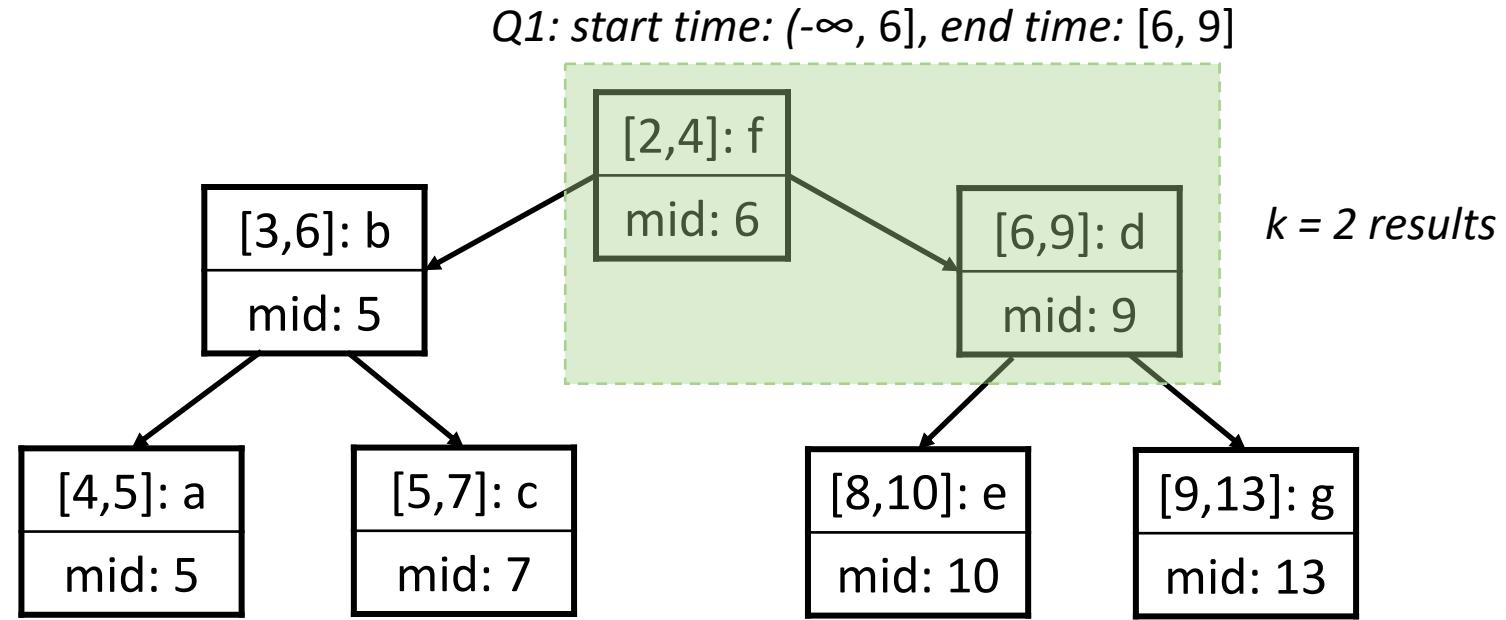
# TGER supports 3-sided queries in $\log(n) + k$



3-sided:

- 1-side (upper range) for heap dimension
- 2-sides (lower + upper range) for BST dimension

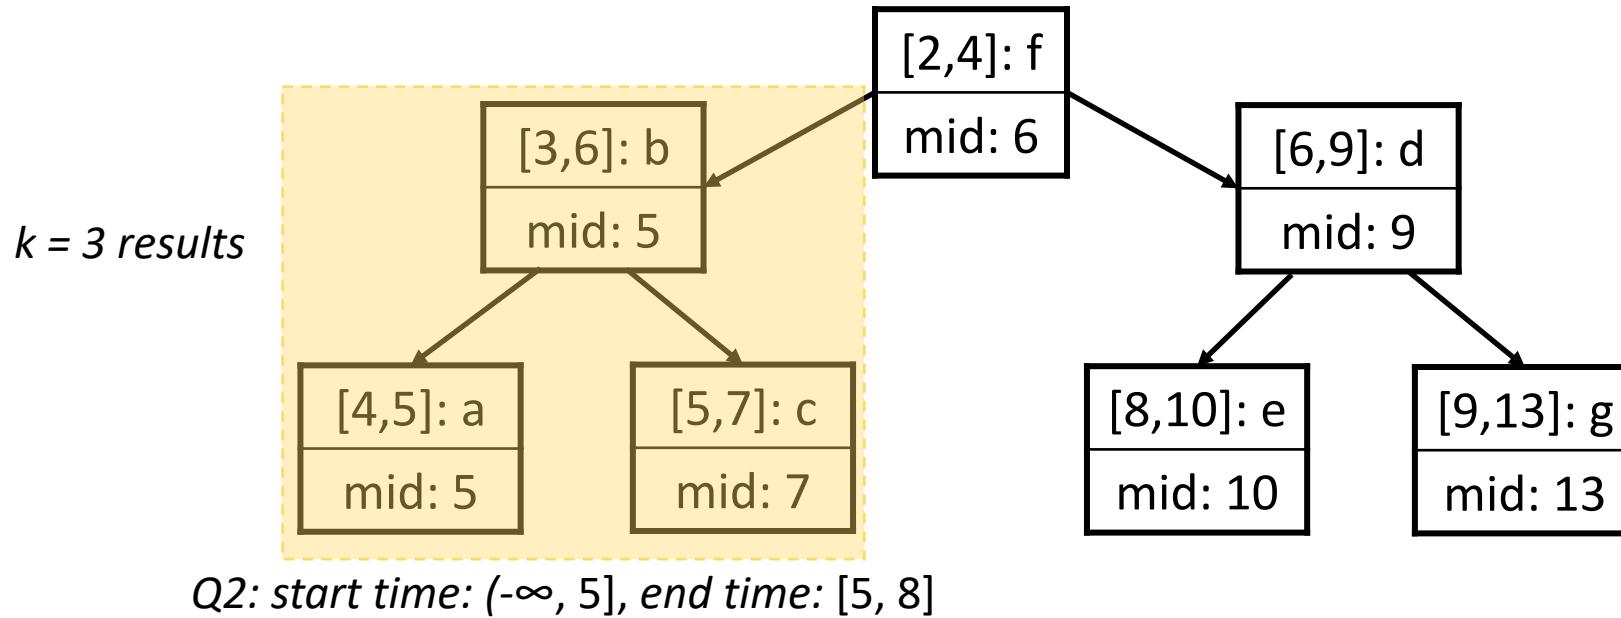
# TGER supports 3-sided queries in $\log(n) + k$



3-sided:

- 1-side (upper range) for heap dimension
- 2-sides (lower + upper range) for BST dimension

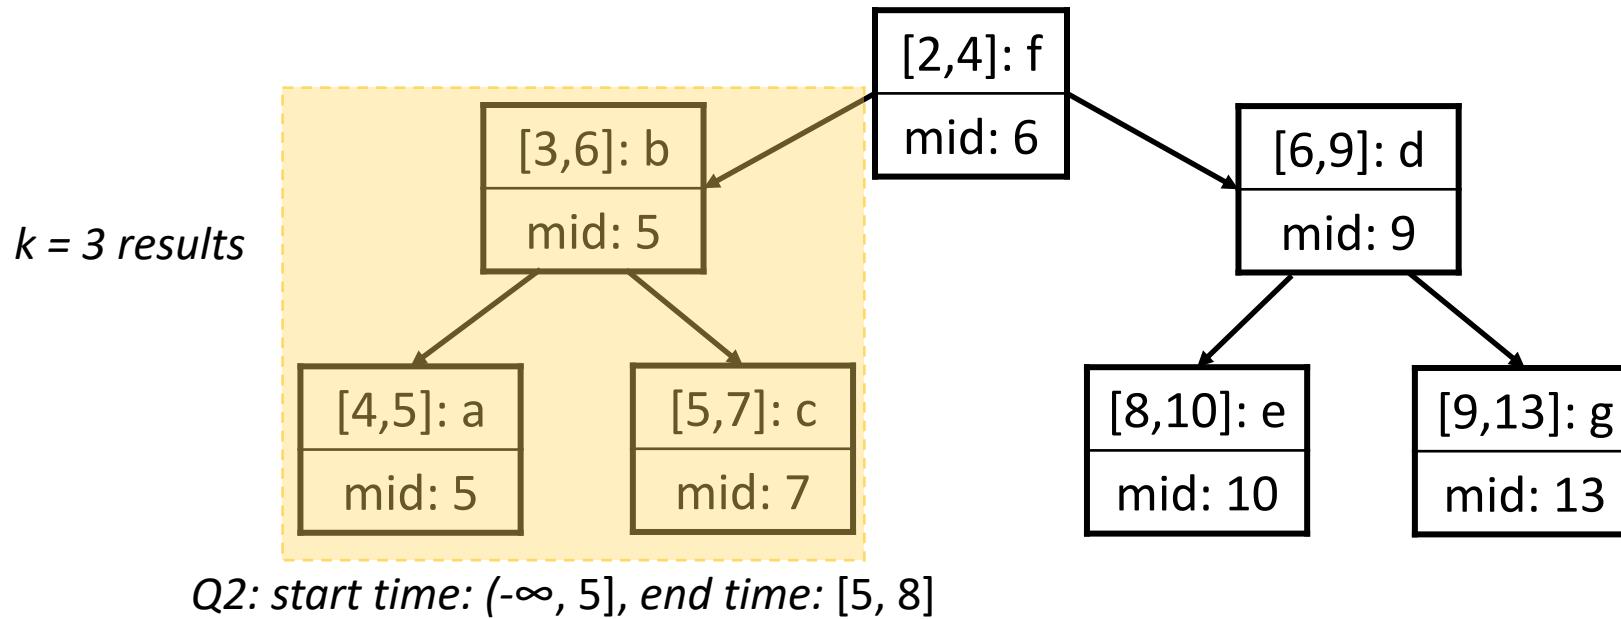
# TGER supports 3-sided queries in $\log(n) + k$



3-sided:

- 1-side (upper range) for heap dimension
- 2-sides (lower + upper range) for BST dimension

# TGER supports 3-sided queries in $\log(n) + k$



3-sided:

- 1-side (upper range) for heap dimension
- 2-sides (lower + upper range) for BST dimension

parallel query processing (see thesis for details):

- per worker: local array stores results
- final merge: prefix sum

# Problem: scan vs index trade-off

## T-CRS

Good parallelism: scan over arrays

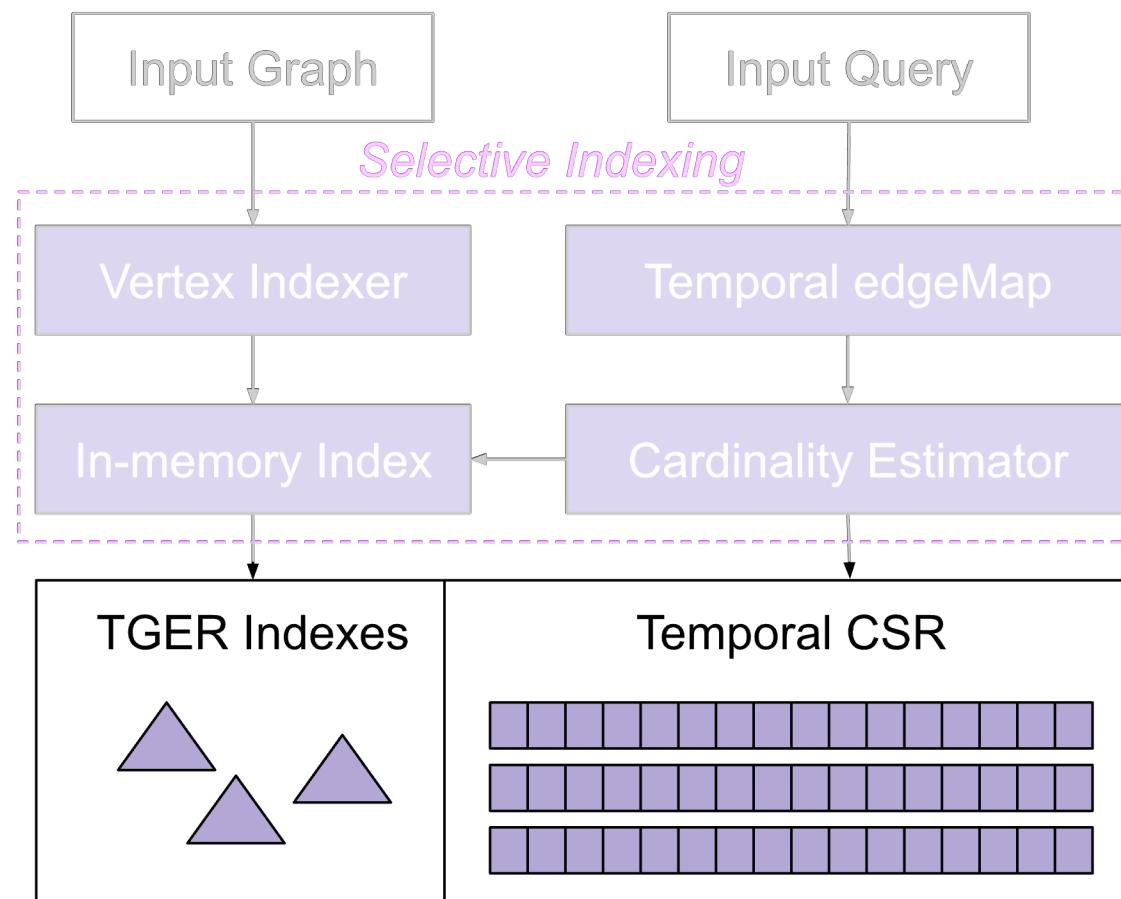
$O(n)$  query complexity: not great when  $n$  is large

## TGER

$O(\log n + k)$  query complexity: beats T-CRS when  $k \lll n$

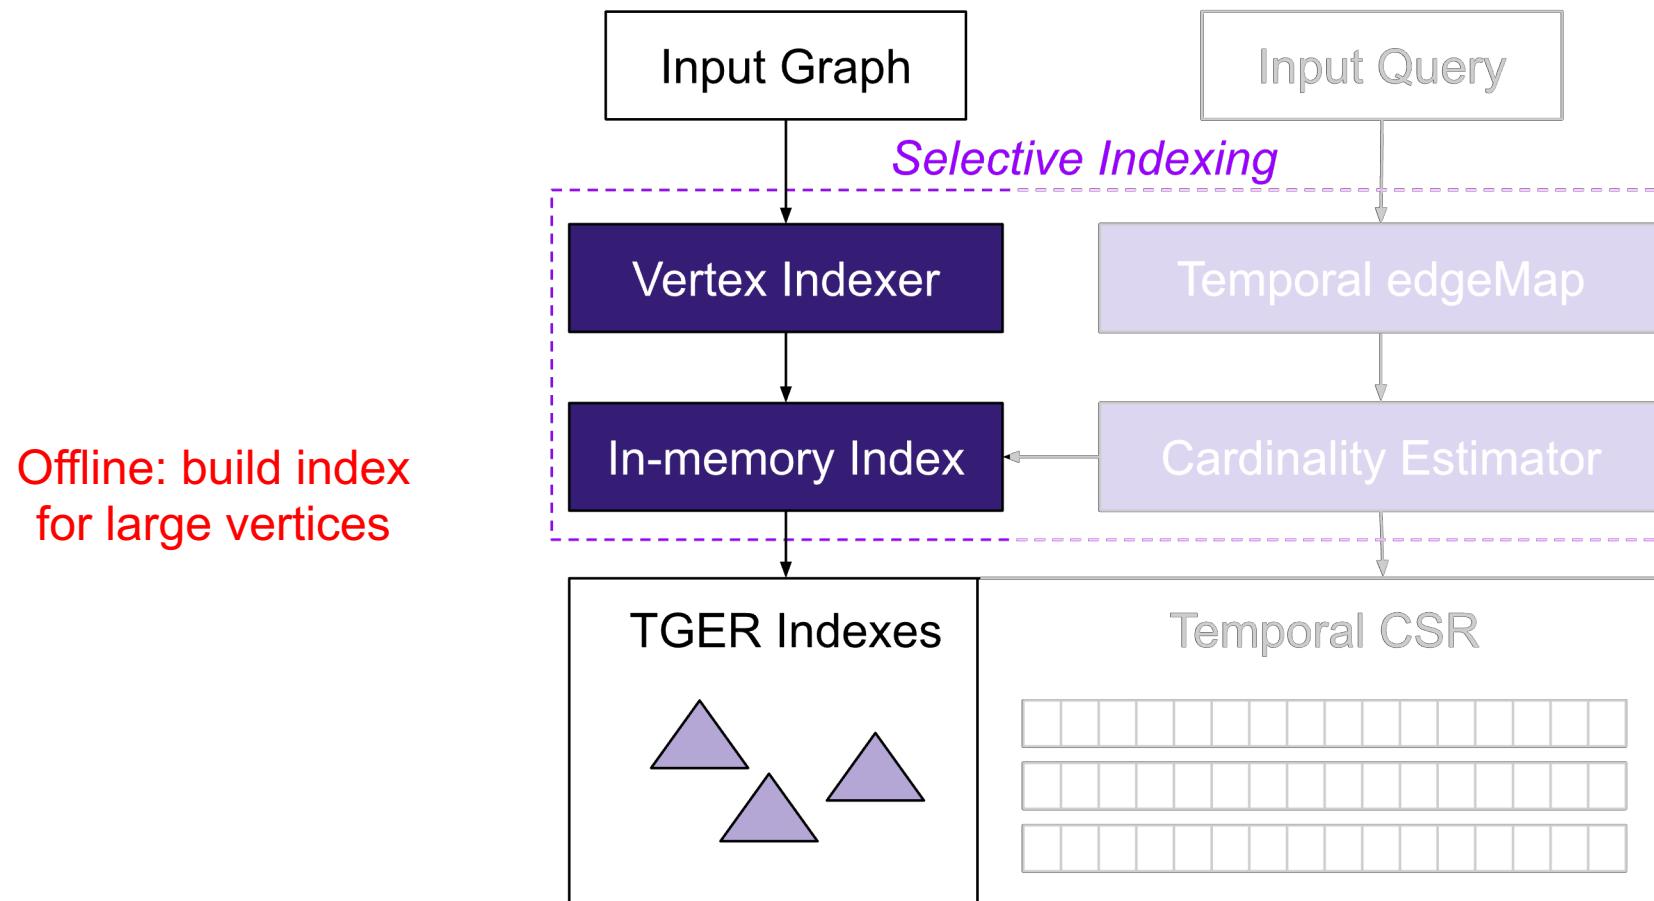
Tree-like data structure: parallelism is not as efficient

# Solution: Selective Indexing

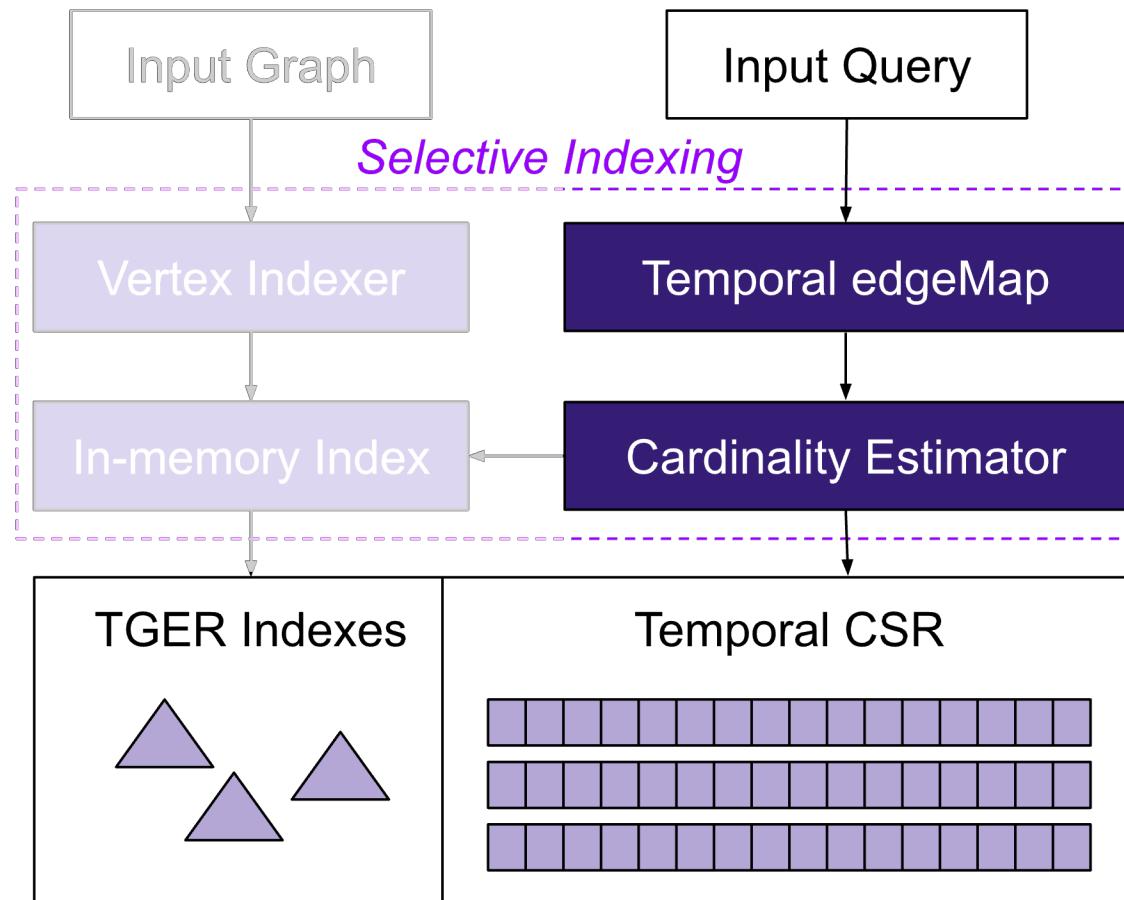


Two data representations for  
in-memory temporal edges

# Solution: Selective Indexing



# Solution: Selective Indexing



At runtime, cost model guides access method decision (index vs scan)

# Selective Indexing: Cost Model

$$T_v = c \cdot [\log(\deg(v)) + k]$$

cost of accessing nbors using TGER

$$S_v = c' \cdot \deg(v)$$

cost of accessing nbors using T-CSR

$$C_v = \begin{cases} T_v & \text{if } \beta \leq \theta_{\text{sel}} \\ S_v & \text{otherwise} \end{cases}$$

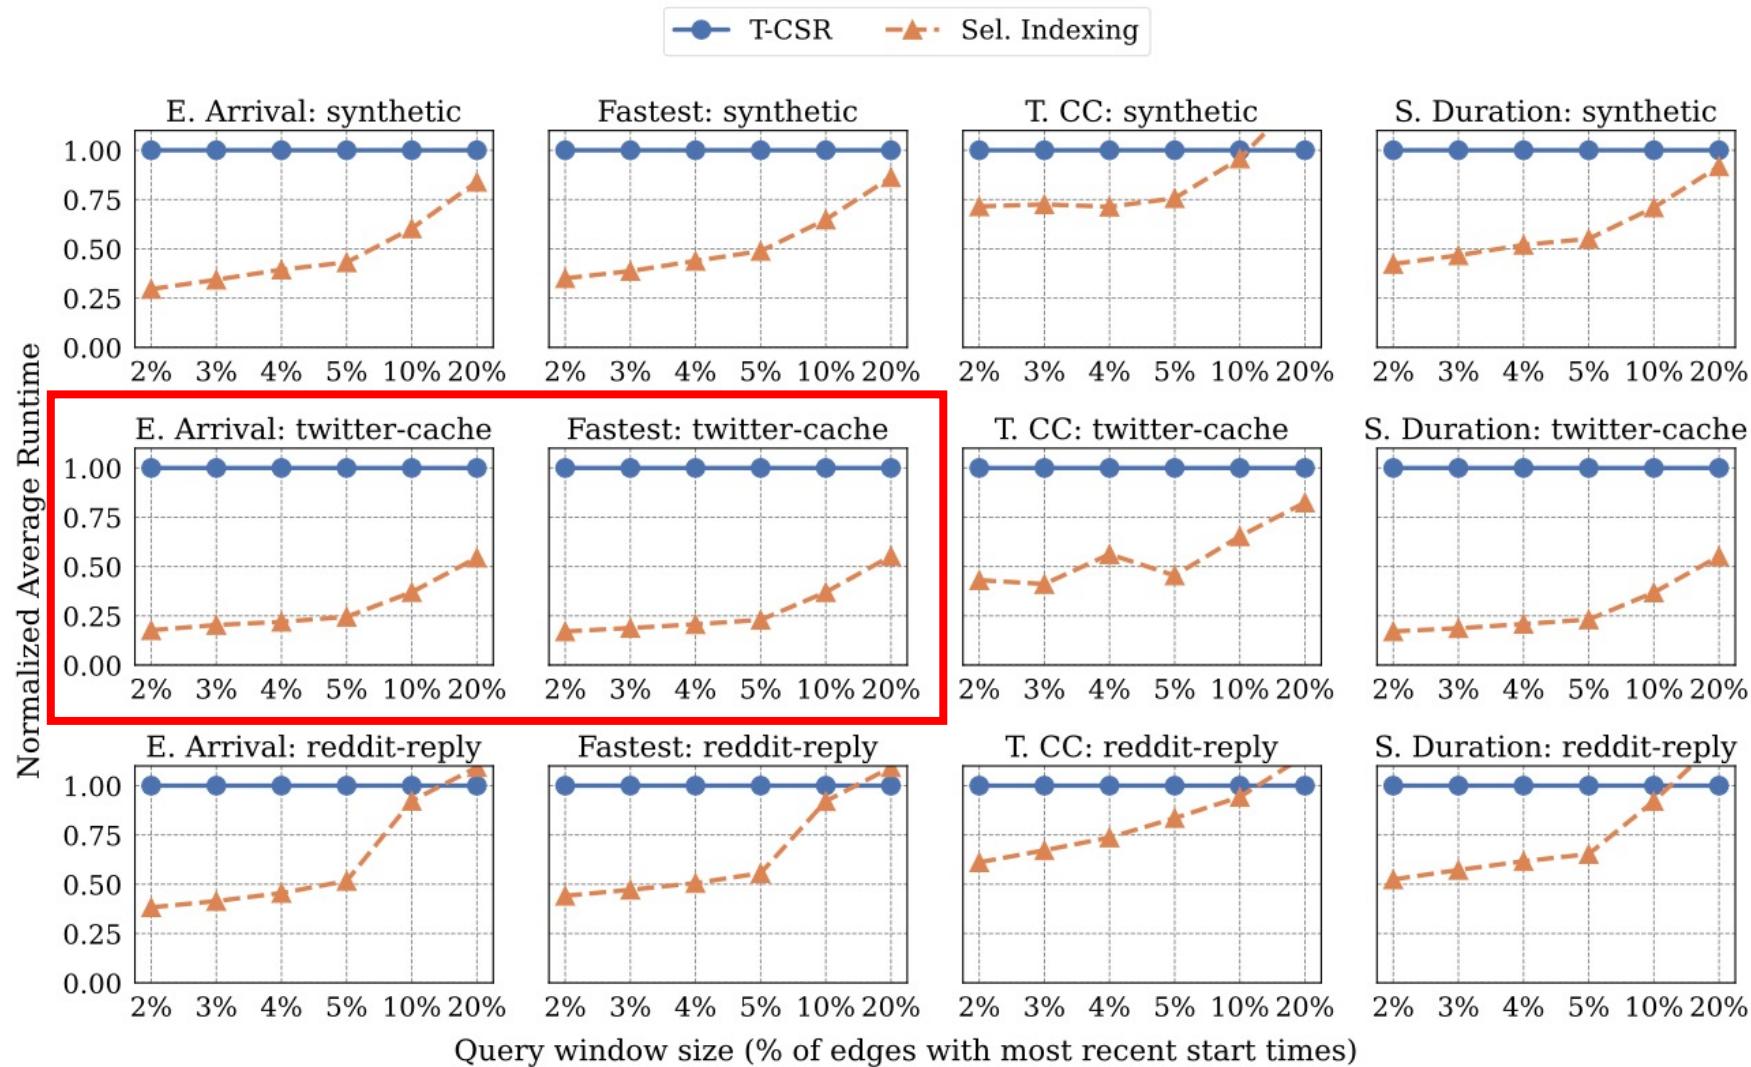
## Impacting factors

Vertex degree distribution

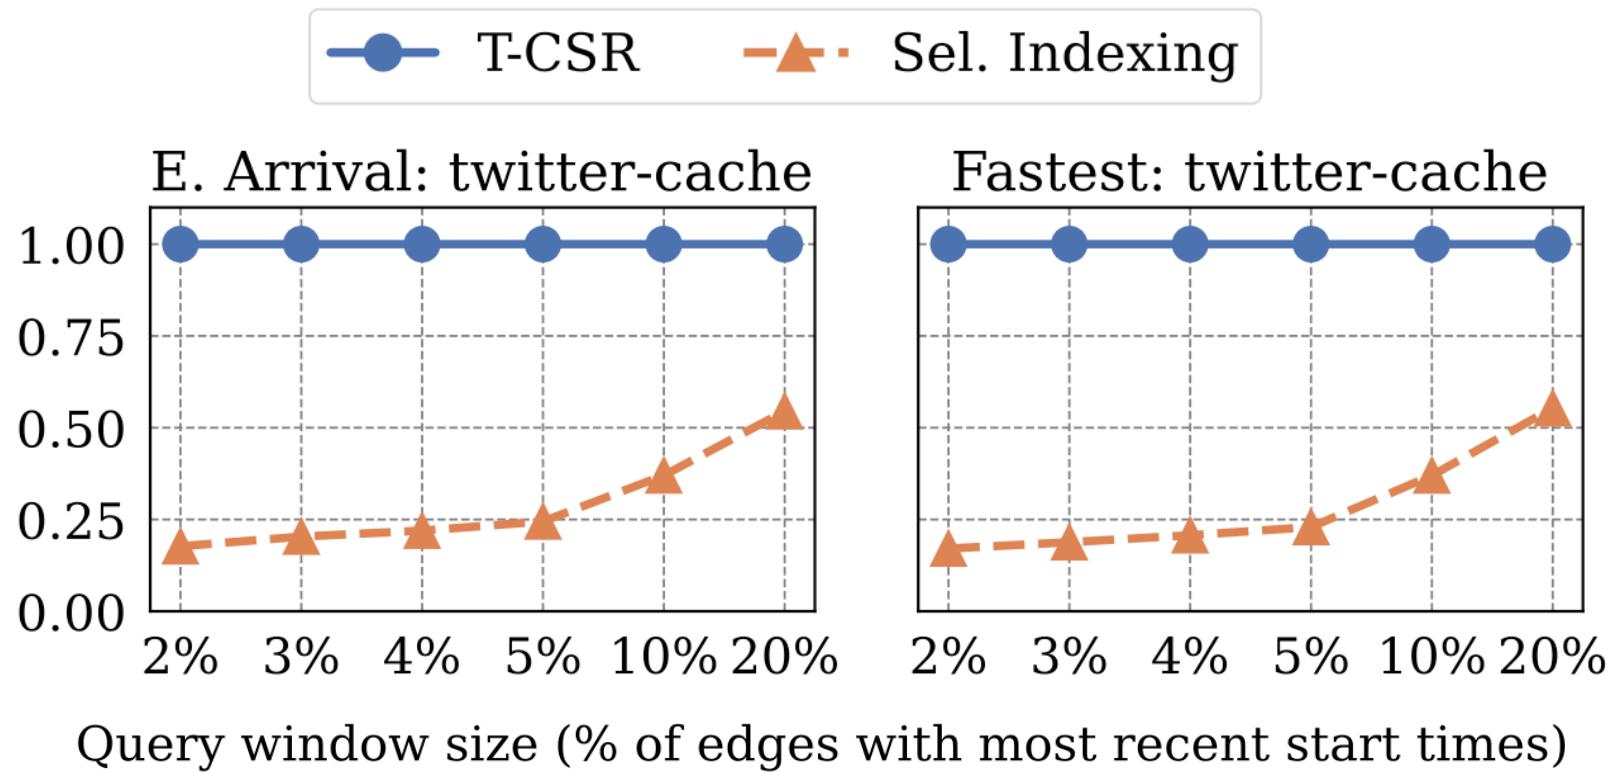
Temporal edges distribution

Query workload (e.g., selectivity)

# Experimental Results (see thesis for more)

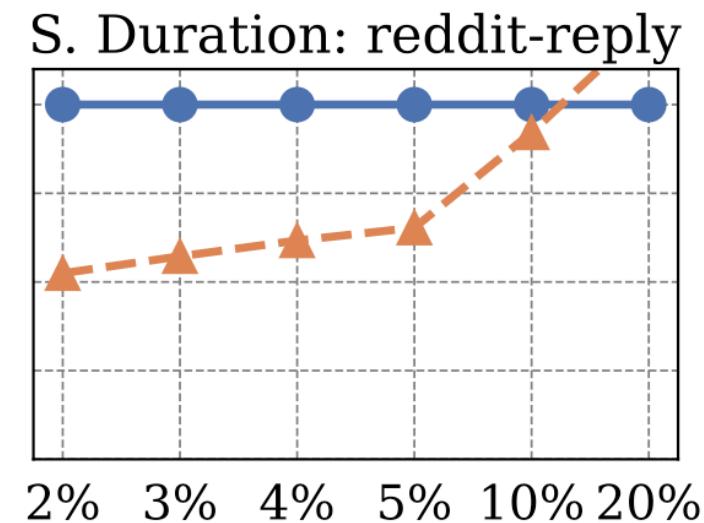
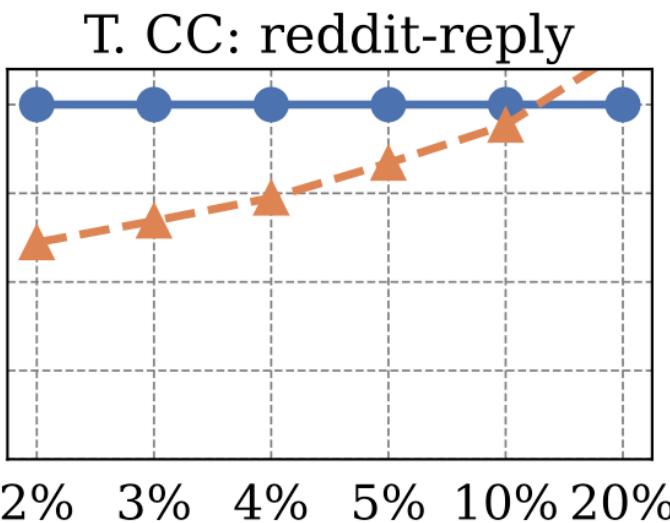


# Experimental Results (see thesis for more)



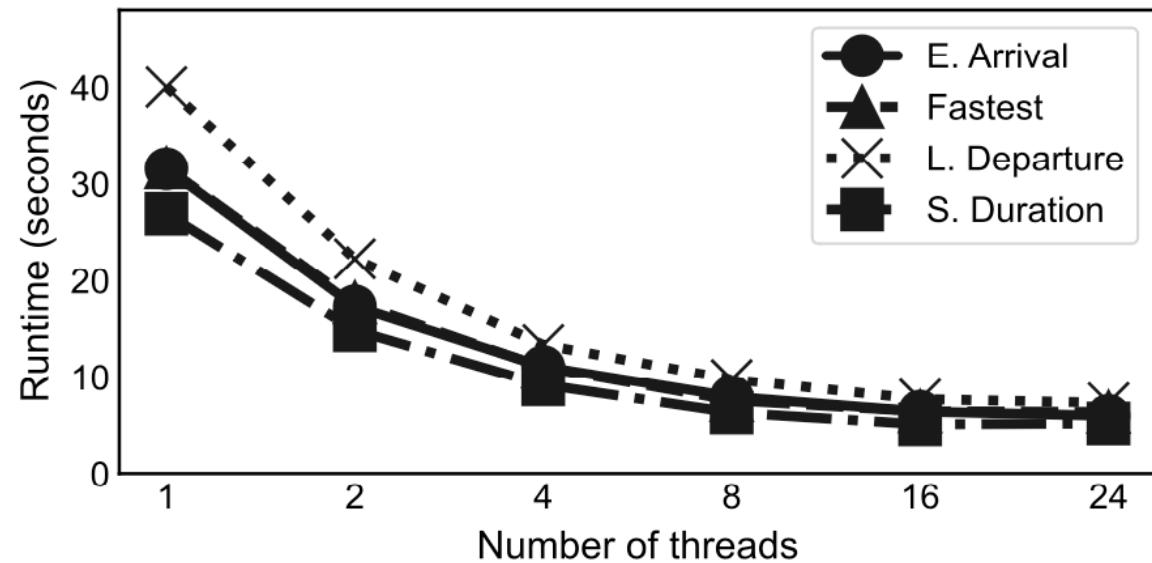
# Experimental Results (see thesis for more)

T-CSP      Sel. Indexing



Query window size (% of edges with most recent start times)

# Experimental Results (see thesis for more)



500M edges synthetic data set (aggressively skewed)

O(secs): orders of magnitude faster than distributed engines

# Kairos Contributions

**Selective Indexing** chooses best access method to retrieve vertex's nbors

TGER: highly optimized **parallel index** for temporal edges

Efficient parallel implementation of algorithms over temporal graphs

Good performance gains

Orders of magnitude vs distributed alternatives for graphs that fit within a single server  
T-CSR baseline already good, but TGER is better for highly selective queries

Other contributions

Work w/ Intel collaborators: Optane NVRAM for storing TGER indexes

Open-sourced T-CSR baseline last year; open source remaining after paper acceptance

# Future Work

## Kaskade

- Expand view types

- Incremental view maintenance + automated view management

## Kairos

- TGER and T-CSR dynamism: support edge removal / insertion

- Data compression + adaptive indexing instead of binary decision

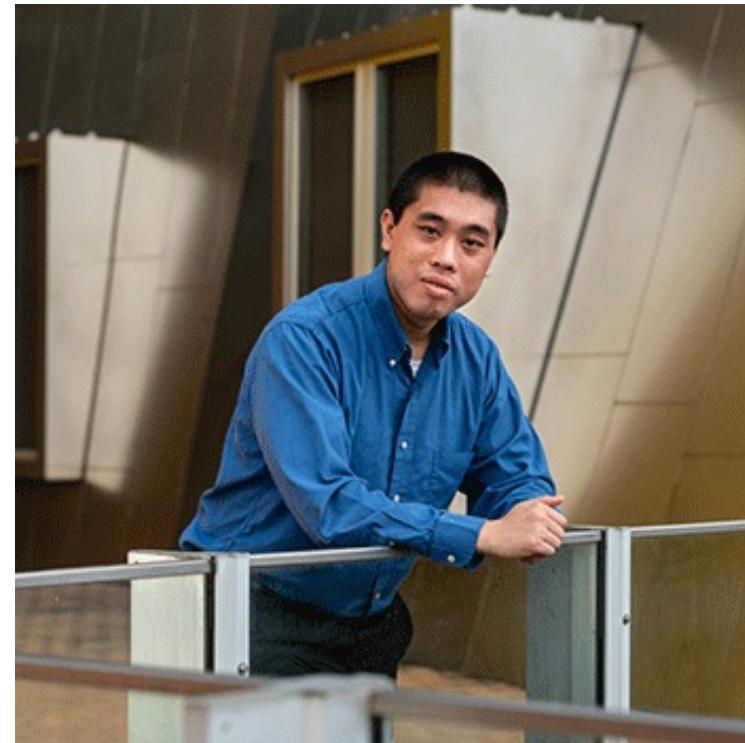
- Distributed processing: requires skew-aware partitioning

## Integrate both systems

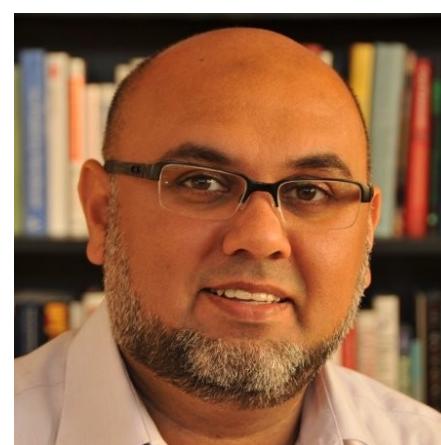
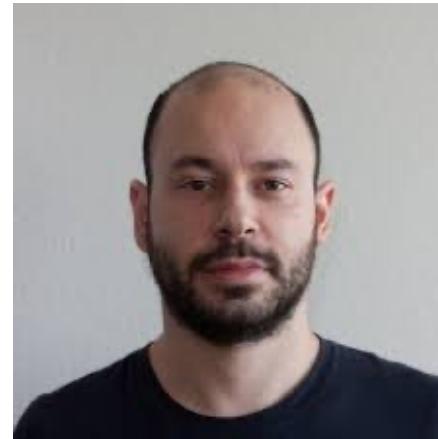
- Temporal views: time-respecting path contractions can lead to view explosion

- Selective indexing naturally extends to non-temporal settings

# Acknowledgements (Sam and Julian)



# Acknowledgements (Committee + Collabs)



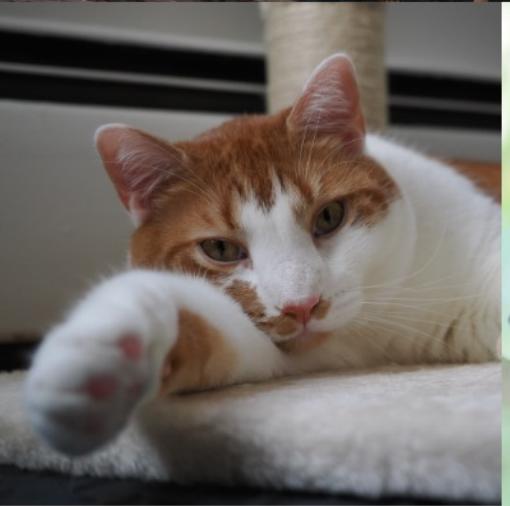
# Acknowledgements (MIT DSG)



# Acknowledgements (MIT + Friends)



# Acknowledgements (Family)

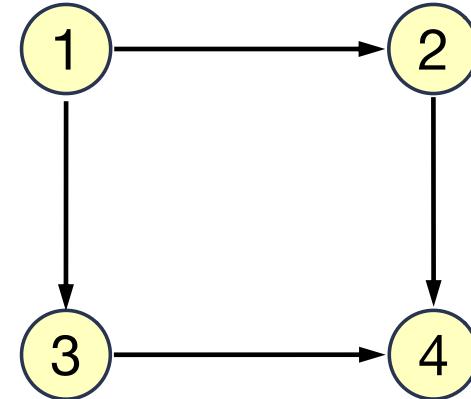


**THANK YOU!**

# Problem formulation: finding graph views as a satisfiability problem

Logical clause example (Prolog syntax):

```
path(X, Y) :- edge(X, Y).  
path(X, Y) :- path(X, Z), path(Z, Y).
```



Facts extracted from input queries and schema are axioms (*constraints*).

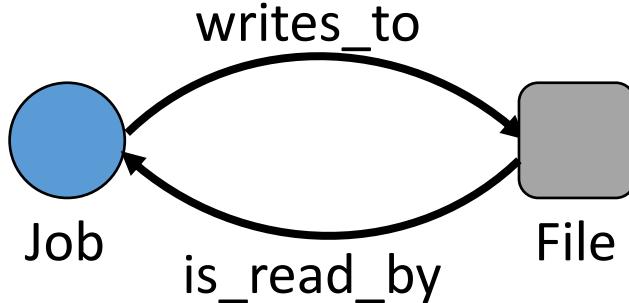
```
edge(1, 2). edge(1, 3). edge(2, 4). edge(3, 4).
```

Graph view enumerators (*templates*) are encoded as logical clauses that are **satisfiable iff view is feasible** given input constraints.

A single template can **enumerate multiple views**, but only most beneficial are chosen, e.g., smaller in size, useful to more than one query, etc.

# Constraint-based view enumeration: Query and schema fact extraction

Schema:



Query:

```
MATCH  
(j1:Job)-[:writes_to]->(f1:File)  
(f1:File)-[r*0..8]->(f2:File)  
(f2:File)-[:is_read_by]->(j2:Job)  
RETURN j1, j2
```

Explicit facts extracted from graph's schema:

```
schemaVertex('Job'). schemaVertex('File').  
schemaEdge('Job', 'File', 'writes_to').  
schemaEdge('File', 'Job', 'is_read_by').
```

Explicit facts extracted from query's graph pattern:

```
queryVertex(j1). queryVertexType(j1, 'Job').  
queryVertex(f1). queryVertexType(f1, 'File').  
queryVertex(f2). queryVertexType(f2, 'File').  
queryVertex(j2). queryVertexType(j1, 'Job').  
queryEdge(j1, f1).  
queryEdgeType(j1, f1, 'writes_to').  
queryEdge(f2, j2).  
queryEdgeType(f2, j2, 'is_read_by').  
queryVariableLengthPath(f1, f2, 0, 8).
```

# Constraint-based view enumeration: Inferring view candidates

Each **view template** yields **set of candidate views** that are feasible given input schema and query constraints.

**Explicit facts extracted from graph's schema:**

```
schemaVertex('Job'). schemaVertex('File').  
schemaEdge('Job', 'File', 'writes_to').  
schemaEdge('File', 'Job', 'is_read_by').
```

**Explicit facts extracted from query's graph pattern:**

```
queryVertex(j1). queryVertexType(j1, 'Job').  
queryVertex(f1). queryVertexType(f1, 'File').  
queryVertex(f2). queryVertexType(j1, 'File').  
queryVertex(j2). queryVertexType(j1, 'Job').  
queryEdge(j1, f1).  
queryEdgeType(j1, f1, 'writes_to').  
queryEdge(f2, j2).  
queryEdgeType(f2, j2, 'is_read_by').  
queryVariableLengthPath(f1, f2, 0, 8).
```



**Feasible assignments for example view template  
"kHopConnectorSameVertexType":**

```
(X='j1', Y='j2', XTYPE='Job', K=2)  
(X='j1', Y='j2', XTYPE='Job', K=4)  
(X='j1', Y='j2', XTYPE='Job', K=6)  
(X='j1', Y='j2', XTYPE='Job', K=8)  
(X='j1', Y='j2', XTYPE='Job', K=10)
```

# Cost Model for Graph Views: Structural Properties Guide View Size Estimation

Number of potential  
k-length paths

$$\hat{E}(G, k) = \binom{n}{k+1} \cdot \left[ \frac{m}{\binom{n}{2}} \right]^k$$

Example: expected number of k-length simple paths in an Erdos-Renyi random graph

# Cost Model for Graph Views: Structural Properties Guide View Size Estimation

$$\hat{E}(G, k) = \binom{n}{k+1} \cdot \left[ \frac{m}{\binom{n}{2}} \right]^k$$

Probability that it  
is an actual path

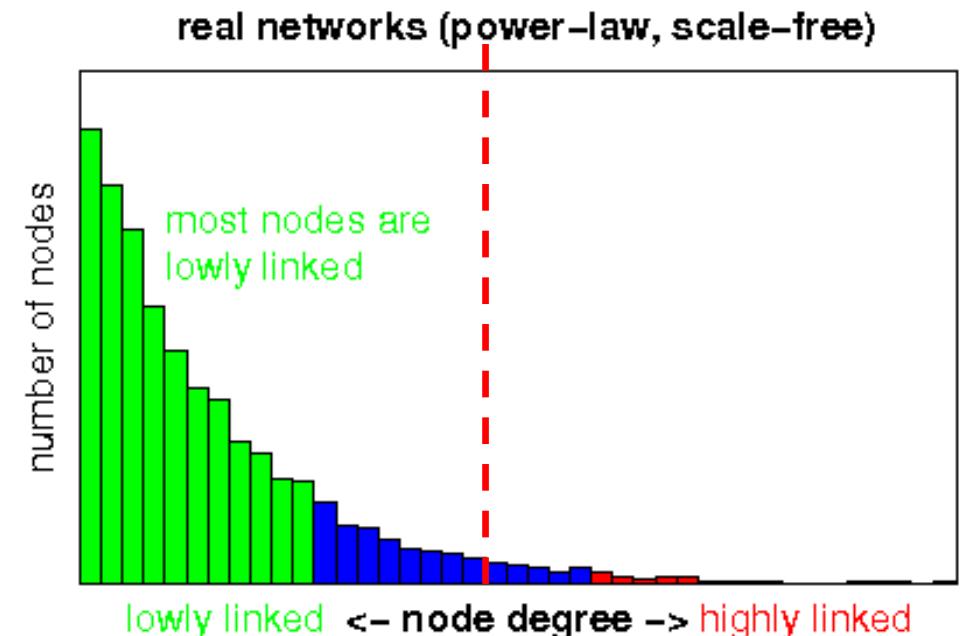
Example: expected number of k-length simple paths in an Erdos-Renyi random graph

# Cost Model for Graph Views: Structural Properties Guide View Size Estimation

Parameterized by percentile out-degree

$$\hat{E}(G, k, \alpha) = n \cdot \deg_{\alpha}^k$$

Example: estimating number of directed  
k-length paths in *homogeneous* graph



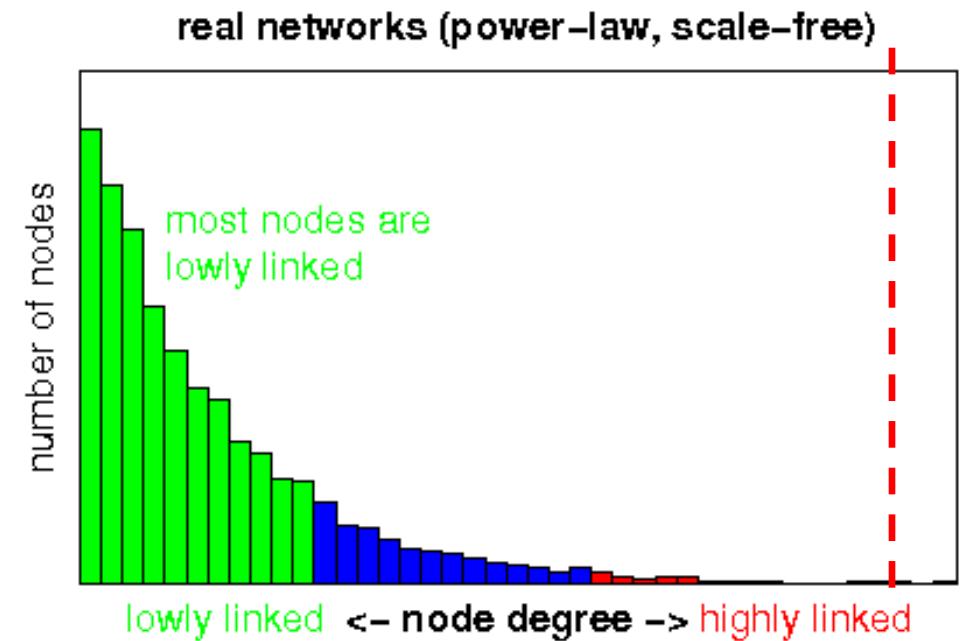
e.g.,  $\alpha = 50$  (median)

# Cost Model for Graph Views: Structural Properties Guide View Size Estimation

Parameterized by percentile out-degree

$$\hat{E}(G, k, \alpha) = n \cdot \deg_{\alpha}^k$$

Example: estimating number of directed  
k-length paths in *homogeneous* graph



e.g.,  $\alpha = 95$  (95<sup>th</sup> %ile)

# Cost Model for Graph Views: Structural Properties Guide View Size Estimation

$$\hat{E}(G, k, \alpha) = \sum_{t \in T_G} n_t \cdot (\deg_\alpha(n_i))^k$$

Example: size estimator for k-hop connector views

# Cost Model for Graph Views: Structural Properties Guide View Size Estimation

$$\hat{E}(G, k, \alpha) = \sum_{t \in T_G} n_t \cdot (\deg_\alpha(n_i))^k$$

For each type of vertex  
that can be an edge source

Example: size estimator for k-hop connector views

# Cost Model for Graph Views: Structural Properties Guide View Size Estimation

How many vertices of that  
type do we have?

$$\hat{E}(G, k, \alpha) = \sum_{t \in T_G} n_t \cdot (\deg_\alpha(n_i))^k$$

Example: size estimator for k-hop connector views

# Cost Model for Graph Views: Structural Properties Guide View Size Estimation

$$\hat{E}(G, k, \alpha) = \sum_{t \in T_G} n_t \cdot (\deg_\alpha(n_i))^k$$

Structural Property:  
out-degree percentile

Example: size estimator for k-hop connector views

# Cost Model for Graph Views: Structural Properties Guide View Size Estimation

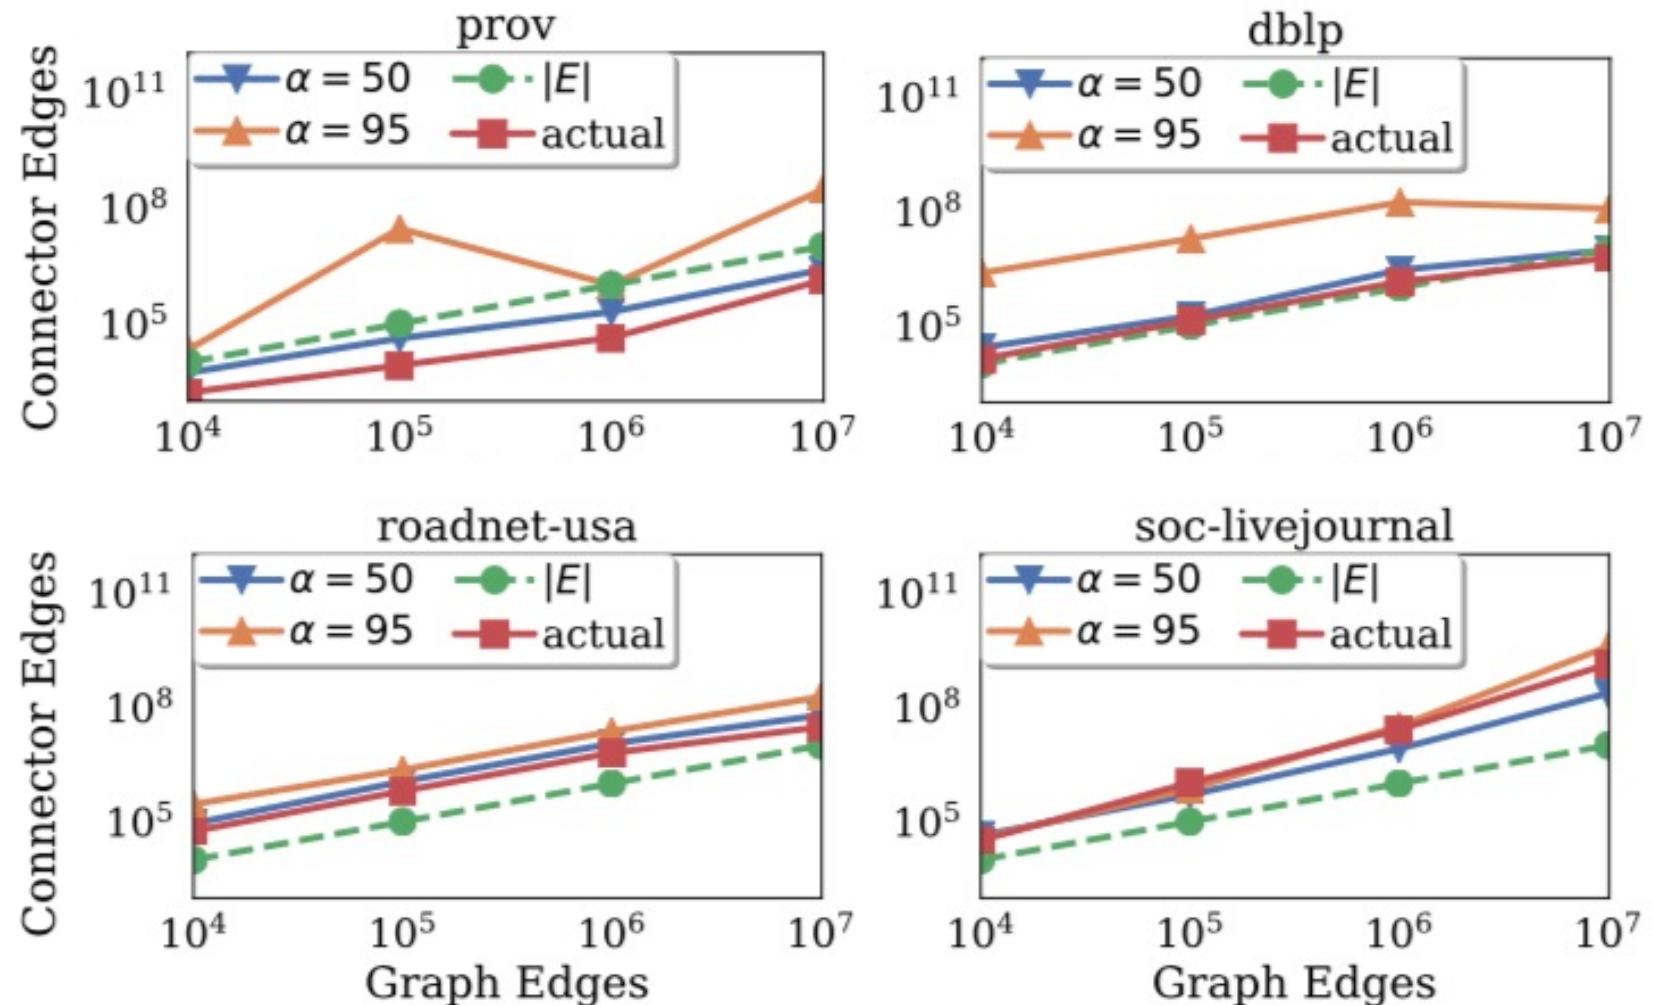
$$\hat{E}(G, k, \alpha) = \sum_{t \in T_G} n_t \cdot (\deg_\alpha(n_i))^k$$

Number of likely  
k-length paths in G

Example: size estimator for k-hop connector views

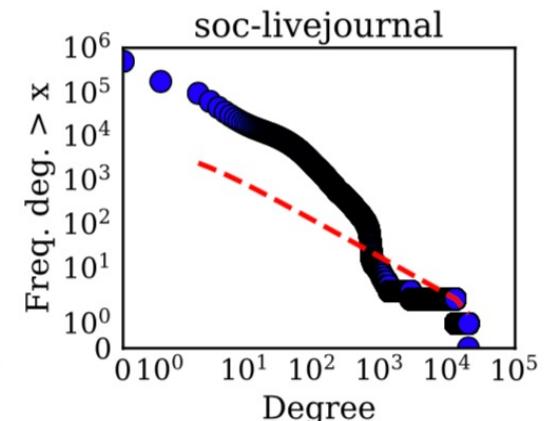
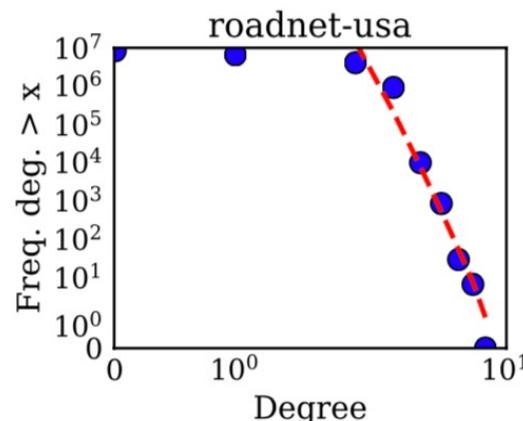
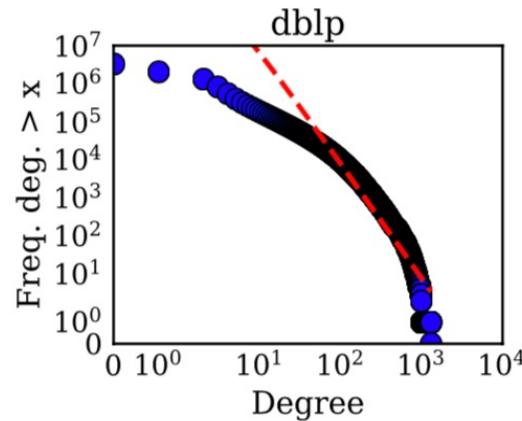
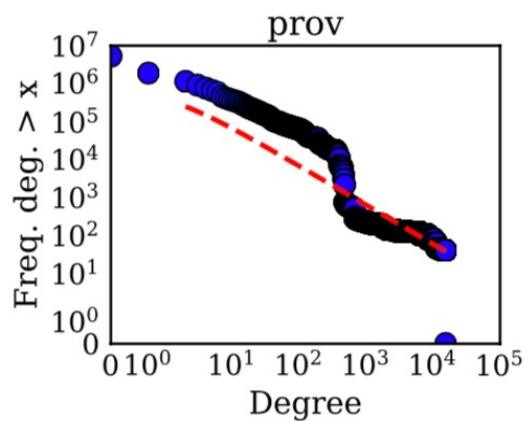
# Cost Model: K-Hop Connector Size Estimator

We find that  $\alpha = 95$  provides a reasonable upper bound



# Datasets

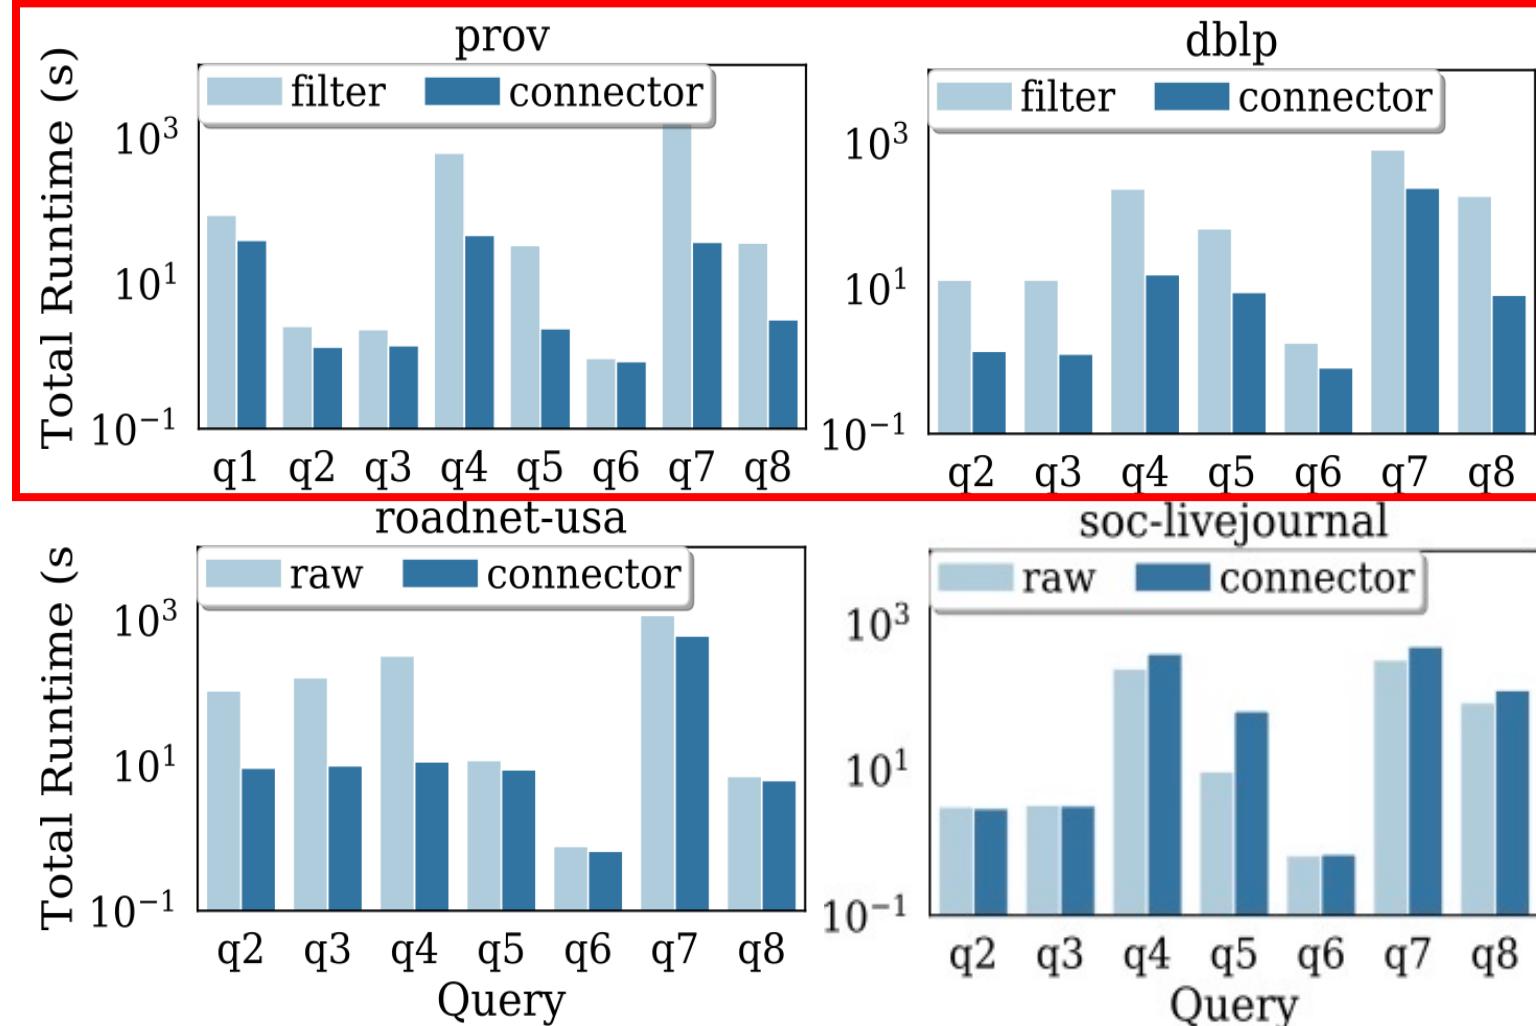
Short Name	Type	No. Vertices	No. Edges
prov (raw)	Data lineage	3.2 Billion	16.4 Billion
prov (summarized)	Data lineage	7 Million	34 Million
dblp-net	Co-authorship	5.1 Million	24.7 Million
soc-livejournal	Social network	4.8 Million	68.9 Million
roadnet-usa	Road network	23.9 Million	28.8 Million



# Query Workload

Query	Operation Type	Result Type
Q1: Job Blast Radius	Retrieval	Subgraph
Q2: Ancestors	Retrieval	Set of Vertices
Q3: Descendants	Retrieval	Set of Vertices
Q4: Path lengths	Retrieval	Bag of scalars
Q5: Edge Count	Retrieval	Single scalar
Q6: Vertex Count	Retrieval	Single scalar
Q7: Community Detection	Update (property)	N/A
Q8: Largest Community	Retrieval	Subgraph

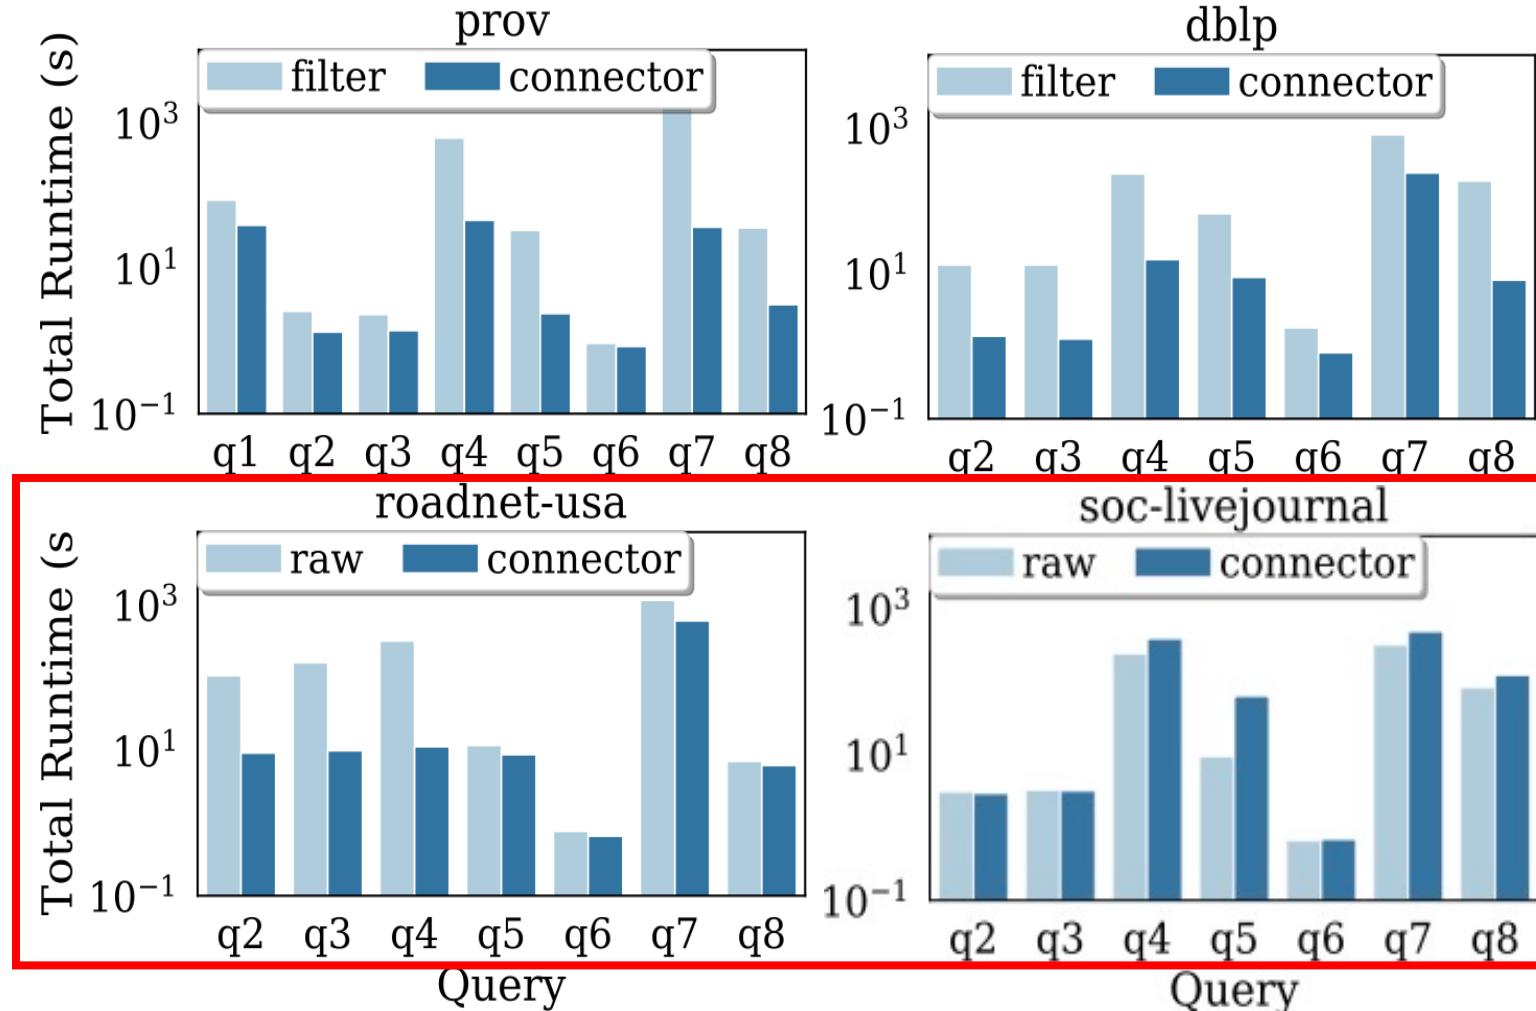
# Query runtimes for rewrites over summarizer and and over 2-hop connector views



- Q1: Job Blast Radius
- Q2: Ancestors
- Q3: Descendants
- Q4: Path lengths
- Q5: Edge Count
- Q6: Vertex Count
- Q7: Community Detection
- Q8: Largest Community

*Highly-structured graphs:  
virtually every query  
benefits from views*

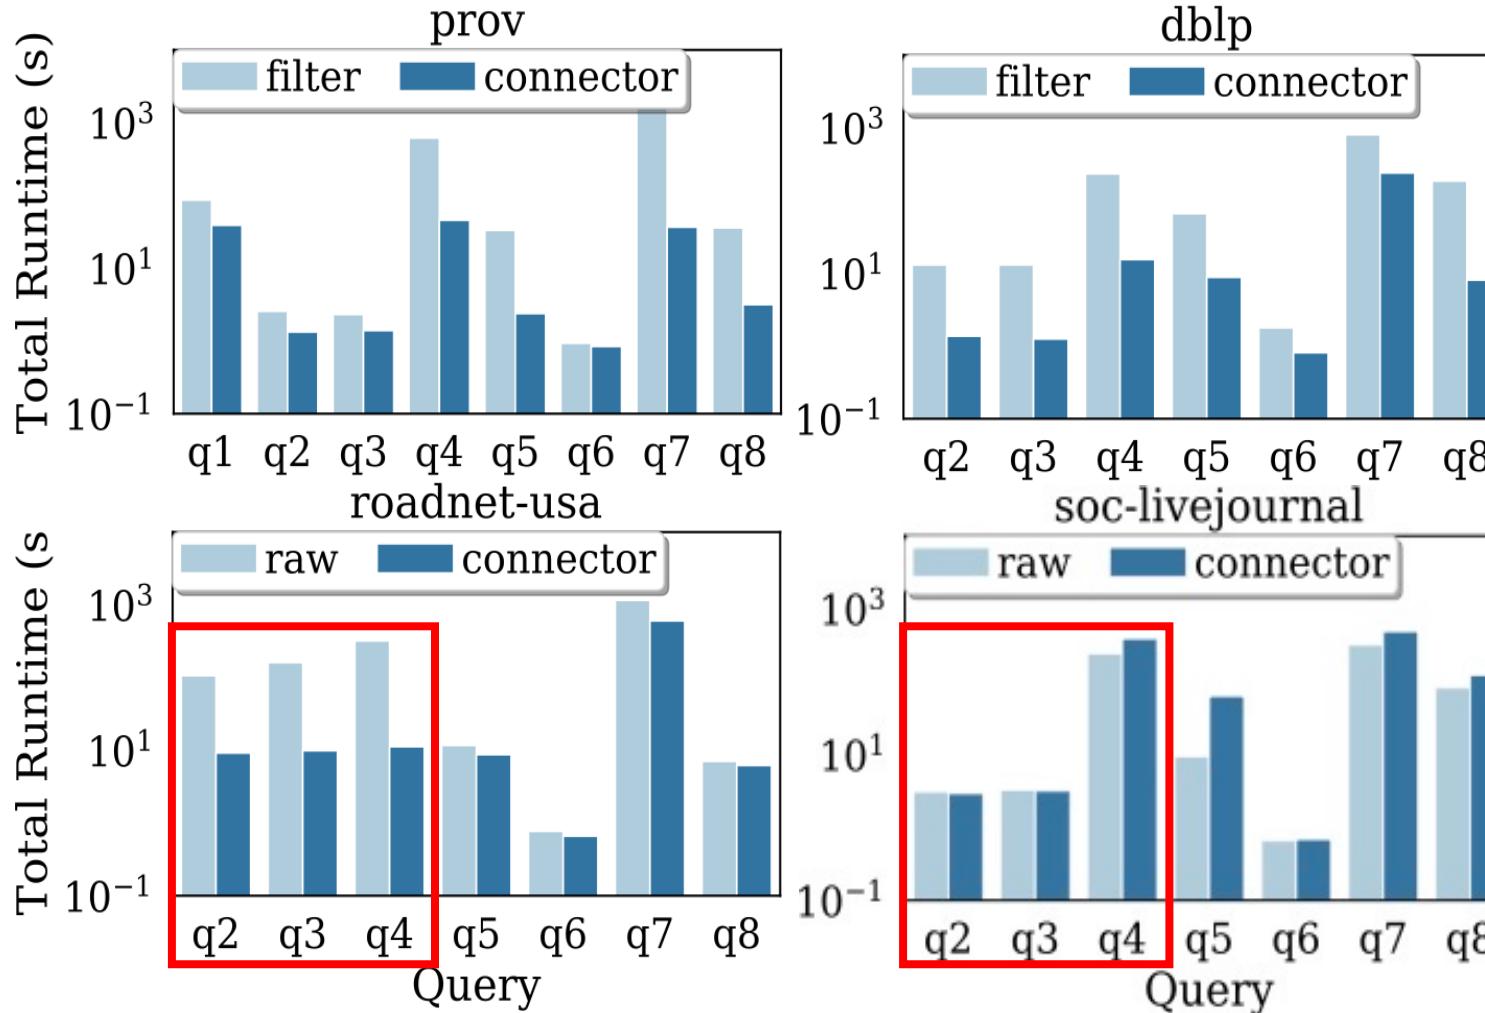
# Query runtimes for rewrites over summarizer and and over 2-hop connector views



- Q1: Job Blast Radius
- Q2: Ancestors
- Q3: Descendants
- Q4: Path lengths
- Q5: Edge Count
- Q6: Vertex Count
- Q7: Community Detection
- Q8: Largest Community

*Less structured graphs:  
vertex-to-vertex connector is  
larger than original graph.*

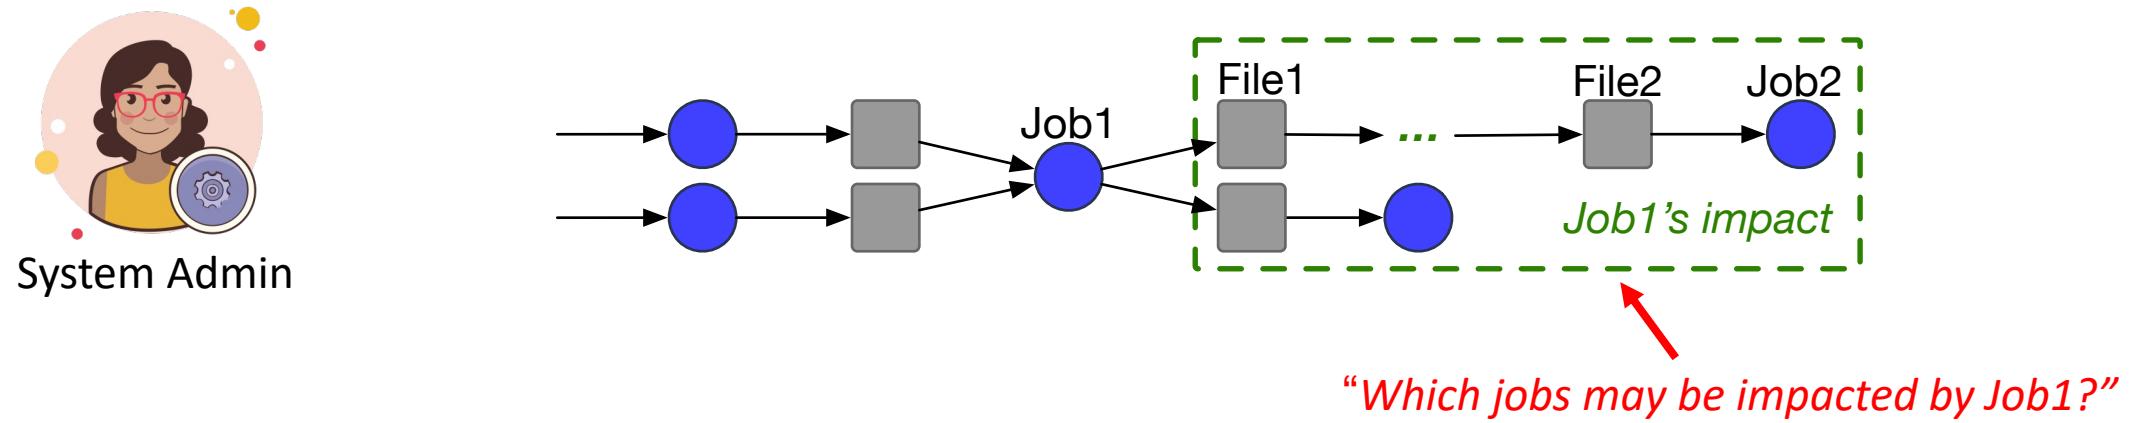
# Query runtimes for rewrites over summarizer and and over 2-hop connector views



- Q1: Job Blast Radius
- Q2: Ancestors
- Q3: Descendants
- Q4: Path lengths
- Q5: Edge Count
- Q6: Vertex Count
- Q7: Community Detection
- Q8: Largest Community

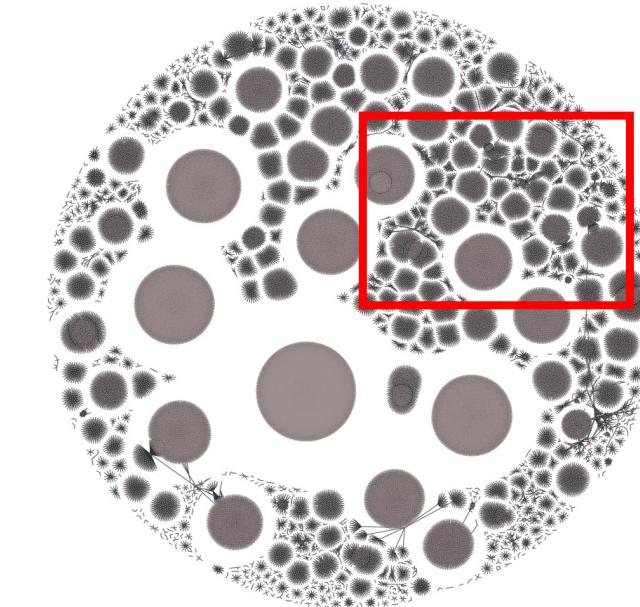
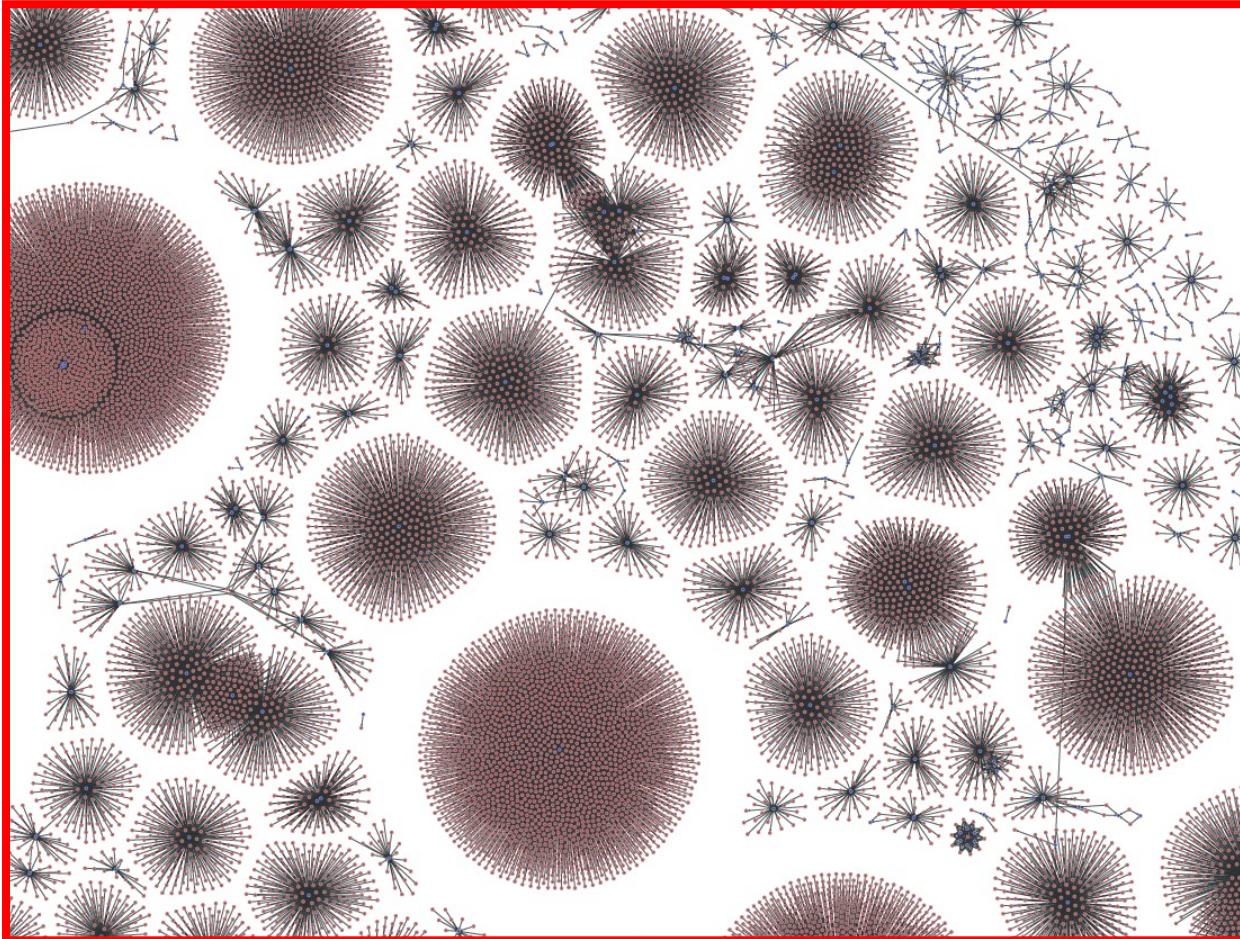
*Non-intuitive result due to larger fraction of long paths in road network.*

# K-hop path query example: Job “blast radius” over machine generated lineage graphs



Several incident response use cases (e.g., identifying privacy violations) in datacenters consist of answering  $k$ -hop path queries.

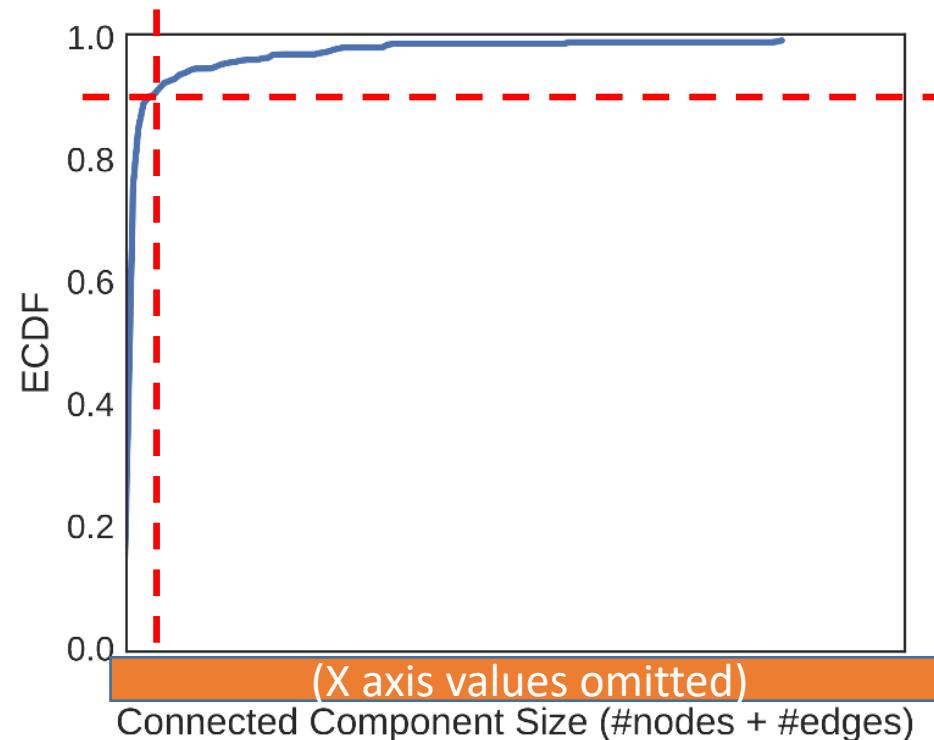
# Real-world graphs are skewed (degree distr.)



Sample from Microsoft prod cluster data lineage graph: each CC is a pipeline

# Real-world graphs are skewed (degree distr.)

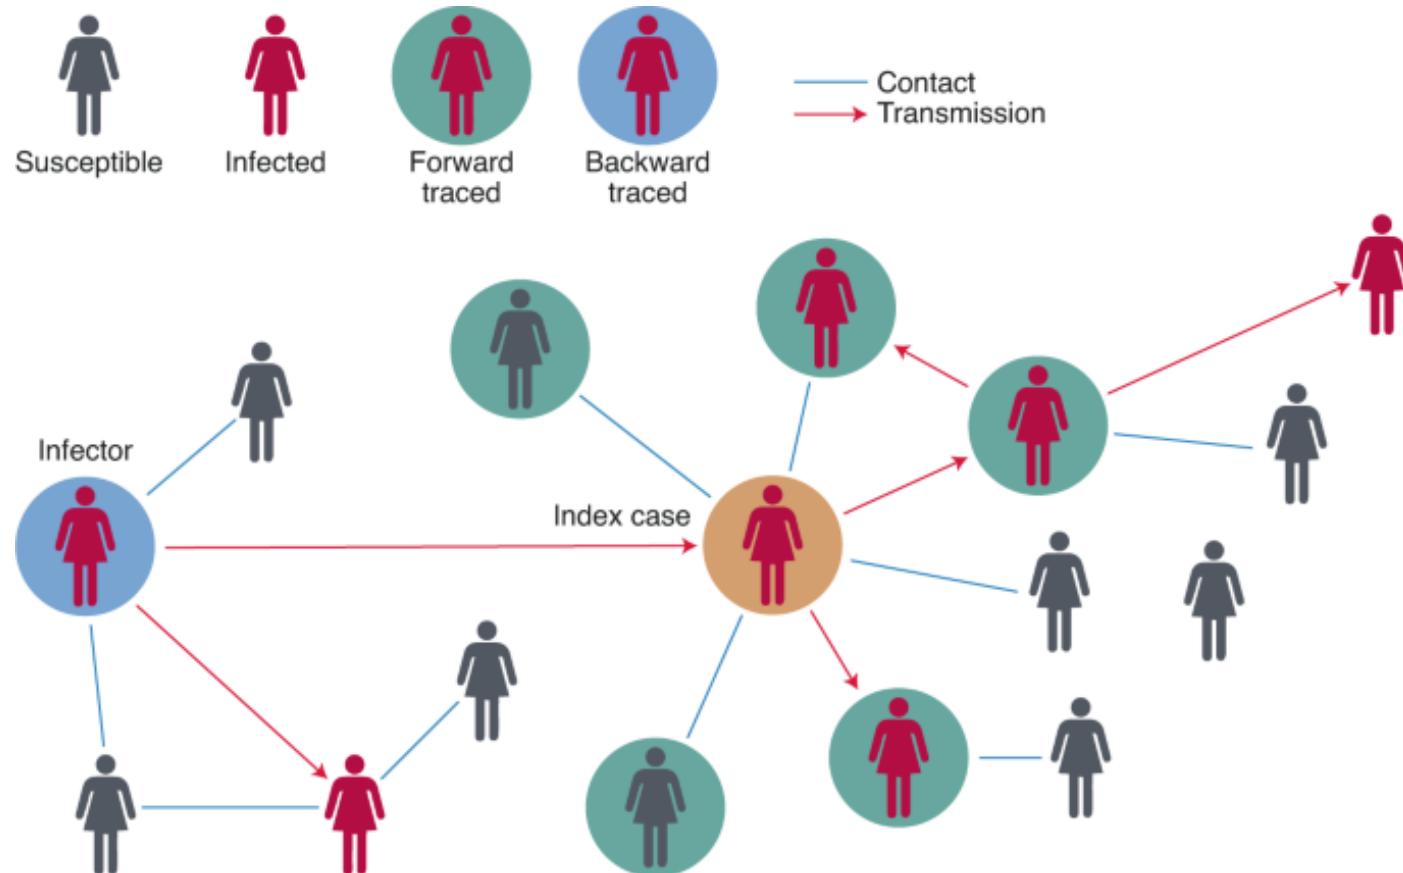
90<sup>th</sup> %ile CC size is  
orders of magnitude  
smaller than largest CC



*e.g., a single job produced output consumed by almost all pipelines*

Sample from Microsoft prod cluster data lineage graph: each CC is a pipeline

# Why do we care about temporal graphs? Certain questions answered incorrectly without time info



*Direction is incorrect if edge start time is missing.*

*False positives / negatives if edge duration is missing.*

Is TGER sufficient for retrieving vertex's neighbors in all temporal walks of interest?

